
Evaluation of Machine Learning in Empirical Asset Pricing

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Several recent studies have claimed that machine learning methods provide superior
2 predictive accuracy of asset returns, relative to simpler modelling approaches, and
3 can correctly identify factors needed to price portfolio risk. Herein, we demonstrate
4 that this performance is critically dependent on several features of the data being
5 analyzed; including, the training/test sample split, the frequency at which the data
6 is observed, and the chosen loss-function. In contrast to existing studies, which
7 claim that neural nets provide superior predictive accuracy, through a series of
8 realistic examples that mimics the stylized facts of asset returns, we demonstrate
9 that neural methods are easily outperformed by simpler methods, such as random
10 forests and elastic nets.

11 1 Introduction

12 The dominance of machine learning (hereafter, ML) methods in terms of predictive accuracy has
13 begun to filter into the empirical asset pricing literature. Arguably, the most common applications
14 of ML methods in empirical finance are for portfolio construction, asset price prediction, and factor
15 selection.

16 Several studies have now used ML techniques to analyze the cross-section of asset returns and
17 produce portfolios that can capture nonlinear information in the cross-section of asset returns. [15]
18 use tree-based methods to understand which firm-level characteristics best predict the cross-section
19 of stock returns, and use this information to help mitigate portfolio risk. Similarly, [13] uses deep
20 feedforward neural nets (DFNs) to construct portfolios and predict returns across a cross-sections
21 of US asset returns. However, while [13] demonstrates that DFNs can better capture nonlinear
22 information, no claim is made that deep learning methods are the best approach to exploit this
23 information.

24 Several studies have now suggested that ML methods can produce better predictions of asset returns
25 ([7], [10] and [5]). The results of [7] suggest that, in terms of predictive performance, as measured by
26 an out-of-sample R^2 , tree-based methods and shallow neural nets can provide superior predictive
27 accuracy over other ML methods and simpler model-based approaches.

28 Similarly, [12], [6], [4] and [16] demonstrate that ML methods can “systematically evaluate the
29 contribution to asset pricing of any new factor” used within an existing linear asset pricing structure.
30 As such, these authors argue that ML can be used, *en masse*, to consistently evaluate the ability of
31 various factors to help price portfolio risk. Such work is particularly pertinent given the literature’s
32 obsession with constructing such factors: [8] currently document that well over 600 different factors
33 have been suggested in the literature.

34 The above studies all demonstrate the potential benefits of ML methods within empirical finance.
35 However, it is unclear if the above findings generalize to different training and validation periods;

different sampling frequencies; and different loss-measures of predictive accuracy. The answer to such questions in the realm of empirical finance are particularly pertinent given that certain ML methods have known difficulties in dealing with data that display the stylized facts of asset returns, e.g., weak and nonlinear dependence, low signal-to-noise and a lack of conditional independence/sparsity. Moreover, training even standard types of neural networks, such as DFNs, becomes particularly difficult when data displays strong, or nonlinear, dependence ([1]).

In many ways, existing applications of ML to empirical finance have either over-looked, downplayed, or simply ignored the importance of the above issues. [13] and [5] use cross validation as part of their model building procedures, destroying the temporal ordering of data. [7] and [13] produce models using training samples that end much earlier than the data sets which they ultimately produce forecasts. This is particularly worrying as the factors driving returns can be starkly different across different time periods.

The goal of this paper is to provide a systematic, and reproducible study on the ability of ML methods to 1) accurately detect significant factors; and 2) accurately predict returns according to a range of loss measures. It is our belief that any such study is necessary in order for practitioners to reliably apply these methods in their particular problems of interest.

After giving the general setup in Section two, in Section three we conduct a rigorous study that gives an in-depth comparison of several ML methods used in the empirical finance literature. The analysis demonstrates that persistence in features, and different complexities of the return generating process affect ML method's ability to: 1) accurately predict future returns across a range of loss measures; and 2) correctly identify the significant factors driving returns. In contrast to existing findings, in this realistic simulation design, we find that neural network procedures, such as feedforward nets, LSTM, and DeepAR models ([17]), are among the worst performing methods, while simpler tree-based methods and elastic net are among the best performing methods.

In Section four, the above findings are validated in an empirical exercise that considers individual returns data from CRSP for all firms listed in the NYSE, AMEX and NASDAQ over a 60 year period, where a set of 549 possible factors are used to explain the cross-section of returns. Careful attention is given to the training and test split, with only use the last fourteen years of returns data used to evaluate the different ML methods. Across all ML methods considered, neural net based procedure perform the worst, while tree-based methods and elastic net perform the best.

Our results suggest that the efficacy of ML methods in empirical finance depends on several features of the underlying problem, such as sampling frequency, the particular training test split, and the data period under analysis. As such, while potentially useful, ML methods are not a panacea for predicting, or understanding the factors that drive financial returns.

2 Model and Methods

2.1 Statistical Model

We briefly discuss the statistical model considered for asset returns. Excess monthly returns on asset i , $i \leq n$, at time t , $t \leq T$, are assumed to evolve in an additive fashion:

$$r_{i,t+1} = E(r_{i,t+1}|\mathcal{F}_t) + \epsilon_{i,t+1}, \quad E(\epsilon_{i,t+1}|\mathcal{F}_t) = 0, \quad (1)$$

where \mathcal{F}_t denotes the observable information at time t , and $\epsilon_{i,t+1}$ is a martingale difference sequence. The conditional mean of returns is an unknown function of a P -dimensional vector of features, measurable at time t :

$$E(r_{i,t+1}|\mathcal{F}_t) = g(z_{i,t}). \quad (2)$$

The features, or predictors, $z_{i,t}$ are composed of time- t information, and only depend on the characteristics of stock i . The assumption that the information set can be characterized by the variables $z_{i,t}$, without dependence on the $j \neq i$ return units, is reasonable if the collection of $z_{i,t}$ is rich enough.

In what follows, we represent the space of possible features as the Kronecker product of two pieces

$$z_{i,t} = x_t \otimes c_{i,t}, \quad (3)$$

where the variables $c_{i,t}$ represent a $P_c \times 1$ vector of individual-level characteristics for return i , and x_t represents a $P_x \times 1$ vector of macroeconomic predictors, and \otimes represents the Kronecker product.

Thus, for $P = P_c \cdot P_x$, $z_{i,t}$ represents a $P \times 1$ feature space over which we approximate the unknown function $g(\cdot)$.

2.2 Methods to be compared

Given features $z_{i,t}$, the goal of any ML method is to approximate the unknown function $g(\cdot)$ in 1. Broadly speaking, how different ML methods choose to approximate this function depends on three components:

1. the model used to make predictions;¹
2. the regularization mechanism employed to mitigate over-fitting;
3. a loss function that penalized poor predictions.

To ensure the results of ML different methods will be comparable, we fix both the regularization mechanisms and loss functions used within each method, and allow only the models used for prediction to vary. This approach seeks to ensure that performances in one method, relative to another, are based on the model structure and not to some feature of how the models were fit. To this end, we first discuss points 2. and 3. above, and then briefly present the models used for our comparison.

Loss functions: We implement all ML methods using two loss functions: Mean Absolute Error (MAE) and Mean Squared Error (MSE): for $\hat{r}_{i,j}$ denoting the predicted return on asset i at time j ,

$$\text{MAE} = \frac{1}{n} \sum_{j=i}^n |r_{i,j} - \hat{r}_{i,j}| \text{ and } \text{MSE} = \frac{1}{n} \sum_{j=i}^n (r_{i,j} - \hat{r}_{i,j})^2.$$

We consider both loss functions since MAE is less sensitive to outliers in the data which financial returns are known to exhibit, and which are due to extreme market movements. Given this, we expect MAE to produce predictive results that are more robust to such outlier events.

Sample Splitting: Since returns are intrinsically dependent, observed data is split into “training”, “validation” and “test” sets according to a schema that respects this dependence structure. To balance computation and accuracy, we use a hybrid “rolling window” and “recursive” approach to training/validation/test splits: for each model refit, the training set is increased by one year of observations, i.e., 12 monthly observations; the validation set is fixed at one year and moves forward (by one year) with each model refit; predictions are generated using that model for the subsequent year.

Models: In what follows we compare a host of different ML models including elastic net ([19], random forest ([2]), feed-forward neural nets, LSTM, FFORMA ([14]) and DeepAR models ([17]). Details on each model and certain features of its implementation used in this work are given in Appendix A. For each of the different methods, we consider two variants, one based on the MAE loss and one based on the MSE loss.

2.3 Model evaluation measures

Predictive accuracy: Predictive performance is assessed using Mean Absolute Error (MAE), Mean Squared Error (MSE) (evaluated over the test set) and an out-of-sample R^2 measure. While out-of-sample R^2 is a common measure, there is no universally agreed-upon definition. As such, we explicitly state the version employed herein as

$$R_{OOS}^2 = 1 - \frac{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \hat{r}_{i,t+1})^2}{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \bar{r}_{i,t+1})^2}, \quad (4)$$

where \mathcal{T}_3 indicates that the fits are only assessed on the test sub-sample, which is never used for training or tuning.

¹The model used by the ML method need not correspond to the statistical model assumed to describe the data. Herein, our goal will not be to asses the “accuracy” of the statistical model, but to determine how different ML methods accurately determine the salient features of this model.

121 Since R^2 is based on in-sample-fit of a linear model, this measure is less meaningful for most of the
 122 ML methods considered in in this paper. However, we report this measure since this measure has also
 123 been considered in other applications of ML to empirical finance (see, e.g., [7]).

124 **Factor Selection:** An important aspect of empirical finance is the knowledge of which features
 125 drive risk, i.e., which features are explicitly represented within $z_{i,t}$. To this end, we follow [7] and
 126 construct a variable importance (VI) measure to compare the different ML methods. The importance
 127 of variable j , VI_j , is defined as the reduction in predictive R^2 from setting all values of predictor j
 128 to 0, while holding the remaining model estimates fixed. Each VI_j is then normalized to sum to 1.

129 However, as VI_j can sometimes be negative, we shift VI_j by the smallest VI_j plus a small constant,
 130 then dividing by this sum to alleviate numerical issues². The resulting VI measure is then:

$$VI_{j,norm} = \frac{VI_j + \min(VI_j) + o}{\sum VI_j + \min(VI_j) + o} \quad ; \quad o = 10^{-100} \quad (5)$$

131 3 Preliminary Results

132 We first explore how ML methods perform in terms of prediction and factor selection for data
 133 that exhibit the stylized facts of empirical returns. We simulate data according to a design which
 134 incorporates a low signal-to-noise ratio, stochastic volatility, low persistence in levels and cross-
 135 sectionally correlated features. Data is generated from a latent factor volatility model for excess
 136 returns r_{t+1} , for $t = 1, \dots, T$:

$$\begin{aligned} r_{i,t+1} &= g(z_{i,t}) + \beta_{i,t+1}v_{t+1} + e_{i,t+1}; \quad z_{i,t} = (1, x_t)' \otimes c_{i,t}, \quad \beta_{i,t} = (c_{i1,t}, c_{i2,t}, c_{i3,t}) \\ e_{i,t+1} &= \sigma_{i,t+1}\varepsilon_{i,t+1}; \\ \log(\sigma_{i,t+1}^2) &= \omega + \gamma \log(\sigma_t^2) + \sigma_u u; \quad u \sim N(0, 1) \end{aligned}$$

137 where v_{t+1} is a 3×1 vector of errors, $w_{t+1} \sim N(0, 1)$, $\varepsilon_{i,t+1} \sim N(0, 1)$ scalar error terms, matrix
 138 C_t is an $N \times P_c$ matrix of latent factors, where the first three columns correspond to $\beta_{i,t}$, across the
 139 $1 \leq i \leq N$ dimensions, while the remaining $P_c - 3$ factors do not enter the return equation. The
 140 $P_x \times 1$ vector x_t is a 3×1 multivariate time series that captures for macroeconomic factors, and
 141 ε_{t+1} is a $N \times 1$ vector of idiosyncratic errors. These parameters were tuned such that the annualized
 142 volatility of each return series was approximately 22% when viewed as monthly returns, as is often
 143 observed empirically.

144 We consider three different functions for $g(z_{i,t})$:

$$\begin{aligned} (1) \quad g_1(z_{i,t}) &= (c_{i1,t}, c_{i2,t}, c_{i3,t} \times x_t'[3,]) \theta_0, \\ (2) \quad g_2(z_{i,t}) &= (c_{i1,t}^2, c_{i1,t} \times c_{i2,t}, \text{sgn}(c_{i3,t} \times x_t'[3,])) \theta_0, \\ (3) \quad g_3(z_{i,t}) &= (1[c_{i3,t} > 0], c_{i2,t}^3, c_{i1,t} \times c_{i2,t} \times 1[c_{i3,t} > 0], \text{logit}(c_{i3,t})) \theta_0, \end{aligned}$$

145 where $x_t'[3,]$ denotes the third element of the x_t' vector. $g_1(z_{i,t})$ allows the characteristics to enter
 146 the return equation linearly, and $g_2(z_{i,t})$ and $g_3(z_{i,t})$ allow the characteristics to enter the return
 147 equation interactively and non-linearly.³ The value of θ^0 was tuned such that the predictive R^2 was
 148 approximately 5%.


149 We consider two different levels of cross-sectional correlation for the N factors, $c_{i,t}$, which correspond
 150 to a small amount of 0.01 and a large amount, 1.0, or cross-sectional correlation. The specific details
 151 regarding the level of cross-sectional correlation, and how it is introduced, is given in Appendix
 152 B.1. The macroeconomic factors, x_t , a 3×1 vector, is generated according to a stationary Vector
 153 Autoregression (VAR) model with a high-degree of persistence (0.95 for each series) and a diagonal
 154 coefficient matrix. See Appendix B.1 for more details.

155 The simulation design results in 9 different data generating process (DGP). For each DGP we fix
 156 with $N = 200$ stocks, $T = 180$ time periods and $P_c = 100$ characteristics. Each DGP was simulated
 157 10 times to assess the robustness of ML algorithms, with the number of simulations kept low for
 158 computational feasibility. We employ the hybrid data splitting approach with a training:validation
 159 length ratio of approximately 1.5 and a test set that is 1 year in length.

²This mechanism was chosen because the other popular normalization mechanism "softmax" was observed to be unable to preserve the distances between each original VI_j , making discernment between each VI_j difficult.

³(g_1, g_2 correspond to the simulation design used by [7].)

Figure 1: Simulation variable importance, faceted by simulation specification



3.1 Simulation Study Results

Prediction Performance: For brevity we only remark on the most interesting findings in the main paper, which are given in the table below (the complete set of simulation results are detailed in Appendix B.2). In contrast to existing studies, we find that elastic nets are the best performing model, followed closely by random forests, then neural networks. Interestingly, the ML models were largely unaffected by the level of cross-sectional correlation in terms of prediction performance.

Generally, ML models fitted with respect to minimizing MAE (quantile loss) generally perform better, even when evaluated against MSE loss metrics. Although the actual level difference between the loss metrics across the different methods is small, the results are remarkably consistent across the various Monte Carlo designs.

Table 1: Top Models in Simulation Study

Corr	model	Test MAE			Test MSE		
		g1	g2	g3	g1	g2	g3
0.01	ELN.MAE	0.0345786	0.0361950	0.0353345	0.0025652	0.0026882	0.0026210
	RF.MAE	0.0354594	0.0354204	0.0355399	0.0026434	0.0026305	0.0026446
	NN2.MAE	0.0359604	0.0369206	0.0363047	0.0026786	0.0027474	0.0026996
	NN1.MAE	0.0358939	0.0368335	0.0363352	0.0026718	0.0027396	0.0027028
	NN3.MAE	0.0358164	0.0369345	0.0364712	0.0026697	0.0027491	0.0027181
1	ELN.MSE	0.0346142	0.0362761	0.0354437	0.0025676	0.0026980	0.0026300
	RF.MAE	0.0359158	0.0356434	0.0360529	0.0026747	0.0026445	0.0026786
	NN5.MAE	0.0370087	0.0372705	0.0374132	0.0027744	0.0027832	0.0027916
	NN4.MSE	0.0373820	0.0368966	0.0373542	0.0028051	0.0027505	0.0027970
	NN3.MAE	0.0372849	0.0370382	0.0371925	0.0027940	0.0027652	0.0027753

Factor Importance: The factor importance results are presented graphically in Figure 1, and demonstrates that overall elastic net outperforms all other models consistently in terms of assigning the correct relative importance to the true underlying features.⁴ However, the performance of elastic net does degrade as the data generating process becomes more non-linear.

Random forests, and to a lesser extent neural networks, can also correctly identified the correct underlying regressors, but struggled with adequately discerning relative importance among correlated regressors. This behavior becomes more pronounced as the degree of cross-sectional correlation increases (see decreasing relative importance of true underlying regressors in Figure 1).

⁴(c_1 .constant, c_2 .constant and $c_3.x_3$ for g_1 and g_2 specifications, and c_1 .constant, c_2 .constant and c_3 .constant for g_3)

The overall high performance of the elastic net models is surprising given its relative simplicity compared to other machine learning models. However, when recalling that elastic net models are the only machine learning models which are specifically noted to perform well on datasets with high degrees of multicollinearity, the result is perhaps less surprising. Indeed, random forest and neural network models are only noted to be better at capturing non-linear relationships in independent and identically distributed data, a property which we observe on the non-linear specification, and only for the random forests.

4 Empirical analysis

We now investigate the performance of ML methods across a large sample of asset returns. As we shall see later, the results obtained in Section 3.1 are largely borne out in this empirical exercise.

4.1 Data

We use the universe of firms listed in the NYSE, AMEX and NASDAQ, starting from 1957 (starting date of the S&P 500) and ending in December 2016, totalling 60 years, that have a quarterly return over this period. This approach allows firms to enter and exit the dataset and helps alleviate the problem of survivorship bias in the dataset. Individual cross-sectional factors, $c_{i,t}$, are constructed following the approach of [7]. To alleviate data quality issues within the factor data, from this universe of firms we restrict our dataset (both returns and factors) to begin in 1993 Q3 and end on 2016 Q4. Our individual factor set contains 94 characteristics: 61 updated annually, 13 updated quarterly and 20 updated monthly.⁵⁶ Complete details of the data and the cleaning procedures employed are detailed in Appendix C.1.

Following [18] (see Table 7) we consider eight macroeconomic factors. These factors were lagged by one period so as to be used to predict one period ahead quarterly returns. The 3-month Treasury Bill rate was also used from this source to proxy for the risk-free rate in order to construct excess quarterly returns. The two sets of factors, $c_{i,t}$ and x_t , are then used to build the baseline set of factors, which we defined as in equation (3); i.e., $z_{i,t} = (1, x_t')' \otimes c_{i,t}$. The total number of features in this baseline set is $61 \times (8 + 1) = 549$.

The final dataset contains 202, 066 individual observations. We note that due to data quality issues, in particular, missingness, LSTMs, FFORMA and DeepAR could not be reliably implemented on the empirical data.⁷

We mimic the sample splitting procedure used in the simulation study: the dataset was split such that the training and validation sets were split such that the training set was approximately 1.5 times the length of the validation set, in order to predict a test set that is one year in length.

4.2 Results

Prediction Accuracy The predictive results for the five best methods, according to the various loss measures, are displayed below. In general, the same pattern of results in Section 3.1 is again in evidence: elastic net performs best, followed by the random forests, then the DFNs. We note that the differences between each model using the MSE and MAE loss metrics are much more pronounced for this empirical data. In addition, the ML models perform better when fitted with respect to quantile loss instead of MSE. Most notably, the lack of robustness for the DFNs observed in Section 3.1 is amplified on the empirical dataset, which directly contradicts existing results already reported in the literature.

⁵The dataset also included 74 Standard Industrial Classification (SIC) codes, but these were omitted due to their inconsistency, and inadequateness at classifying companies, as noted by WRDS

⁶To deal with missing data, any characteristics that had over 20% of their data missing were omitted. Remaining missing data were then imputed using their cross sectional medians for each year. See Appendix for more details.

⁷The dataset was not normalized for all methods, as only penalized regression and neural networks are sensitive to normalization. For these two methods, the dataset was normalized such that each predictor column had mean zero and unit variance.


Table 2: Top 5 models in empirical study

model	Sample 1			Sample 2			Sample 3		
	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
ELN_MAE	0.131369	0.040718	0.014306	0.137092	0.041892	0.017875	0.146251	0.045207	0.000835
RF_MAE	0.126703	0.036785	0.109505	0.173721	0.057546	-0.349132	0.14692	0.046037	-0.01752
NN5_MAE	0.146411	0.044901	-0.086967	0.18499	0.06461	-0.514744	0.184986	0.063861	-0.411475
NN4_MAE	0.157301	0.050286	-0.217308	0.168815	0.055711	-0.306102	0.167998	0.055129	-0.218463
NN3_MAE	0.140781	0.042832	-0.036882	0.181096	0.06216	-0.4573	0.164896	0.053458	-0.181528

219 That being said, we do observe some evidence that deeper neural networks perform better, though
220 this result is less apparent due to the lack of robustness of these methods on empirical data (see C.2
221 in Appendix C for results).

222 **Factor Importance** As the data generating process for empirical returns is unknown, the variable
223 importance results cannot be directly compared with those of the simulation study. Even so, we
224 see similar results: the elastic net and random forest models tend to agree on the same subset of
225 predictors. Similar to the prediction performance results, neural networks perform poorly.

Figure 2: Empirical individual and macroeconomic factor importance, averaged over all samples



Individual factors shown on x axis (see Table C.1 in Appendix for definitions)

226 The elastic net, random forest and to a lesser extent DFNs tend to pick out the max return and 1 month
227 momentum factors out of the individual characteristics as important, and the book-to-market factor
228 out of the macroeconomic factors are important. In general, the variable importance metrics are less
229 consistent for the random forests, and it should be noted in particular that the random forest tends
230 to determine factors highly correlated with momentum as important, such as change in momentum,
231 dollar trading volume and return volatility. Within the macroeconomic factors, penalized linear
232 models tend to identify the average book to market ratio and the default spread as the most important.

233 The overall results of this analysis again question existing results already reported in the literature,
234 which conclude that all ML methods tend to agree on the same subset of important factors (see, e.g.,
235 [7]). In our context, we see, at best, only mild agreement between the various ML methods in regards
236 to individual factor selection.

237 Interestingly, the linear models assign the controversial dividend price ratio macroeconomic factor
238 as highly important, a result which was somewhat mirrored only with the neural networks. Their
239 variable importance for individual factors across different training samples is non-robust, with the
240 important variables almost completely changing year to year.

241 Before concluding this section, we note that a more comprehensive set of empirical results, including
242 robustness checks, are given in Appendix C.2, with the majority of the results presented in Appendix
243 C.2 being omitted due to space restrictions. That being said, the findings of this expanded analysis
244 are perfectly congruent to the results presented in this section.

5 Conclusion

Our findings demonstrate that the field of ML may offer certain tools to improve stock prediction and identification of underlying factors. This study suggest that penalized linear models and random forests are the most robust methods for data displaying the stylized facts of asset returns. In contrast to existing results, we find that DFNs fail in the context of return prediction, and variable importance analysis. This result is consistent across a variety of simulated data sets, as well as empirical data.

Therefore, the overall findings of this research differs from the sparse literature on ML methods in empirical finance. However, the performance of the penalized linear models with respect to both out of sample prediction performance and variable importance analysis is promising, and our findings show that ML provides some tools which may aid in the problems of stock return prediction and risk factor selection in the financial world.

Broader Impact

This research calls into question the broad applicability of machine learning methods within empirical finance, at least in the context of return prediction and factor selection. In contrast to existing studies, we find that more complex machine learning methods, such as deep feedforward neural nets, do not perform as well as simpler penalized linear methods and random forest. As such, this research suggests that ML methods are not a panacea for empirical finance, and that great care and diligence is needed in the application of these methods within any financial decision making process.

References

- [1] Bengio, Y., Simard, P., and Frasconi, P. (1994). Learning long-term dependencies with gradient descent is difficult. *IEEE transactions on neural networks*, 5(2):157–166.
- [2] Breiman, L. (2001). Random forests. *Machine learning*, 45(1):5–32.
- [3] Fama, E. F. and French, K. R. (2015). A five-factor asset pricing model. *Journal of financial economics*, 116(1):1–22.
- [4] Feng, G., Giglio, S., and Xiu, D. (2019). Taming the Factor Zoo: A Test of New Factors. Technical Report w25481, National Bureau of Economic Research, Cambridge, MA.
- [5] Feng, G., He, J., and Polson, N. G. (2018). Deep Learning for Predicting Asset Returns. *arXiv:1804.09314 [cs, econ, stat]*. arXiv: 1804.09314.
- [6] Freyberger, J., Neuhierl, A., and Weber, M. (2017). Dissecting characteristics nonparametrically. Technical report, National Bureau of Economic Research.
- [7] Gu, S., Kelly, B. T., and Xiu, D. (2019). Empirical Asset Pricing via Machine Learning. SSRN Scholarly Paper ID 3159577, Social Science Research Network, Rochester, NY.
- [8] Harvey, C. R. and Liu, Y. (2019). A Census of the Factor Zoo. *Social Science Research Network*, page 7.
- [9] Hochreiter, S. and Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation*, 9(8):1735–1780.
- [10] Hsu, J. and Kalesnik, V. (2014). Finding smart beta in the factor zoo. *Research Affiliates (July)*.
- [11] Ishwaran, H., Kogalur, U. B., Blackstone, E. H., and Lauer, M. S. (2008). Random survival forests. *Annals of Applied Statistics*, 2(3):841–860.
- [12] Kozak, S., Nagel, S., and Santosh, S. (2017). Shrinking the cross section. Technical report, National Bureau of Economic Research.
- [13] Messmer, M. (2017). Deep Learning and the Cross-Section of Expected Returns. SSRN Scholarly Paper ID 3081555, Social Science Research Network, Rochester, NY.

- [14] Montero-Manoso, P., Athanasopoulos, G., Hyndman, R. J., and Talagala, T. S. (2020). FFORMA: Feature-based forecast model averaging. *International Journal of Forecasting*, 36(1):86–92.
- [15] Moritz, B. and Zimmermann, T. (2016). Tree-Based Conditional Portfolio Sorts: The Relation between Past and Future Stock Returns. SSRN Scholarly Paper ID 2740751, Social Science Research Network, Rochester, NY.
- [16] Rapach, D. and Zhou, G. (2013). Forecasting Stock Returns. In *Handbook of Economic Forecasting*, volume 2, pages 328–383. Elsevier.
- [17] Salinas, D., Flunkert, V., Gasthaus, J., and Januschowski, T. (2019). DeepAR: Probabilistic forecasting with autoregressive recurrent networks. *International Journal of Forecasting*.
- [18] Welch, I. and Goyal, A. (2008). A Comprehensive Look at The Empirical Performance of Equity Premium Prediction. *Review of Financial Studies*, 21(4):1455–1508.
- [19] Zou, H. and Hastie, T. (2005). Regularization and variable selection via the Elastic Net. *Journal of the Royal Statistical Society, Series B*, 67:301–320.

A Additional details: models

In this section, we give a brief overview of all the models considered in the simulation and empirical study.

A.1 Linear models

Linear models model the conditional expectation $g^*(z_{i,t})$ as a linear function of the predictors and the parameter vector θ :

$$g(z_{i,t}; \theta) = z'_{i,t} \theta \quad (6)$$

This yields the OLS estimator when optimized with respect to MSE, and the LAD estimator when optimized with respect to MAE.

A.2 Elastic nets

Elastic Nets are similar to linear models but differ via the addition of a penalty term in the loss function:

$$\mathcal{L}(\theta; \cdot) = \underbrace{\mathcal{L}(\theta)}_{\text{Loss Function}} + \underbrace{\phi(\theta; \cdot)}_{\text{Penalty Term}} \quad (7)$$

where the elastic net penalty [19] is:

$$\phi(\theta; \lambda, \rho) = \lambda(1 - \rho) \sum_{j=1}^P |\theta_j| + \frac{1}{2} \lambda \rho \sum_{j=1}^P \theta_j^2 \quad (8)$$

Further details are given in [19].

A.3 Random forests

Random Forests are an extension of Classification and Regression Trees (CART) proposed by [2] (see for more comprehensive details). CART are fully non-parametric models that can capture complex multi-way interactions. A tree "grows" in a series of iterations. With each iteration, a split ("branch") is made along one predictor such that it is the best split available at that stage with respect to minimizing the loss function. These steps are continued until each observation is its own node, or more commonly until the stopping criterion is met. The eventual model slices the predictor space into rectangular partitions, and predicts the unknown function $g^*(z_{i,t})$ with the average value of the outcome variable in each partition. The prediction of a tree, \mathcal{T} , with K "leaves" (terminal nodes), and depth L is

$$g(z_{i,t}; \theta, K, L) = \sum_{k=1}^K \theta_k \mathbf{1}_{z_{i,t} \in C_k(L)} \quad (9)$$

where $C_k(L)$ is one of the K partitions in the model. For this study, only recursive binary trees were considered. Though trees were originally proposed and fit with respect to minimizing MSE, they can be grown with respect to a variety of loss functions, where the loss within each C partition is denoted by $H(\theta, C)$:

$$H(\theta, C) = \frac{1}{|C|} \sum_{z_{i,t} \in C} L(r_{i,t+1} - \theta) \quad (10)$$

where $|C|$ denotes the number of observations in set C (partition). Given C , it is clear that the optimal choice for minimising the loss is simply the average of the partition for MSE, and the median of the partition for MAE.

A.4 Feed forward neural networks

A feed forward neural network consists of layers denoted by $l = 0, 1, \dots, L$, with $l = 0$ denoting the input layer and $l = L$ denoting the output layer. The input layer is defined by the scaled predictor set, $x^{(0)} = (1, z_1, \dots, z_N)'$. The model adds complexity through the use of one or more hidden layer, each containing $K^{(l)}$ "neurons". Each neuron linearly aggregates the values of the previous layer, and applies some non-linear "activation function" which we denote as α to its aggregated signal before sending its output to the next layer. The output of neuron k in layer l is then $x_k^{(l)}$. Next, define the vector of outputs for this layer as $x^{(l)} = (1, x_1^{(l)}, \dots, x_{K^{(l)}}^{(l)})'$. The recursive output formula for the neural network at each neuron in layer $l > 0$ is then:

$$x_k^{(l)} = \alpha(x^{(l-1)'} \theta_k^{l-1}), \quad (11)$$

where $\alpha()$ represents the activation function for that layer with the final output ⁸

$$g(z; \theta) = x^{(L-1)'} \theta^{L-1} \quad (12)$$

The neural network's weight and bias parameters for each layer are estimated by minimizing the loss function with respect to the parameters, i.e. by calculating the partial derivative with respect to a specific weight or bias element.

Due to the complexity and hence non-existent analytical form for this solution, this is typically found via backpropagation, an algorithm which exploits the chain rule of the partial derivative and iteratively finds a local optimum using a first order gradient based algorithm, also known as "gradient descent." The gradient descent algorithm minimizes some function (such as the loss function in the context of machine learning) by iteratively moving in the direction of steepest descent, defined as the negative gradient. Formally, for a loss function $L(x)$ that is defined and has a gradient defined in the neighbourhood of the parameter set a , the updating algorithm is:

$$a_{n+1} = a_n - \gamma \Delta F(a_n) \quad (13)$$

where γ controls the size of each update. This γ parameter is known as the learning rate in neural network training, and controlling this is critical for good performance. As the loss functions of neural networks can be very complex with many local minima, the learning rate should be high enough such that the parameter updates are large enough to skip or jump over them. Too large of a learning rate however, and the neural may fail to converge to a solution at all. Due to computational limitations, we tune the learning rate manually, and consider a variety of different "optimizers", or algorithms which adapt the learning rate in different ways (see Appendix for computational details).

For our application, we considered the following grid of hyperparameters:

⁸Note that the specification of a constant "1" at the beginning of each layer is the same as specifying a bias term as is popular in other parametrizations.

Table 3: Hyperparameters considered for feed forward neural networks

Hyperparameter	Grid
Activation Function	ReLU, Leaky ReLU, tanh
Optimizer	ADAM, NADAM, SGD, RMSPROP
Learning Rate	(0.1, 0.01, 0.001, 0.0001)
L_1 Penalty	(1, 0.1, 0.01, 0.001)
L_2 Penalty	(1, 0.1, 0.01, 0.001)
Batch Size	(32, 64, 128, 256, 512, 1024, 2048)
Early Stopping Patience	(10, 20, 30, 40, 50)

359 A.5 Long short term memory networks

360 Long short term memory (LSTM) networks are special case of Recurrent Neural Networks (RNNs)
 361 which utilize LSTM “cells” to address some of their shortcomings, initially proposed by [9] (see for
 362 more comprehensive details). The general model setup for an RNN introduces intermediate “latent”
 363 states h_t which are simply the output for each time t and adds them as extra features for next period’s
 364 prediction:

$$h_t = \sigma_h(U_h x_t + V_h h_{t-1} + b_h) \quad (14)$$

365 where U_h , V_h and b_h are two parameter matrices and bias vectors to be learned, respectively. The
 366 resulting output from each time step is therefore:

$$y_t = \sigma_y(W_y h_t + b_y) \quad (15)$$

367 where W_y and b_y are similarly parameter matrices and bias vectors to be learned, respectively.

368 An LSTM “cell” is a specific structure of calculating and training latent states and their parameters.
 369 Let h_t , C_t be hidden layer vectors, x_t the original input vector at each time t , b_f , b_i , b_c and b_o be bias
 370 vectors, W_f , W_i , W_c and W_o parameter matrices. Then the output for each latent h_t is defined as:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) \quad (16)$$

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \quad (17)$$

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) \quad (18)$$

$$\tilde{C}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c) \quad (19)$$

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \quad (20)$$

$$h_t = o_t \odot \tanh(C_t) \quad (21)$$

371 where, for two vectors x , y of the same dimension, $x \odot y$ denotes their Hadamard product.

For our application, we considered the following grid of hyperparameters:

Table 4: Hyperparameters considered for LSTMs

Hyperparameter	Grid
Activation Function	ReLU, Leaky ReLU, tanh
Optimizer	ADAM, NADAM, SGD, RMSPROP
Learning Rate	(0.1, 0.01, 0.001, 0.0001)
L_1 Penalty	(1, 0.1, 0.01, 0.001)
L_2 Penalty	(1, 0.1, 0.01, 0.001)
Batch Size	(32, 64, 128, 256, 512, 1024, 2048)
Early Stopping Patience	(10, 20, 30, 40, 50)
Layer 1 Units	(1, 2, 4, 8)
Layer 2 Units	(1, 2, 4, 8)

A.6 FFORMA

Feature-based Forecast Model Averaging, ([14]) is an automated method for obtaining weighted forecast combinations for time series. We provide a brief overview of the two phases in this methodology.

We follow [14] selection of time series features as inputs to the meta-learner.

To incorporate all regressors in each individual time series model, we applied dimensional reduction techniques of PCA and UMAP to generate new feature mappings for use in GARCH (1, 1) models (generally the best performing of the constituent models). It was noted that none of the new external regressors as generated by these feature mappings improved fit, however.

The constituent models we considered are:

- Naive
- Random walk with drift
- Theta method
- ARIMA
- ETS
- TBATS
- Neural network auto-regressive model
- ARMA (1, 1) with Generalized Error Distribution GARCH(1, 1) errors
- ARMA (1, 1) with Generalized Error Distribution GARCH(1, 1) errors and UMAP external regressors

The time series features used to train the meta-model are detailed in [14], with the addition of realized volatility.

Note that because financial returns data does not typically exhibit seasonality, features and constituent models related which utilized seasonality were omitted.

A.7 DeepAR

DeepAR is a generalization of traditional Auto Regressive (AR) models to include additional layers into order to introduce non-linearities into the model. We provide the general formulation here - for full details see [17].

DeepAR aims to model the conditional distribution

$$P(\mathbf{r}_{i,t_0:T} | \mathbf{r}_{i,1:t_0-1}, \mathbf{x}_{i,1:T})$$

of the future of each time series $[r_{i,t_0}, r_{i,t_0+1}, \dots, r_{i,T}] := \mathbf{r}_{i,t_0:T}$ given its past $[r_{i,1}, \dots, r_{i,t_0-2}, r_{i,t_0-1}] := \mathbf{r}_{i,1:t_0-1}$, where t_0 denotes the time point from which we assume $r_{i,t}$ to be unknown at prediction time, and $\mathbf{x}_{i,1:T}$ are covariates that are assumed to be known for all time points. During training, both ranges have to lie in the past so that the $r_{i,t}$ are observed, but during prediction $r_{i,t}$ is only available in the conditioning range. DeepAR assumes that the model distribution $Q_{\Theta}(\mathbf{r}_{i,t_0:T} | \mathbf{r}_{i,1:t_0-1}, \mathbf{x}_{i,1:T})$ consists of a product of likelihood factors

$$Q_{\Theta}(\mathbf{r}_{i,t_0:T} | \mathbf{r}_{i,1:t_0-1}, \mathbf{x}_{i,1:T}) = \prod_{t=t_0}^T Q_{\Theta}(r_{i,t} | \mathbf{r}_{i,1:t-1}, \mathbf{x}_{i,1:T}) = \prod_{t=t_0}^T \ell(r_{i,t} | \theta(\mathbf{h}_{i,t}, \Theta))$$

parametrized by the output $\mathbf{h}_{i,t}$ of an autoregressive RNN

$$\mathbf{h}_{i,t} = h(\mathbf{h}_{i,t-1}, r_{i,t-1}, \mathbf{x}_{i,t}, \Theta), \quad (22)$$

where h is a function implemented by a multi-layer recurrent neural network with LSTM cells. The model is autoregressive, in the sense that it consumes the observation at the last time step $r_{i,t-1}$ as an input, as well as recurrent, i.e. the previous output of the network $\mathbf{h}_{i,t-1}$ is fed back as an input at the next time step. The likelihood $\ell(r_{i,t} | \theta(\mathbf{h}_{i,t}))$ is a fixed distribution whose parameters are given by a function $\theta(\mathbf{h}_{i,t}, \Theta)$ of the network output $\mathbf{h}_{i,t}$ (see below).

414 Information about the observations in the conditioning range $\mathbf{r}_{i,1:t_0-1}$ is transferred to the prediction
 415 range through the initial state \mathbf{h}_{i,t_0-1} . In the sequence-to-sequence setup, this initial state is the
 416 output of an *encoder network*. We follow [17] and opt for using the same architecture for the model
 417 in the conditioning range and the prediction range (corresponding to the *encoder* and *decoder* in
 418 a sequence-to-sequence model), and share weights between them, so that the initial state for the
 419 decoder \mathbf{h}_{i,t_0-1} is obtained by computing (22) for $t = 1, \dots, t_0 - 1$, where all required quantities
 420 are observed. The initial state of the encoder $\mathbf{h}_{i,0}$ as well as $r_{i,0}$ are initialized to zero.

421 Given the model parameters Θ , we can directly obtain joint samples $\tilde{\mathbf{z}}_{i,t_0:T} \sim$
 422 $Q_{\Theta}(\mathbf{r}_{i,t_0:T} | \mathbf{r}_{i,1:t_0-1}, \mathbf{x}_{i,1:T})$ through ancestral sampling: First, we obtain \mathbf{h}_{i,t_0-1} by computing
 423 (22) for $t = 1, \dots, t_0$. For $t = t_0, t_0 + 1, \dots, T$ we sample $\tilde{r}_{i,t} \sim \ell(\cdot | \theta(\tilde{\mathbf{h}}_{i,t}, \Theta))$ where
 424 $\tilde{\mathbf{h}}_{i,t} = h(\mathbf{h}_{i,t-1}, \tilde{r}_{i,t-1}, \mathbf{x}_{i,t}, \Theta)$ initialized with $\tilde{\mathbf{h}}_{i,t_0-1} = \mathbf{h}_{i,t_0-1}$ and $\tilde{r}_{i,t_0-1} = r_{i,t_0-1}$.

425 The parameters of the model denoted by Θ can be learned given $\mathbf{r}_{i,1:T}$ and associated covariates
 426 $\mathbf{x}_{i,1:T}$ by maximizing the log-likelihood:

$$\mathcal{L} = \sum_{i=1}^N \sum_{t=t_0}^T \log(\ell r_{i,t} | \theta(\mathbf{h}_{i,t})) \quad (23)$$

For our application, we utilized AWS' Bayesian grid search within the following ranges:

Table 5: Hyperparameters considered for DeepAR

Hyperparameter	From	To
Dropout Rate	0	0.2
Learning rate	0.0001	1
Likelihood	Gaussian	student-T
Mini batch size	16	128
Number of Cells	10	100
Number of Layers	1	4

427

B Additional details: simulation design

In this section, we give additional features of the simulation design required to implement our results. All code and data needed to replicate the results of this analysis are available at:

<https://github.com/Meron35/Evaluation-of-Machine-Learning-in-Asset-Pricing>.

B.1 Simulation Design

Our specification is a latent factor model for excess returns r_{t+1} , for $t = 1, \dots, T$:

$$\begin{aligned} r_{i,t+1} &= g(z_{i,t}) + \beta_{i,t+1}v_{t+1} + e_{i,t+1}; \quad z_{i,t} = (1, x_t)' \otimes c_{i,t}, \quad \beta_{i,t} = (c_{i1,t}, c_{i2,t}, c_{i3,t}) \\ e_{i,t+1} &= \sigma_{i,t+1}\varepsilon_{i,t+1}; \\ \log(\sigma_{i,t+1}^2) &= \omega + \gamma \log(\sigma_t^2) + \sigma_u u; \quad u \sim N(0, 1) \end{aligned}$$

where v_{t+1} is a 3×1 vector of errors, $w_{t+1} \sim N(0, 1)$, $\varepsilon_{i,t+1} \sim N(0, 1)$ scalar error terms, matrix C_t is an $N \times P_c$ matrix of latent factors, where the first three columns correspond to $\beta_{i,t}$, across the $1 \leq i \leq N$ dimensions, while the remaining $P_c - 3$ factors do not enter the return equation. The $P_x \times 1$ vector x_t is a 3×1 multivariate time series, and ε_{t+1} is a $N \times 1$ vector of idiosyncratic errors. The parameters of these were tuned such that the annualized volatility of each return series was approximately 22% when viewed as monthly returns, as is often observed empirically.

Simulating characteristics We build in correlation across time among factors by drawing normal random numbers for each $1 \leq i \leq N$ and $1 \leq j \leq P_c$, according to :

$$\bar{c}_{ij,t} = \rho_j \bar{c}_{ij,t-1} + \epsilon_{ij,t}; \quad \rho_j \sim \mathcal{U}(0.5, 1) \quad (24)$$

We then build in cross sectional correlation:

$$\hat{C}_t = L\bar{C}_t; \quad B = LL' \quad (25)$$

$$B := \Lambda\Lambda' + 0.1\mathbb{I}_n, \quad \Lambda_i = (\lambda_{i1}, \dots, \lambda_{i4}), \quad \lambda_{ik} \sim N(0, \lambda_{sd}), \quad k = 1, \dots, 4 \quad (26)$$

where B serves as a variance covariance matrix with λ_{sd} its density, and L represents the lower triangle matrix of B via the Cholesky decomposition. λ_{sd} values of 0.01, 0.1 and 1 were used to explore increasing degrees of cross sectional correlation. Characteristics are then normalized to be within $[-1, 1]$ for each $1 \leq i \leq N$ and for $j = 1, \dots, P_c$ via:

$$c_{ij,t} = \frac{2}{n+1} \text{rank}(\hat{c}_{ij,t}) - 1. \quad (27)$$

Simulating macroeconomic series We consider a Vector Autoregression (VAR) model for x_t , a 3×1 multivariate time series ⁹:

$$x_t = Ax_{t-1} + u_t; \quad A = 0.95I_3; \quad u_t \sim N(\mu = (0, 0, 0)', \Sigma = I_3)$$

Simulating return series We consider three different functions for $g(z_{i,t})$:

$$(1) \quad g_1(z_{i,t}) = (c_{i1,t}, c_{i2,t}, c_{i3,t} \times x'_t[3,]) \theta_0 \quad (28)$$

$$(2) \quad g_2(z_{i,t}) = (c_{i1,t}^2, c_{i1,t} \times c_{i2,t}, \text{sgn}(c_{i3,t} \times x'_t[3,])) \theta_0 \quad (29)$$

$$(3) \quad g_3(z_{i,t}) = (1[c_{i3,t} > 0], c_{i2,t}^3, c_{i1,t} \times c_{i2,t} \times 1[c_{i3,t} > 0], \text{logit}(c_{i3,t})) \theta_0 \quad (30)$$

where $x'_t[3,]$ denotes the third element of the x'_t vector. $g_1(z_{i,t})$ allows the characteristics to enter the return equation linearly, and $g_2(z_{i,t})$ and $g_3(z_{i,t})$ allow the characteristics to enter the return equation interactively and non-linearly. ¹⁰ θ^0 was tuned such that the predictive R^2 was approximately 5%.

The simulation design results in $3 \times 3 = 9$ different simulated datasets, each with $N = 200$ stocks, $T = 180$ periods and $P_c = 100$ characteristics. Each design was simulated 10 times to assess the robustness of machine learning algorithms, with the number of simulations kept low for computational

⁹More complex specifications for A were briefly explored, but these did not have a significant impact on results.

¹⁰(g_1, g_2 correspond to the simulation design used by [7].)

456 feasibility. We employ the hybrid data splitting approach with a training:validation length ratio of
457 approximately 1.5 and a test set that is 1 year in length. Other schemes in the forecasting literature such
458 as using an “inner” rolling window validation loop to find the best hyperparameters on average, finally
459 aggregating them in an “outer” loop for a more robust error were considered but not implemented
460 due to a) computational feasibility and b) the relative instability of optimal hyperparameters across
461 different different windows.

Table 6: Simulation Study Loss Statistics

model	Corr	g1			g2			g3		
		Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
LM.MSE	0.01	0.0366775	0.0027400	0.0082732	0.0382548	0.0028801	-0.1117880	0.0373098	0.0027954	-0.0320680
	0.10	0.0369652	0.0027653	-0.0110198	0.0385796	0.0029144	-0.1429443	0.0375694	0.0028168	-0.0549404
	1.00	0.0429486	0.0034141	-0.4387965	0.0453765	0.0037172	-0.7809535	0.0434339	0.0034688	-0.4887785
LM.MAE	0.01	0.0366417	0.0027373	0.0090496	0.0383478	0.0028862	-0.1163694	0.0373235	0.0027967	-0.0351619
	0.10	0.0368113	0.0027555	0.0029188	0.0387449	0.0029275	-0.1525797	0.0374894	0.0028098	-0.0476746
	1.00	0.0423399	0.0033445	-0.3930442	0.0453420	0.0036847	-0.7699555	0.0435349	0.0034682	-0.5445237
ELN.MSE	0.01	0.0345878	0.0025663	0.1403351	0.0362229	0.0026898	0.0368766	0.0353534	0.0026227	0.0991416
	0.10	0.0345630	0.0025643	0.1442376	0.0361830	0.0026860	0.0372585	0.0352923	0.0026167	0.1002410
	1.00	0.0346142	0.0025676	0.1671841	0.0362761	0.0026980	0.0378391	0.0354437	0.0026300	0.1198755
ELN.MAE	0.01	0.0345786	0.0025652	0.1409821	0.0361950	0.0026882	0.0391694	0.0353345	0.0026210	0.1004424
	0.10	0.0345582	0.0025637	0.1446272	0.0361730	0.0026877	0.0388747	0.0352851	0.0026167	0.1009186
	1.00	0.0345989	0.0025667	0.1677712	0.0363047	0.0027028	0.0365834	0.0354652	0.0026310	0.1180225
RF.MSE	0.01	0.0357752	0.0026710	0.0634257	0.0357179	0.0026571	0.0676147	0.0358032	0.0026613	0.0702977
	0.10	0.0357695	0.0026649	0.0667382	0.0356845	0.0026525	0.0691389	0.0358666	0.0026704	0.0628386
	1.00	0.0362325	0.0026977	0.0687741	0.0359893	0.0026833	0.0571035	0.0362129	0.0026952	0.0698868
RF.MAE	0.01	0.0354594	0.0026434	0.0833385	0.0354204	0.0026305	0.0876529	0.0355399	0.0026446	0.0865291
	0.10	0.0355153	0.0026489	0.0814253	0.0354894	0.0026345	0.0834048	0.0355688	0.0026438	0.0816426
	1.00	0.0359158	0.0026747	0.0870806	0.0356434	0.0026445	0.0809651	0.0360529	0.0026786	0.0753573
NN1.MSE	0.01	0.0364516	0.0027219	0.0163443	0.0367677	0.0027319	-0.0039174	0.0366874	0.0027384	0.0093355
	0.10	0.0364624	0.0027191	0.0204223	0.0367762	0.0027345	-0.0072588	0.0367326	0.0027372	0.0029550
	1.00	0.0375452	0.0028206	-0.0144520	0.0370492	0.0027638	-0.0146973	0.0374589	0.0027975	-0.0124689
NN1.MAE	0.01	0.0359604	0.0026786	0.0558139	0.0369206	0.0027474	-0.0151053	0.0363047	0.0026996	0.0393707
	0.10	0.0360823	0.0026866	0.0506976	0.0370100	0.0027503	-0.0205616	0.0363220	0.0027022	0.0323034
	1.00	0.0378894	0.0028338	-0.0431818	0.0379790	0.0028445	-0.0840747	0.0373056	0.0027926	0.0021783
NN2.MSE	0.01	0.0370187	0.0027850	-0.0217869	0.0373197	0.0027752	-0.0433537	0.0370890	0.0027745	-0.0173037
	0.10	0.0369775	0.0027651	-0.0212763	0.0370088	0.0027478	-0.0275384	0.0369898	0.0027584	-0.0206446
	1.00	0.0375360	0.0028138	-0.0139783	0.0369035	0.0027518	-0.0058664	0.0375157	0.0028087	-0.0169336
NN2.MAE	0.01	0.0358939	0.0026718	0.0577427	0.0368335	0.0027396	-0.0071579	0.0363352	0.0027028	0.0363052
	0.10	0.0358898	0.0026681	0.0603096	0.0369367	0.0027503	-0.0170774	0.0362701	0.0026960	0.0371567
	1.00	0.0374795	0.0028142	-0.0095290	0.0377146	0.0028226	-0.0653904	0.0374711	0.0028038	-0.0101183
NN3.MSE	0.01	0.0367827	0.0027568	-0.0067616	0.0368397	0.0027379	-0.0075249	0.0370360	0.0027644	-0.0200783
	0.10	0.0369384	0.0027613	-0.0153994	0.0368517	0.0027384	-0.0151060	0.0368743	0.0027573	-0.0044063
	1.00	0.0374242	0.0028081	-0.0129638	0.0369376	0.0027543	-0.0063529	0.0374202	0.0027991	-0.0103479
NN3.MAE	0.01	0.0358164	0.0026697	0.0654321	0.0369345	0.0027491	-0.0163983	0.0364712	0.0027181	0.0299484
	0.10	0.0358935	0.0026771	0.0620017	0.0368590	0.0027406	-0.0118497	0.0362000	0.0026932	0.0406114
	1.00	0.0370087	0.0027744	0.0213288	0.0372705	0.0027832	-0.0296437	0.0374132	0.0027916	-0.0083067
NN4.MSE	0.01	0.0368808	0.0027586	-0.0206197	0.0368555	0.0027423	-0.0077152	0.0371255	0.0027752	-0.0265634
	0.10	0.0368772	0.0027610	-0.0145791	0.0372207	0.0027615	-0.0487112	0.0368718	0.0027480	-0.0088940
	1.00	0.0373820	0.0028051	-0.0064811	0.0368966	0.0027505	-0.0053689	0.0373542	0.0027970	-0.0077389
NN4.MAE	0.01	0.0359348	0.0026782	0.0577196	0.0368974	0.0027487	-0.0109166	0.0367079	0.0027376	0.0070464
	0.10	0.0358281	0.0026651	0.0650415	0.0369333	0.0027494	-0.0191117	0.0362730	0.0026954	0.0377039
	1.00	0.0370948	0.0027786	0.0198663	0.0373230	0.0027947	-0.0293767	0.0373013	0.0027871	-0.0018876
NN5.MSE	0.01	0.0372306	0.0027846	-0.0499701	0.0369309	0.0027474	-0.0170017	0.0371140	0.0027720	-0.0218954
	0.10	0.0370264	0.0027669	-0.0321897	0.0371758	0.0027623	-0.0394362	0.0369093	0.0027565	-0.0113522
	1.00	0.0373642	0.0027949	-0.0104952	0.0369277	0.0027552	-0.0053762	0.0374751	0.0028071	-0.0149737

Table 6: Simulation Study Loss Statistics

model	Corr	g ¹			g ²			g ³		
		Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
NN5.MAE	0.01	0.0358880	0.0026693	0.0585792	0.0368354	0.0027380	-0.0086455	0.0366851	0.0027371	0.0046430
	0.10	0.0360381	0.0026803	0.0509764	0.0367451	0.0027273	-0.0049349	0.0364843	0.0027103	0.0181920
	1.00	0.0372849	0.0027940	0.0025412	0.0370382	0.0027652	-0.0127290	0.0371925	0.0027753	0.0025723
LSTM.MSE	0.01	0.0372963	0.0027982	-0.0432886	0.0372268	0.0027764	-0.0447640	0.0375909	0.0028180	-0.0625164
	0.10	0.0372369	0.0027946	-0.0319550	0.0371342	0.0027674	-0.0382547	0.0371984	0.0027845	-0.0303936
	1.00	0.0381282	0.0028506	-0.0820266	0.0373821	0.0027921	-0.0442426	0.0377803	0.0028300	-0.0443304
LSTM.MAE	0.01	0.0374310	0.0028046	-0.0564056	0.0373372	0.0027801	-0.0518537	0.0376270	0.0028169	-0.0674327
	0.10	0.0374461	0.0028036	-0.0629523	0.0371178	0.0027679	-0.0325442	0.0372409	0.0027931	-0.0333196
	1.00	0.0380266	0.0028456	-0.0614833	0.0374152	0.0027902	-0.0455057	0.0377435	0.0028252	-0.0458837
FFORMA.MSE	0.01	0.0382767	0.0028820	-0.1326717	0.0384600	0.0028893	-0.1473902	0.0424656	0.0033108	-0.4861451
	0.10	0.0383581	0.0028947	-0.1407652	0.0384795	0.0028912	-0.1600616	0.0423231	0.0032914	-0.4739906
	1.00	0.0388747	0.0029647	-0.1312392	0.0388080	0.0029331	-0.1659900	0.0430130	0.0033713	-0.4709541
FFORMA.MAE	0.01	0.0387548	0.0029387	-0.1797483	0.0387472	0.0029178	-0.1740938	0.0429893	0.0033651	-0.5279094
	0.10	0.0389359	0.0029511	-0.1927930	0.0387959	0.0029457	-0.1759939	0.0430966	0.0034057	-0.5863752
	1.00	0.0392468	0.0029721	-0.1636559	0.0393873	0.0029960	-0.2116186	0.0437090	0.0034483	-0.5260813
DeepAR	0.01	0.0382993	0.0029000	-0.1289295	0.0384895	0.0029121	-0.1325183	0.0393898	0.0030161	-0.2049803
	0.10	0.0388318	0.0029353	-0.1816633	0.0384345	0.0029045	-0.1318744	0.0391770	0.0029932	-0.1905583
	1.00	0.0405348	0.0031590	-0.2391417	0.0387870	0.0029524	-0.1440285	0.0396918	0.0030422	-0.1823646

463 B.3 Random Forest VIMPs

464 We note that random forest methods typically have their own methodologies to calculate variable
 465 importance (VIMP) which are different to the VIMP metric presented in the main body of the paper.
 466 We provide two popular schemes of calculating random forest VIMP - Breiman-cutler VIMP, [2] and
 467 Ishwaran-Kogalur VIMP, [11], and show that importantly, the overall conclusion regarding factor
 468 selection does not change with respect to which methodology employed.

Figure 3: Simulation Breiman-Cutler VIMP

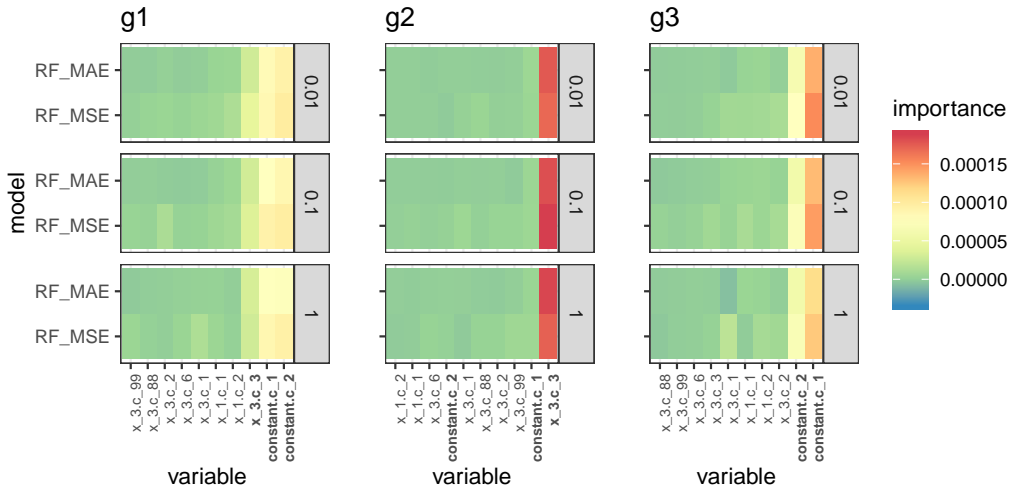



Figure 4: Simulation Ishwaran-Kogalur VIMP



C Additional details: Empirical analysis

C.1 Data & cleaning

We begin by obtaining monthly individual price data from CRSP for all firms listed in the NYSE, AMEX and NASDAQ, starting from 1957 (starting date of the S&P 500) and ending in December 2016, totalling 60 years. To build individual factors, we construct a factor set based on the cross section of returns literature. This is the same data used in [7]. We restrict our dataset to begin from 1993 Q3 and end on 2016 Q4 to alleviate data quality issues. Our individual factor set contains 94 characteristics: 61 updated annually, 13 updated quarterly and 20 updated monthly.¹¹ We detail our cleaning procedure of this dataset. To reduce the size of the dataset and increase feasibility, for the NASDAQ index we only consider equities with a traded share price larger than \$5. To achieve a balance between having a dataset with enough data points and variability among factors, the dataset was converted to a quarterly format. Quarterly returns were then constructed using the PRC variable according to:

$$RET_t = (PRC_t - PRC_{t-1})/PRC_{t-1} \quad (31)$$

We allow all stocks which have a quarterly return to enter the dataset, even if they disappear from the dataset for certain periods. To deal with missing data, any characteristics that had over 20% of their data missing were omitted. Remaining missing data were then imputed using their cross-sectional medians for each year.

We then follow [7] and construct eight macroeconomic factors following the variable definitions in [18] (see Table 7). These factors were lagged by one period so as to be used to predict one period ahead quarterly returns. The 3-month Treasury Bill rate was also used from this source to proxy for the risk free rate in order to construct excess quarterly returns. The two sets of factors were then combined to form a baseline set of covariates, which we define throughout all methods and analysis as:

$$z_{i,t} = (1, x_t)' \otimes c_{i,t} \quad (32)$$

where $c_{i,t}$ is a P_c matrix of characteristics for each stock i , and $(1, x_t)'$ is a $P_x \times 1$ vector of macroeconomic predictors, \otimes represents the Kronecker product. $z_{i,t}$ is therefore a $P_x P_c$ vector of features for predicting individual stock returns and includes interactions between stock level characteristics and macroeconomic variables. The total number of covariates in this baseline set is $61 \times (8 + 1) = 549$. The final dataset contains 202, 066 individual observations.


¹¹The dataset also included 74 Standard Industrial Classification (SIC) codes, but these were omitted due to their inconsistency, and inadequateness at classifying companies, as noted by WRDS

Table 7: Macroeconomic Factors, (see [18] for a detailed description).

No.	Acronym	Macroeconomic Factor
1	macro_dp	Dividend Price Ratio
2	macro_ep	Earnings Price Ratio
3	macro_bm	Book to Market Ratio
4	macro_ntis	Net Equity Expansion
5	macro_tbl	Treasury Bill Rate
6	macro_tms	Term Spread
7	macro_dfy	Default Spread
8	macro_svar	Stock Variance

497 We mimic the sample splitting procedure used in the simulation study: the dataset was split such that
 498 the training and validation sets were split such that the training set was approximately 1.5 times the
 length of the validation set, in order to predict a test set that is one year in length (see Figure 5).

Figure 5: Empirical Data Sample Splitting Procedure.



499

Table 8: Individual Factors

Acronym	Firm Characteristic	Data Source	Frequency
ind_absacc ¹¹	Absolute Accruals	Compustat	Annual
ind_acc ¹¹	Working capital accruals	Compustat	Annual
ind_aeavol	Abnormal earnings announcement volume	Compustat	Quarterly
ind_age	# years since first Compustat coverage	Compustat	Annual
ind_agr	Asset growth	Compustat	Annual
ind_baspread	Bid-ask spread	Compustat	Monthly
ind_beta	Beta	Compustat	Monthly
ind_betasq	Beta squared	Compustat	Monthly
ind_bm	Book-to-market	Compustat	Annual
ind_bm_ia	Industry-adjusted book to market	Compustat	Quarterly
ind_cash	Cash holdings	Compustat	Annual
ind_cashdebt	Cashflow to debt	Compustat	Annual
ind_cashpr	Cash productivity	Compustat	Annual
ind_cfp ¹¹	Cashflow to price ratio	Compustat	Annual
ind_cfp_ia ¹¹	Industry-adjusted cashflow to price ratio	Compustat	Annual
ind_chatoia ¹¹	Industry-adjusted change in asset turnover	Compustat	Annual
ind_chcsho	Change in shares outstanding	Compustat	Annual
ind_chempia	Industry-adjusted change in employee	Compustat	Annual
ind_chinv ¹¹	Change in inventory	Compustat	Annual
ind_chmom	Change in 6-month momentum	Compustat	Monthly
ind_chpmia ¹¹	Industry-adjusted change in profit margin	Compustat	Annual
ind_chtx	Change in tax expense	Compustat	Quarterly
ind_cinvest	Corporate investment	Compustat	Quarterly
ind_convind	Convertible debt indicator	Compustat	Annual
ind_currat	Current ratio	Compustat	Annual
ind_depr	Depreciation / PP&E	Compustat	Annual
ind_divi	Dividend initiation	Compustat	Annual
ind_divo	Dividend omission	Compustat	Annual
ind_dolvol	Dollar trading volume	Compustat	Monthly
ind_dy	Dividend to price	Compustat	Annual
ind_ear	Earnings announcement return	Compustat	Quarterly
ind_egr	Growth in common shareholder eq	Compustat	Annual
ind_ep	Earnings to price	Compustat	Annual
ind_gma	Gross profitability	Compustat	Annual
ind_grCAPX ¹¹	Growth in capital expenditures	Compustat	Annual
ind_grltnoa ¹¹	Growth in long term net operating assets	Compustat	Annual
ind_herf	Industry sales concentration	Compustat	Annual
ind_hire	Employee growth rate	Compustat	Annual
ind_idiovola	Idiosyncratic return volatility	Compustat	Monthly

¹¹The factor was included in the original dataset provided by [7], but was not used due to missing data issues.

ind_ill	Illiquidity	Compustat	Monthly
ind_indmom	Industry momentum	Compustat	Monthly
ind_invest ¹¹	Capital expenditures and inventory	Compustat	Annual
ind_lev	Leverage	Compustat	Annual
ind_lgr	Growth in long-term debt	Compustat	Annual
ind_maxret	Maximum daily return	Compustat	Monthly
ind_mom12m	12-month momentum	Compustat	Monthly
ind_mom1	1-month momentum	Compustat	Monthly
ind_mom36m ¹¹	36-month momentum	Compustat	Monthly
ind_mom6m	6-month momentum	Compustat	Monthly
ind_ms	Financial statement score	Compustat	Quarterly
ind_mvel1	Size	Compustat	Monthly
ind_mve_ia	Industry-adjusted size	Compustat	Annual
ind_nincr	Number of earnings increases	Compustat	Quarterly
ind_operprof ¹¹	Operating profitability	Compustat	Annual
ind_orgcap ¹¹	Organizational capital	Compustat	Annual
ind_pchcapx_ia ¹¹	Industry adjusted % change in capital expenditures	Compustat	Annual
ind_pchcurrat ¹¹	% change in current ratio	Compustat	Annual
ind_pchdepr ¹¹	% change in depreciation	Compustat	Annual
ind_pchgm_pchsale ¹¹	% change in gross margin - % change in sales	Compustat	Annual
ind_pchquick ¹¹	% change in quick ratio	Compustat	Annual
ind_pchsale_pchinvt ¹¹	% change in sales - % change in inventory	Compustat	Annual
ind_pchsale_pchrect ¹¹	% change in sales - % change in A/R	Compustat	Annual
ind_pchsale_pchxsga ¹¹	% change in sales - % change in SG	Compustat	Annual
ind_pchsaleinv ¹¹	% change sales-to-inventory	Compustat	Annual
ind_pctacc ¹¹	Percent accruals	Compustat	Annual
ind_pricedelay	Price delay	Compustat	Monthly
ind_ps	Financial statements score	Compustat	Annual
ind_quick	Quick ratio	Compustat	Annual
ind_rd	R&D increase	Compustat	Annual
ind_rd_mve ¹¹	R&D to market capitalization	Compustat	Annual
ind_rd_sale ¹¹	R&D to sales	Compustat	Annual
ind_realestate ¹¹	Real estate holdings	Compustat	Annual
ind_retvol	Return volatility	Compustat	Monthly
ind_roaq	Return on assets	Compustat	Quarterly
ind_roavol ¹¹	Earnings volatility	Compustat	Quarterly
ind_roeq	Return on equity	Compustat	Quarterly
ind_roic	Return on invested capital	Compustat	Annual
ind_rsup	Revenue surprise	Compustat	Quarterly
ind_salecash	Sales to cash	Compustat	Annual
ind_saleinv ¹¹	Sales to inventory	Compustat	Annual
ind_salerec	Sales to receivables	Compustat	Annual
ind_secured ¹¹	Secured debt	Compustat	Annual
ind_securedind	Secured debt indicator	Compustat	Annual
ind_sgr ¹¹	Sales growth	Compustat	Annual
ind_sin	Sin stocks	Compustat	Annual
ind_sp	Sales to price	Compustat	Annual
ind_std_dolvol	Volatility of liquidity (dollar trading volume)	Compustat	Annual
ind_std_turn	Volatility of liquidity (share turnover)	Compustat	Monthly
ind_stdacc ¹¹	Accrual volatility	Compustat	Monthly
ind_stdcf ¹¹	Cashflow volatility	Compustat	Quarterly
ind_tang	Debt capacity/rm tangibility	Compustat	Quarterly
ind_tb ¹¹	Tax income to book income	Compustat	Annual
ind_turn	Share turnover	Compustat	Monthly
ind_zerotrade	Zero trading days	Compustat	Monthly

C.2 Empirical study comprehensive results & robustness checks

We now provide and discuss the comprehensive set of results for the empirical analysis in Section 4 of the main paper. These results have been omitted from the main body of the paper due to space restrictions. The predictive accuracy results across all models considered in the empirical analysis are provided in Table 9, and represent an expanded version of Table 2 in the main text. These expanded results echo the empirical conclusions discussed in Section 4.

In addition, we also provide four robustness checks to ascertain the reliability of our empirical results. We reconsider the empirical analysis using different training/validation splitting schemes, our missing data imputation scheme and the impact of additional features. Importantly, our overall results are consistent across all checks.

First we reconsider the analysis using a missing data threshold of 10% - that is, any regressors with over 10% missing data were omitted before being imputed. The results for predictive accuracy are displayed in Table 10, while the factor importance results are given in Figures 7-8. In terms of both prediction performance and factor importance, the results are very similar to those obtained with a 20% missing threshold. Indeed, most models display a lower value of the overall loss, indicating that removing many of the poorer quality factors may actually improve predictive performance.

We then consider training:validation length ratios of 1:1 and 1:2 in addition to 1:1.5 in the main study.

A training:validation length ratio of 1:1 saw the elastic net models perform the best, with the random forests fitted with respect to mean squared loss suffering noticeably in performance; see Table 11 for full details. The factor importance results were subdued across all models except for the elastic nets - random forests and neural networks struggled to discern any importance individual factors, and occasionally picked out the dividend price ratio and earnings price ratio macroeconomic factors as important; see Figures 9-10 for details.

A training:validation length ratio of 1:2 again produced similar results to those in the main study - quantile forests performed best in the first sample, and were then beaten by elastic nets. Interestingly, it was noted that quantile elastic nets started to suffer from convergence issues, and hence had very poor performance in the last sample, though their mean squared variant was still the best performing model; see Table 12 and Figures 11-12 for full details.

Finally, we consider supplementing our macroeconomic feature set with the five Fama-French factors (see [3]). The empirical analysis was then re-run using these additional features and a missingness threshold of 20%. Similarly to the initial analysis, elastic net models dominate in terms of prediction performance, though their quantile versions were again noted to suffer from convergence issues; see Table 13 for details. In addition, while introducing these additional factors has led to a sizeable increase in dimensionality of the feature space,¹² it has not yielded any significantly new insights into the factors driving asset returns; see Figures 13-14 for details.

¹²We recall that, since the features are constructed using the Kronecker product, adding the five Fama-French factors leads to $61 \times (13 + 1) = 854$ overall factors, which is a sizeable increase from the previous $61 \times (8 + 1) = 549$ factors considered in the empirical analysis in Section 4.

Table 9: Empirical Study Loss Statistics

model	Sample 1			Sample 2			Sample 3		
	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
LM.MSE	0.229034	0.116015	-1.808481	0.397573	0.312653	-6.329935	0.566307	0.83804	-17.522476
LM.MAE	0.273452	0.15894	-2.8476	0.555673	0.742223	-16.400898	0.651614	1.225121	-26.077774
ELN.MSE	0.133887	0.039947	0.032956	0.140402	0.04277	-0.002712	0.14433	0.043761	0.032789
ELN.MAE	0.131369	0.040718	0.014306	0.137092	0.041892	0.017875	0.146251	0.045207	0.000835
RF.MSE	0.130366	0.036629	0.113289	0.195817	0.070642	-0.656158	0.157934	0.05122	-0.132066
RF.MAE	0.126703	0.036785	0.109505	0.173721	0.057546	-0.349132	0.14692	0.046037	-0.01752
NN1.MSE	0.169127	0.057044	-0.380909	0.207662	0.074751	-0.752494	0.192125	0.069738	-0.541369
NN1.MAE	0.157324	0.050418	-0.22052	0.191762	0.066746	-0.564818	0.18547	0.063053	-0.393606
NN2.MSE	0.168773	0.059436	-0.43883	0.181808	0.063232	-0.482433	0.180584	0.062745	-0.386797
NN2.MAE	0.162667	0.055447	-0.342256	0.194277	0.069386	-0.626702	0.185173	0.065186	-0.440746
NN3.MSE	0.154784	0.050152	-0.21408	0.180103	0.060193	-0.411175	0.177604	0.060404	-0.335065
NN3.MAE	0.146411	0.044901	-0.086967	0.18499	0.06461	-0.514744	0.184986	0.063861	-0.411475
NN4.MSE	0.153802	0.048641	-0.177503	0.193066	0.067515	-0.582833	0.172707	0.057774	-0.276929
NN4.MAE	0.157301	0.050286	-0.217308	0.168815	0.055711	-0.306102	0.167998	0.055129	-0.218463
NN5.MSE	0.149436	0.047279	-0.14452	0.183584	0.064137	-0.503653	0.170238	0.056992	-0.259652
NN5.MAE	0.140781	0.042832	-0.036882	0.181096	0.06216	-0.4573	0.164896	0.053458	-0.181528

Figure 6: Empirical study random forest vimps




Table 10: Missing Data Threshold Robustness Check Loss Statistics

model	Sample 1			Sample 2			Sample 3		
	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
LM.MSE	0.247457	0.130166	-2.151058	0.541089	0.700574	-15.424468	0.615714	1.188991	-25.279238
LM.MAE	0.214055	0.102848	-1.489727	0.372683	0.259976	-5.094954	0.507397	0.766373	-15.93847
ELN.MSE	0.133887	0.039947	0.032956	0.140402	0.04277	-0.002712	0.14433	0.043761	0.032789
ELN.MAE	0.131338	0.040465	0.020421	0.137083	0.041804	0.019938	0.146589	0.045362	-0.002596
RF.MSE	0.129226	0.035869	0.131692	0.198914	0.072749	-0.705542	0.168068	0.05777	-0.276838
RF.MAE	0.124319	0.035103	0.150229	0.167845	0.053578	-0.256106	0.15463	0.051594	-0.140342
NN1.MSE	0.153785	0.048726	-0.179553	0.221019	0.084867	-0.98964	0.172557	0.058354	-0.289742
NN1.MAE	0.154534	0.048854	-0.18266	0.199647	0.073699	-0.727823	0.176348	0.061359	-0.356155
NN2.MSE	0.158513	0.057061	-0.381324	0.233631	0.095004	-1.227299	0.154083	0.048353	-0.068708
NN2.MAE	0.138489	0.043364	-0.049759	0.215253	0.078792	-0.847234	0.164459	0.055049	-0.216706
NN3.MSE	0.167392	0.058508	-0.416345	0.19754	0.071293	-0.671422	0.156873	0.049602	-0.096299
NN3.MAE	0.144457	0.045293	-0.096445	0.210372	0.077747	-0.822723	0.159841	0.05152	-0.138704
NN4.MSE	0.147989	0.047211	-0.142888	0.184277	0.064247	-0.506225	0.152214	0.048185	-0.064987
NN4.MAE	0.15851	0.052021	-0.259326	0.18643	0.063032	-0.477746	0.177651	0.064046	-0.415562
NN5.MSE	0.153187	0.050053	-0.211683	0.181622	0.060313	-0.413989	0.161028	0.051221	-0.132095
NN5.MAE	0.149496	0.050779	-0.229251	0.165726	0.053988	-0.265712	0.156151	0.049772	-0.100061

Figure 7: Missing Data Threshold Robustness Check Individual Factor Importance




Figure 8: Missing Data Threshold Robustness Check RF VIMP




Table 11: Train:Validation 1:1 Robustness Check Loss Statistics

model	Sample 1			Sample 2			Sample 3		
	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
LM.MSE	0.915703	2.495094	-59.401029	0.717	1.553454	-35.419641	0.451206	0.375505	-7.299459
LM.MAE	0.751551	1.583265	-37.32754	0.469831	0.524686	-11.300895	0.675112	1.105759	-23.43964
ELN.MSE	0.134609	0.040072	0.029933	0.141434	0.043169	-0.012055	0.144375	0.043705	0.034019
ELN.MAE	0.131668	0.040748	0.013583	0.137494	0.042135	0.012178	0.146776	0.045753	-0.01123
RF.MSE	0.155282	0.046655	-0.129427	0.210936	0.078006	-0.828784	0.229147	0.092622	-1.047155
RF.MAE	0.13882	0.04016	0.027805	0.185338	0.063217	-0.482087	0.182753	0.063873	-0.411736
NN1.MSE	0.218129	0.087699	-1.123002	0.238606	0.110201	-1.583582	0.260721	0.120908	-1.672321
NN1.MAE	0.202259	0.072844	-0.763409	0.205092	0.073567	-0.724721	0.239051	0.096477	-1.132346
NN2.MSE	0.239446	0.101312	-1.452556	0.206109	0.078412	-0.838305	0.228591	0.095126	-1.102488
NN2.MAE	0.19141	0.068261	-0.652455	0.184095	0.062366	-0.462125	0.220087	0.086888	-0.920403
NN3.MSE	0.193117	0.069206	-0.675336	0.193859	0.070747	-0.658609	0.205093	0.076497	-0.690745
NN3.MAE	0.191596	0.066926	-0.620138	0.176555	0.060022	-0.407183	0.234768	0.091003	-1.011359
NN4.MSE	0.191361	0.07068	-0.71101	0.175311	0.059253	-0.389136	0.18148	0.061718	-0.364096
NN4.MAE	0.139659	0.041096	0.005158	0.179318	0.05976	-0.401027	0.188921	0.066144	-0.461932
NN5.MSE	0.17209	0.056982	-0.379418	0.164756	0.054398	-0.275325	0.202012	0.074051	-0.636691
NN5.MAE	0.170945	0.056029	-0.356356	0.180669	0.059697	-0.399552	0.189149	0.065921	-0.456988

Figure 9: Train:Validation = 1:1 Robustness Check Individual Factor Importance




Figure 10: Train:Validation = 1:1 Robustness Check RF VIMP




Table 12: Train:Validation 2:1 Robustness Check Loss Statistics

model	Sample 1			Sample 2			Sample 3		
	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
LM.MSE	0.277087	0.164599	-2.98459	0.383421	0.31299	-6.337839	0.523418	0.740288	-15.361936
LM.MAE	0.246936	0.147979	-2.582262	0.277044	0.161215	-2.779579	0.487285	0.631575	-12.95915
ELN.MSE	0.133715	0.039919	0.033647	0.139723	0.042525	0.003028	0.145034	0.044306	0.020752
ELN.MAE	0.131237	0.040361	0.022952	0.137205	0.041858	0.018674	0.174408	0.064513	-0.425873
RF.MSE	0.130808	0.036982	0.104754	0.162762	0.051118	-0.198417	0.155264	0.048661	-0.075516
RF.MAE	0.127013	0.036722	0.111033	0.146758	0.043961	-0.030633	0.168905	0.055983	-0.237348
NN1.MSE	0.155088	0.050284	-0.217281	0.165871	0.053459	-0.253309	0.181984	0.064621	-0.428262
NN1.MAE	0.159797	0.050566	-0.224107	0.163397	0.052329	-0.226828	0.181636	0.062407	-0.379326
NN2.MSE	0.155815	0.050954	-0.233492	0.168576	0.055738	-0.306745	0.170991	0.057453	-0.269824
NN2.MAE	0.148149	0.047617	-0.152709	0.166334	0.054058	-0.26734	0.163141	0.052639	-0.163436
NN3.MSE	0.154141	0.04976	-0.204586	0.166218	0.053402	-0.251967	0.169539	0.05661	-0.251204
NN3.MAE	0.142464	0.043771	-0.059594	0.154233	0.048682	-0.141321	0.184217	0.064175	-0.418401
NN4.MSE	0.166547	0.056184	-0.360092	0.150748	0.047566	-0.115162	0.168447	0.056575	-0.250437
NN4.MAE	0.150167	0.046919	-0.135802	0.16197	0.05226	-0.225199	0.171676	0.057352	-0.267598
NN5.MSE	0.155784	0.052258	-0.265047	0.139699	0.043082	-0.010018	0.166166	0.055027	-0.216219
NN5.MAE	0.161161	0.053216	-0.28825	0.149207	0.046344	-0.086511	0.149424	0.047544	-0.050824

Figure 11: Train:Validation = 1:2 Robustness Check Individual Factor Importance




Figure 12: Train:Validation = 1:2 Robustness Check RF VIMP




Table 13: Fama French Factor Robustness Check Loss Statistics

model	Sample 1			Sample 2			Sample 3		
	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2	Test MAE	Test MSE	Test R^2
LM.MSE	0.288636	0.182966	-3.42923	0.367636	0.264918	-5.210825	1.101604	5.012469	-109.78624
LM.MAE	0.280535	0.179777	-3.352038	0.376163	0.279476	-5.552114	1.25341	7.06036	-155.048996
ELN.MSE	0.13383	0.039956	0.032746	0.14022	0.0427	-0.00107	0.144472	0.043852	0.030769
ELN.MAE	0.128936	0.039665	0.039798	0.13716	0.042144	0.011965	0.172148	0.063154	-0.395841
RF.MSE	0.146318	0.042607	-0.031434	0.151137	0.047091	-0.104011	0.177125	0.064664	-0.429221
RF.MAE	0.138266	0.04005	0.030475	0.138714	0.042246	0.009583	0.152068	0.048488	-0.071698
NN1.MSE	0.168063	0.055354	-0.340017	0.192143	0.068904	-0.61541	0.275195	0.138165	-2.053731
NN1.MAE	0.161596	0.051507	-0.246873	0.199416	0.068181	-0.598444	0.23054	0.093434	-1.065082
NN2.MSE	0.169842	0.056899	-0.377415	0.179733	0.058966	-0.382416	0.252929	0.117102	-1.588199
NN2.MAE	0.155816	0.046809	-0.133147	0.185008	0.060854	-0.426679	0.219342	0.085115	-0.881213
NN3.MSE	0.1621	0.053165	-0.287008	0.182996	0.059643	-0.398278	0.232226	0.099353	-1.195903
NN3.MAE	0.161255	0.050737	-0.228237	0.191625	0.064676	-0.516291	0.218355	0.085297	-0.885238
NN4.MSE	0.166036	0.055575	-0.345349	0.191589	0.066207	-0.552182	0.23417	0.097348	-1.151607
NN4.MAE	0.148375	0.045227	-0.094843	0.168623	0.054176	-0.270114	0.20837	0.077667	-0.7166
NN5.MSE	0.147379	0.044503	-0.077315	0.166006	0.054935	-0.287914	0.20667	0.077866	-0.721013
NN5.MAE	0.150541	0.045723	-0.106868	0.172466	0.055402	-0.298865	0.218796	0.084938	-0.877301

Figure 13: Fama French Factors Robustness Check Individual Factor Importance

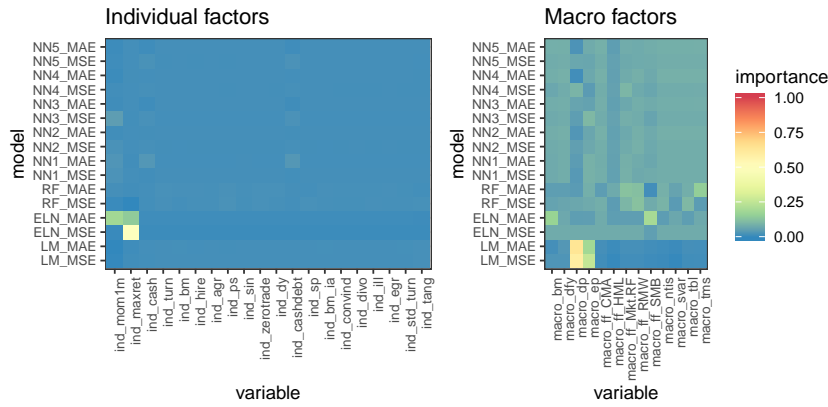


Figure 14: Fama French Factors Robustness Check RF VIMP

