
Evaluation of Machine Learning in Empirical Asset Pricing

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Several recent studies have claimed that machine learning methods provide superior
2 predictive accuracy of asset returns, relative to simpler modelling approaches, and
3 can correctly identify factors needed to price portfolio risk. Herein, we demonstrate
4 that this performance is critically dependent on several features of the data being
5 analysed; including, the training/test sample split, the frequency at which the data
6 is observed, and the chosen loss-function. In contrast to existing studies, which
7 claim that neural nets provide superior predictive accuracy, through a series of
8 realistic examples that mimics the stylized facts of asset returns, we demonstrate
9 that neural methods are easily outperformed by simpler methods, such as random
10 forests and elastic nets.

11 1 Introduction

12 The dominance of machine learning (hereafter, ML) methods in terms of predictive accuracy has now
13 begun to filter into the empirical asset pricing literature. The most common application of machine
14 learning methods within finance are for portfolio construction, asset price prediction, and factor
15 selection.

16 Several studies have now used ML techniques to analyze the cross-section of asset returns and
17 produce portfolios that can capture nonlinear information in the cross-section of asset returns. Mortiz
18 and Zimmermann (2016) use tree-based methods to understand which firm-level characteristics best
19 predict the cross-section of stock returns, and use this information to help mitigate portfolio risk.
20 Similarly, Messemmer (2017) uses deep feedforward neural nets (DFNs) to construct portfolios and
21 predict the returns across a cross-sections of US asset returns. However, while Messemmer (2017)
22 demonstrates that DFNs can better capture nonlinear information, no claim is made that deep learning
23 methods are the best approach to exploit this information.

24 In addition, several studies have now suggested that machine learning methods can produce better
25 predictions of asset returns ([?], [?] and [?]). In particular, the results of Gu et al. (2019) suggest that,
26 in terms of predictive performance, as measured by an out-of-sample R^2 , tree-based methods and
27 shallow neural nets can provide superior predictive accuracy over other machine learning methods
28 and simpler model-based approaches. This finding is born out both in terms of simulated data, and
29 an empirical example with monthly returns data from 1957 to 2016. [?] attribute this to machine
30 learning’s ability to evaluate and consider non-linear complexities among factors that cannot be
31 feasibly achieved using traditional techniques.

32 Similarly, work by Kozak et al, (2018), Freyberger et al. (2018), Feng et al., (2019) and Rapach
33 and Zhou (2013), demonstrate that machine learning methods can “systematically evaluate the
34 contribution to asset pricing of any new factor” used within an existing linear asset pricing structure.
35 In addition, Gu et al. (2019) use variable importance metrics to quantify the differential impact of

36 factors across a large set of possible factors available for asset pricing. As such, these authors argue
37 that ML can be used, *en masse*, to consistently evaluate the ability of various factors to help price
38 portfolio risk. Such work is particularly pertinent given the literature’s obsession with constructing
39 such factors: as of 2014, quantitative trading firms were using 81 factor models (Hsu and Kalesnik,
40 2014), while Harvey and Liu (2019) currently document that well over 600 different factors have
41 been suggested in the literature.

42 The above studies all demonstrate the potential benefits of ML methods within empirical finance.
43 However, it is unclear if the above findings generalize to: one, different training and validation periods;
44 two, different sampling frequencies, which results in returns with significant different characteristics;
45 three, different loss-measures of predictive accuracy. The answer to such questions in the realm of
46 empirical finance are particularly pertinent given that certain ML methods, including those references
47 above, have known difficulties in dealing with data that display the stylized facts of asset returns; e.g.,
48 penalized regression and tree-based models assume a form of conditional independence that is clearly
49 violated by the state dependence within, and across, asset returns. Moreover, training even standard
50 types of neural networks, such as DFNs, becomes particularly difficult when data displays strong, or
51 nonlinear, dependence ([?]).

52 In some ways, existing applications of ML to empirical finance have either over-looked, downplayed,
53 or simply ignored the importance of the above (and other) issues. Messemer (2017) and [?] use cross
54 validation as part of their model building procedures, destroying the temporal ordering of data. [?]
55 and Messemer (2017) produce models using training samples that end much earlier than the data sets
56 which they ultimately produce forecasts for: in the case of Messemer (2017), the training period ends
57 in 1981, and forecasts are produced for the most recent 30 years of data; in [?], the training set ends in
58 the 1970s, with predictions ultimately produced only for the period of returns from 1987-2016. This
59 is particularly worrying as the factors driving returns in the 1980s, are starkly different than those
60 driving returns in, say, 2001 onwards. However, both papers suggest that their training/validation
61 split does not impact the test set results.

62 The goal of this paper is to provide a systematic, and reproducible study on the ability of ML methods
63 to 1) accurately detect significant factors; and 2) accurately predict returns according to a range of
64 loss measures. It is our belief that any such study is necessary in order for practitioners to reliably
65 apply these methods in their problems of interest.

66 After stating the general setup in Section two, in Section three we conduct a rigorous simulation
67 study, which captures the stylized facts of asset returns, we give an in-depth comparison of several
68 machine learning methods used in the literature. The simulation study explores how persistence
69 in features, cross sectional correlation and different complexities of data generating process affect
70 ML method’s ability to: 1) accurately predict future returns across a range of loss measures; and 2)
71 correctly identify the significant factors driving returns. In contrast to existing findings, in this realistic
72 simulation design, we find that neural network procedures, such as feedforward nets, LSTM (CITE),
73 and DeepAR models (CITE), are among the worst performing methods, while simpler tree-based
74 methods and elastic net are among the best performing methods. This result is consistent across
75 various levels of volatility, cross-sectional correlation, return signal, and different loss functions.
76 Elastic net and tree-based methods also outperform other methods in correctly identifying significant
77 factors.

78 In Section four, the above findings are then validated via an empirical exercise that considers quarterly
79 individual returns data from CRSP for all firms listed in the NYSE, AMEX and NASDAQ. The
80 starting period of the data is January first 1957 (starting date of the S&P 500) and the ending date is
81 December 2016, totalling 60 years. A set of 549 possible factors are used to explain the cross-section
82 of returns. We pay careful attention to the training and test split, and only use the last fourteen years
83 of quarterly returns to evaluate the different machine learning methods. The results in this analysis
84 completely agree with those in the aforementioned simulation study: across all machine learning
85 methods, neural net based procedure perform the worst, while tree-based methods and elastic net
86 perform the best.

87 The results of this study suggest that great care and diligence is required if one wishes to implement
88 machine learning methods within empirical finance. Our results suggest that the efficacy of machine
89 learning methods within empirical finance depends on several features of the underlying problem,
90 such as sampling frequency, the particular test training split, and the data period under analysis. As

such, while potentially useful, ML methods are not a panacea for predicting, or understanding the factors that drive, financial returns.

2 Model and Methods

2.1 Statistical Model

In this section we briefly discuss the statistical model considered for asset returns. Excess monthly returns on asset i , $i = 1, \dots, n$, at time t , $t = 1, \dots, T$, are assumed to evolve in an additive fashion:

$$r_{i,t+1} = E(r_{i,t+1}|\mathcal{F}_t) + \epsilon_{i,t+1}, \quad E(\epsilon_{i,t+1}|\mathcal{F}_t) = 0 \quad (1)$$

where \mathcal{F}_t denotes the observable information at time t , and $\epsilon_{i,t+1}$ is a martingale difference sequence (hereafter, mds). We further consider that the conditional mean of returns is an unknown function of a P -dimensional vector of features, assumed measurable at time t , such that

$$E(r_{i,t+1}|\mathcal{F}_t) = g(z_{i,t}) \quad (2)$$

The features, or predictors, $z_{i,t}$ are assumed to be composed of time- t information, and depends only the characteristics of stock i . It is not assumed that all $z_{i,t}$ are present within the function $g(\cdot)$ across all i units. That is, the function $g(\cdot)$ need not depend on the same $z_{i,t}$ as i varies. The assumption that the information set can be characterized by the variables $z_{i,t}$ without dependence on the $j \neq i$ return units, is reasonable given that the collection of $z_{i,t}$ is rich enough.

In what follows, we represent the space of possible features as the Kronecker product of two pieces

$$z_{i,t} = x_t \otimes c_{i,t} \quad (3)$$

where the variables $c_{i,t}$ represent a $P_c \times 1$ vector of individual-level characteristics for return i , and x_t represents a $P_x \times 1$ vector of macroeconomic predictors, and \otimes represents the Kronecker product. Thus, for $P = P_c \cdot P_x$, $z_{i,t}$ represents a $P \times 1$ feature space that can be used to approximate the unknown function $g(\cdot)$.

2.2 Methods

Given features $z_{i,t}$, the goal of any machine learning method is to approximate the unknown function $g(\cdot)$ in 1. Broadly speaking, how different ML methods choose to approximate this function depends on three components:

1. the model used to make predictions,¹
2. the regularization mechanism employed to mitigate over-fitting;
3. a loss function that penalized poor predictions.

To ensure the results of ML different methods will be comparable, we fix both the regularization mechanisms and loss functions used within each method, and allow only the models used for prediction to vary. This approach seeks to ensure that performances in one method, relative to another, are based on the model structure and not to some feature of how the models were fit. To this end, we first discuss points 2. and 3. above, and then briefly present the models used for our comparison.

Loss functions: The choice of loss function used to fit the ML methods is instrumental in the methods' ultimate performance. Herein, we consider two separate loss functions: Mean Absolute Error (MAE) and Mean Squared Error (MSE):

$$\text{MAE} = \frac{1}{n} \sum_{j=i}^n |y_j - \hat{y}_j| \text{ and } \text{MSE} = \frac{1}{n} \sum_{j=i}^n (y_j - \hat{y}_j)^2,$$

We consider both loss functions since MAE is less sensitive to outliers in the data which financial returns are known to exhibit, and which are caused by extreme market movements. Given this, we expect MAE to produce predictive results that are more robust to such outlier events.

¹The model used by the ML method need not correspond to the statical models assumed to describe the data. Herein, our goal will not be to asses the "accuracy" of the statistical model, but to determine how different ML methods accurately determine the salient features of this model.

128 **Sample Splitting:** ML methods guard against over-fitting by emphasizing out-of-sample perfor-
 129 mance. To this end, observed data is split into “training”, “validation” and “test” sets. Since returns
 130 data is intrinsically dependent, when constructing such a split we must consider a schema that respects
 131 this dependence structure.

132 Throughout our experiments/applications, to balance computation and accuracy, we use a hybrid
 133 “rolling window” and “recursive” approach to training/validation/test splits: for each model refit, the
 134 training set is increased by one year observations, i.e., 12 monthly observations; the validation set is
 135 fixed at one year and moves forward (by one year) with each model refit; predictions are generated
 136 using that model for the subsequent year.

137 **Models** The remaining specification for the ML methods is the chosen model used to generate
 138 predictions. Herein, we consider a host of different models: including elastic net (Hastie et al.,
 139 XXX), Random forest (XXX), feed-forward neural nets (XXX), LSTM (XXX), FFORMA (XXX)
 140 and DeepAR models (XXX). To keep the details as brief as possible, we give full details on each
 141 model and certain features of its implementation used in this work in the appendix. For each of the
 142 different methods, we consider two variants, one based on the MAE loss and one based on the MSE
 143 loss.

144 2.3 Model evaluation measures

145 **Predictive accuracy** Predictive performance for individual excess returns are assessed using Mean
 146 Absolute Error (MAE), Mean Squared Error (MSE) (evaluated over the test set) and an out-of-sample
 147 R^2 measure. While out-of-sample R^2 is a common measure, there is no universally agreed-upon
 148 definition. As such, we explicitly state the version employed herein as

$$R_{OOS}^2 = 1 - \frac{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \hat{r}_{i,t+1})^2}{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \bar{r}_{i,t+1})^2} \quad (4)$$

149 where \mathcal{T}_3 indicates that the fits are only assessed on the test sub-sample, which is never used for
 150 training or tuning.

151 Since R^2 is based on in-sample-fit of a linear model, this measure is less meaningful for most of the
 152 ML methods considered in in this paper. However, we report this measure since this measure has also
 153 been considered in other applications of ML to empirical finance (see, e.g., Gu et al., 2019).

154 **Factor Selection** An important aspect of empirical finance is the understanding of which features
 155 drive risk. That is, which features are explicitly represented within $z_{i,t}$ and can thus be used to help
 156 price risk using equation 1. To this end, we define a simple variable importance (VI) measure to be
 157 applied across all ML methods in this research. To this end, we mirror the measure produced in [?]
 158 and define VI_j as the reduction in predictive R^2 from setting all values of predictor j to 0, while
 159 holding the remaining model estimates fixed. Each VI_j is then normalized to sum to 1.

160 However, as VI_j can sometimes be negative, we shift VI_j by the smallest VI_j plus a small constant,
 161 then dividing by this sum to alleviate numerical issues². The resulting VI measure is then.

$$VI_{j,norm} = \frac{VI_j + \min(VI_j) + o}{\sum VI_j + \min(VI_j) + o} \quad ; \quad o = 10^{-100} \quad (5)$$

162 3 Simulation study

163 We begin with the simulation study as a way to explore how machine learning performs with regards
 164 to the stylized facts of empirical returns in a controlled environment. We simulate according to a
 165 design which incorporates low signal to noise ratio, stochastic volatility in errors, persistence and
 166 cross sectional correlation in regressors. Our specification is a latent factor model for excess returns

²This mechanism was chosen because the other popular normalization mechanism “softmax” was observed to be unable to preserve the distances between each original VI_j , making discernment between each VI_j difficult.

167 r_{t+1} , for $t = 1, \dots, T$:

$$r_{i,t+1} = g(z_{i,t}) + \beta_{i,t+1}v_{t+1} + e_{i,t+1}; \quad z_{i,t} = (1, x_t)' \otimes c_{i,t}, \quad \beta_{i,t} = (c_{i1,t}, c_{i2,t}, c_{i3,t}) \quad (6)$$

$$e_{i,t+1} = \sigma_{i,t+1}\varepsilon_{i,t+1}; \quad (7)$$

$$\log(\sigma_{i,t+1}^2) = \omega + \gamma \log(\sigma_t^2) + \sigma_u u; \quad u \sim N(0, 1) \quad (8)$$

168 where v_{t+1} is a 3×1 vector of errors, $w_{t+1} \sim N(0, 1)$, $\varepsilon_{i,t+1} \sim N(0, 1)$ scalar error terms, matrix
 169 C_t is an $N \times P_c$ matrix of latent factors, where the first three columns correspond to $\beta_{i,t}$, across the
 170 $1 \leq i \leq N$ dimensions, while the remaining $P_c - 3$ factors do not enter the return equation. The
 171 $P_x \times 1$ vector x_t is a 3×1 multivariate time series, and ε_{t+1} is a $N \times 1$ vector of idiosyncratic
 172 errors. The parameters of these were tuned such that the annualized volatility of each return series
 173 was approximately 22%, as is often observed empirically.

174 **Simulating characteristics** We build in correlation across time among factors by drawing normal
 175 random numbers for each $1 \leq i \leq N$ and $1 \leq j \leq P_c$, according to :

$$\bar{c}_{ij,t} = \rho_j \bar{c}_{ij,t-1} + \epsilon_{ij,t}; \quad \rho_j \sim \mathcal{U}(0.5, 1) \quad (9)$$

176 We then build in cross sectional correlation:

$$\hat{C}_t = L\bar{C}_t; \quad B = LL' \quad (10)$$

$$B : = \Lambda\Lambda' + 0.1\mathbb{I}_n, \quad \Lambda_i = (\lambda_{i1}, \dots, \lambda_{i4}), \quad \lambda_{ik} \sim N(0, \lambda_{sd}), \quad k = 1, \dots, 4 \quad (11)$$

177 where B serves as a variance covariance matrix with λ_{sd} its density, and L represents the lower
 178 triangle matrix of B via the Cholesky decomposition. λ_{sd} values of 0.01, 0.1 and 1 were used to
 179 explore increasing degrees of cross sectional correlation. Characteristics are then normalized to be
 180 within $[-1, 1]$ for each $1 \leq i \leq N$ and for $j = 1, \dots, P_c$ via:

$$c_{ij,t} = \frac{2}{n+1} \text{rank}(\hat{c}_{ij,t}) - 1. \quad (12)$$

181 **Simulating macroeconomic series** We consider a Vector Autoregression (VAR) model for x_t , a
 182 3×1 multivariate time series ³:

$$x_t = Ax_{t-1} + u_t; \quad A = 0.95I_3; \quad u_t \sim N(\mu = (0, 0, 0)', \Sigma = I_3)$$

183 **Simulating return series** We consider three different functions for $g(z_{i,t})$:

$$(1) \quad g_1(z_{i,t}) = (c_{i1,t}, c_{i2,t}, c_{i3,t} \times x_t'[3,]) \theta_0 \quad (13)$$

$$(2) \quad g_2(z_{i,t}) = (c_{i1,t}^2, c_{i1,t} \times c_{i2,t}, \text{sgn}(c_{i3,t} \times x_t'[3,])) \theta_0 \quad (14)$$

$$(3) \quad g_3(z_{i,t}) = (1[c_{i3,t} > 0], c_{i2,t}^3, c_{i1,t} \times c_{i2,t} \times 1[c_{i3,t} > 0], \text{logit}(c_{i3,t})) \theta_0 \quad (15)$$

184 where $x_t'[3,]$ denotes the third element of the x_t' vector. $g_1(z_{i,t})$ allows the characteristics to enter
 185 the return equation linearly, and $g_2(z_{i,t})$ and $g_3(z_{i,t})$ allow the characteristics to enter the return
 186 equation interactively and non-linearly. ⁴ θ^0 was tuned such that the predictive R^2 was approximately
 187 5%.

188 The simulation design results in $3 \times 3 = 9$ different simulated datasets, each with $N = 200$ stocks,
 189 $T = 180$ periods and $P_c = 100$ characteristics. Each design was simulated 10 times to assess the
 190 robustness of machine learning algorithms, with the number of simulations kept low for computational
 191 feasibility. We employ the hybrid data splitting approach with a training:validation length ratio of
 192 approximately 1.5 and a test set that is 1 year in length.

193 3.1 Simulation Study Results

194 **Prediction Performance** In general, elastic nets are the best performing model, followed closely
 195 by random forests, then neural networks. All machine learning models were unaffected by cross

³More complex specifications for A were briefly explored, but these did not have a significant impact on results.

⁴(g_1, g_2 correspond to the simulation design used by [?].)

Figure 1: Simulation variable importance, faceted by simulation specification

sectional correlation in terms of prediction performance, and typically had better performance when fitted with respect to quantile loss. Random forests only outperformed the elastic nets on highly non-linear specifications. The neural network models were not observed to outperform any of the machine learning models.

This is in stark contrast to the linear models, which are severely affected by both increasing non-linearities cross sectional correlation. This result is consistent across all loss metrics.

Machine learning models fitted with respect to minimizing MAE (quantile loss) generally perform better, even when evaluated against MSE loss metrics. This is not a surprising result, especially considering the stochastic error design which introduces significant shocks to the returns process. Though the actual difference between the loss metrics between the penalized linear models, random forests and neural networks are very small, when considering the consistency of the results across numerous Monte Carlo simulations, the differences in prediction performance, though small, is robust and significant.

Factor Importance We observe that the elastic net outperforms all other models consistently in terms of assigning the correct relative importance to the true underlying regressors,⁵ even in settings with high cross sectional correlation.

Elastic net models perform the best at identifying the true data generating regressors, and this appears to be mostly robust regardless of cross sectional correlation, though their performance worsens as the data generating process becomes more non-linear. On more difficult specifications, the elastic net models are conservative and typically identify a single regressor as importance - most apparent on the g_2 specification. Occasionally, the elastic nets identified the incorrect covariates, assigned them low relative importance.

The random forests and to a lesser extent the neural networks also correctly identified the correct underlying regressors, but struggled with adequately discerning relative importance among correlated regressors. This was became more apparent as the degree of cross sectional correlation increased (see decreasing relative importance of true underlying regressors in Figures ?? and ?? in Appendix).

The linear models unsurprisingly struggled with factor significance analysis with respect to both increasing cross sectional correlation non-linearities. This highlights the non-robustness and ineffectiveness of using traditional linear regression as documented by the literature; linear models were consistently observed to identify irrelevant regressors as important, especially as the degree of cross sectional correlation increased. Considering that the graphs represent the averaged variable importance metrics over different simulation realisations, this means that on a single simulation realization, the performance of linear models is significantly worse.

4 Empirical analysis

We conduct an empirical study as a final way to corroborate the findings of the properties of machine learning models which we observed in the simulation study. Though our simulation study was aimed at capturing the main features of observed data, the underlying data generating process for empirical returns is unknown. This study thus acts as a robustness check as to how machine learning performs on real world data, which can be significantly more complex and noisy than simulated contexts.

Importantly, we find that our findings from the simulation study are largely corroborated for empirical returns data.

4.1 Data

We begin by obtaining monthly individual price data from CRSP for all firms listed in the NYSE, AMEX and NASDAQ, starting from 1957 (starting date of the S&P 500) and ending in December 2016, totalling 60 years. To build individual factors, we construct a factor set based on the cross

⁵(c_1 .constant, c_2 .constant and $c_3.x_3$ for g_1 and g_2 specifications, and c_1 .constant, c_2 .constant and c_3 .constant for g_3)

section of returns literature. This data was sourced from and is the same data used in [?]. We restrict our dataset to begin from 1993 Q3 and end on 2016 Q4 to alleviate data quality issues. Our individual factor set contains 94 characteristics: 61 updated annually, 13 updated quarterly and 20 updated monthly ⁶.

We detail our cleaning procedure of this dataset. To reduce the size of the dataset and increase feasibility, we only consider non-penny equities traded primarily on the NASDAQ. To achieve a balance between having a dataset with enough data points and variability among factors, the dataset was converted to a quarterly format. Quarterly returns were then constructed using the PRC variable according to:

$$RET_t = (PRC_t - PRC_{t-1}) / PRC_{t-1} \quad (16)$$

We allow all stocks which have a quarterly return to enter the dataset, even if they disappear from the dataset for certain periods. This was primarily done to reduce survivorship bias in the dataset, and also allows for stocks which were unlisted and relisted again to feature in the dataset. ⁷

We then follow [?] and construct eight macroeconomic factors following the variable definitions in [?] (see Table ??). These factors were lagged by one period so as to be used to predict one period ahead quarterly returns. The treasury bill rate was also used from this source to proxy for the risk free rate in order to construct excess quarterly returns.

The two sets of factors were then combined to form a baseline set of covariates, which we define throughout all methods and analysis as:

$$z_{i,t} = (1, x_t)' \otimes c_{i,t} \quad (17)$$

where $c_{i,t}$ is a P_c matrix of characteristics for each stock i , and $(1, x_t)'$ is a $P_x \times 1$ vector of macroeconomic predictors, , and \otimes represents the Kronecker product. $z_{i,t}$ is therefore a $P_x P_c$ vector of features for predicting individual stock returns and includes interactions between stock level characteristics and macroeconomic variables. The total number of covariates in this baseline set is $61 \times (8 + 1) = 549$ ⁸. The final dataset contains 202, 066 individual observations. We note that due to data quality issues, LSTMs, FFORMA and DeepAR are not feasible on empirical data, though the results of the simulation study suggest that even if were to be used, their performance would be underwhelming. ⁹

We mimic the sample splitting procedure used in the simulation study: the dataset was split such that the training and validation sets were split such that the training set was approximately 1.5 times the length of the validation set, in order to predict a test set that is one year in length.

4.2 Empirical Data Results

In general, the empirical results are in remarkable agreement with the those obtained in the simulation study: the penalized linear models general perform the best, with the random forest models offering slightly worse performance. Machine learning models fitted with respect to median quantile loss were similarly observed to typically offer improvements across all machine learning models across all loss metrics.

Prediction Accuracy In general the results of the simulation study were repeated: the elastic net models perform the best, followed by the random forests, then the DFNs, and finally the linear models. We note that the differences between each model using the MSE and MAE loss metrics are much more pronounced on empirical data. Even so, the predictive performance between the elastic net models and the quantile random forests is not particularly large, and we observe the quantile random

⁶The dataset also included 74 Standard Industrial Classification (SIC) codes, but these were omitted due to their inconsistency, and inadequateness at classifying companies, as noted by WRDS

⁷To deal with missing data, any characteristics that had over 20% of their data missing were omitted. Remaining missing data were then imputed using their cross sectional medians for each year. See Appendix for more details.

⁸As the individual and macroeconomic factors can have similar names, individual and macroeconomic factors were prefixed with ind_ and macro_ respectively.

⁹The dataset was not normalized for all methods, as only penalized regression and neural networks are sensitive to normalization. For these two methods, the dataset was normalized such that each predictor column had 0 mean and 1 variance.

281 forests outperforming the elastic nets in the first data sample. We similarly see that machine learning
 282 models perform better when fitted with respect to quantile loss instead of MSE. Most notably, we
 283 start to see the neural network models performing poorly on the empirical data, a direct contradiction
 284 to what has been reported in the literature.

285 The non-robustness of DFNs is amplified on the empirical dataset. This was observed to be somewhat
 286 more common on neural networks fitted with respect to MSE, suggesting that they are indeed very
 287 sensitive to outliers in training data. We similarly observe some evidence that deeper neural networks
 288 perform better, though this result is less apparent due to the lower robustness on empirical data (see
 289 ?? in Appendix for results).

290 **Factor Importance** As the data generating process for empirical returns is unknown, the variable
 291 importance results cannot be directly compared with those of the simulation study. Even so, we
 292 see similar results: the elastic net and random forest models tend to agree on the same subset of
 293 predictors, but the random forest struggles to discern between highly correlated regressors. Similar to
 294 the prediction performance results, neural networks perform poorly.

Figure 2: Empirical individual and macroeconomic factor importance, averaged over all samples

295 The elastic net, random forest and to a lesser extent DFNs tend to pick out the max return and 1 month
 296 momentum factors out of the individual characteristics as important, and the book-to-market factor
 297 out of the macroeconomic factors are important. In general, the variable importance metrics are less
 298 consistent for the random forests, and it should be noted in particular that the random forest tends
 299 to determine factors highly correlated with momentum as important, such as change in momentum,
 300 dollar trading volume and return volatility. Within the macroeconomic factors, penalized linear
 301 models tend to identify the average book to market ratio and the default spread as the most important.
 302 The random forests were inconsistent with the elastic nets, and tended to assign very similar variable
 303 importance metrics to most macroeconomic factors.

304 Interestingly, the linear models assign the controversial dividend price ratio macroeconomic factor
 305 as highly important, a result mirrored only with the neural networks. Their variable importance
 306 for individual factors across different training samples is non-robust, with the important variables
 307 almost completely changing year to year. The linear models consistently identified the controversial
 308 dividend-price ratio as important, a result that was somewhat consistent with the neural networks.

309 The overall results again contradict the results of [?], who conclude that all of the machine methods
 310 agree on the same subset of important factors. Indeed, we only see mild consistency in variable
 311 importance between the elastic nets and random forests on the individual factors only - all other
 312 variable importance metrics were either inconsistent between different models, or non-robust.

313 5 Conclusion

314 Our findings demonstrate that the field of machine learning may offer certain tools to improve stock
 315 prediction and identification of true underlying factors. Penalized linear models and to a lesser extent,
 316 random forests are the best performing methods in the analysis undertaken.

317 Importantly, we find that DFNs fail in the context of stock return prediction, at both prediction
 318 performance and variable importance analysis. This result is consistent across a variety of simulated
 319 datasets, as well as empirical data.

320 Lastly, we find that the top performing models - the elastic nets and random forests, tend to agree
 321 and correctly identify the correct underlying regressors in simulated contexts, and agree on the
 322 same subset of factors which are important in empirical contexts. We find that of all the models
 323 considered, the elastic nets are the most consistent at identifying true underlying regressors through
 324 the simulation study. We find that in the empirical setting, among the individual factors the 1 and
 325 6 month momentum factors are the most powerful predictors of stock returns, according to the
 326 penalized linear models and random forests.

327 The overall findings of this paper differ from the sparse literature on machine learning methods in
 328 empirical finance. However, the performance of the penalized linear models with respect to both out
 329 of sample prediction performance and variable importance analysis is promising, and our findings

330 show that machine learning provides some tools which may aid in the problems of stock return
331 prediction and risk factor selection in the financial world.