

---

# Evaluation of Machine Learning in Empirical Asset Pricing

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1        Several recent studies have claimed that machine learning methods provide superior  
2        predictive accuracy of asset returns, relative to simpler modelling approaches, and  
3        can correctly identify factors needed to price portfolio risk. Herein, we demonstrate  
4        that this performance is critically dependent on several features of the data being  
5        analysed; including, the training/test sample split, the frequency at which the data  
6        is observed, and the chosen loss-function. In contrast to existing studies, which  
7        claim that neural nets provide superior predictive accuracy, through a series of  
8        realistic examples that mimics the stylized facts of asset returns, we demonstrate  
9        that neural methods are easily outperformed by simpler methods, such as random  
10       forest and elastic net

## 11    1    Introduction

12    The dominance of machine learning methods in terms of predictive accuracy has now begun to filter  
13    into the application and assessment of asset pricing. The most common application of machine  
14    learning methods within finance are for portfolio construction, asset price prediction, and factor  
15    selection.

16    Several studies have now used machine learning techniques to analyze the cross-section of asset  
17    returns and produce portfolios that can capture nonlinear information in the cross-section of asset  
18    returns. Mortiz and Zimmermann (2016) use tree-based methods in an attempt to understand which  
19    firm-level characteristics best predict the cross-section of stock returns, and where this information  
20    can then be used within portfolio sorting to help mitigate risk. Similarly, Messemmer (2017) uses deep  
21    feedforward neural nets (DFNs) to construct portfolios and predict the returns across a cross-sections of  
22    US asset returns. Similarly, While Messemmer (2017) demonstrates that such DFNs can better capture  
23    nonlinear information, and outperform portfolios generated from linear benchmarks, the author does  
24    claim that deep learning methods are the best methods to exploit these nonlinear interactions.

25    In addition, several studies have now suggested that machine learning methods can produce better  
26    predictions of asset returns ([?], [?] and [?]). In particular, the results of Gu et al. (2019) suggest that,  
27    in terms of predictive performance, as measures by an out-of-sample  $R^2$ , tree-based methods and  
28    shallow neural nets can provide superior predictive accuracy over other machine learning methods,  
29    and simpler model-based approaches. This finding is born out both in terms of simulated data, and  
30    an empirical example with month returns data from 1957 to 2016. [?] attribute this to machine  
31    learning’s ability to evaluate and consider non-linear complexities among factors that cannot be  
32    feasibly achieved using traditional techniques.

33    Similarly, work by Kozak et al, (2018), Freyberger et al. (2018), Feng et al., (2019) and Rapach  
34    and Zhou (2013), demonstrate that machine learning methods can “systematically evaluate the  
35    contribution to asset pricing of any new factor” used within an existing linear asset pricing structure.

36 In addition, Gu et al. (2019) use variable importance metrics to demonstrate the quantify the  
37 differential impact of factors across a large set of possible factors available for asset pricing. As such,  
38 machine learning methods can be used, *en mass*, to consistently evaluate the ability of various factors  
39 to help price portfolio risk. Such work is particularly useful given the literatures seeming obsession  
40 with the XXX and constructing such factors: as of 2014, quantitative trading firms were using 81  
41 factor models (Hsu and Kalesnik, 2014), while Harvey and Liu (2019) currently document that well  
42 over 600 different factors have been suggested in the literature.

43 While the above studies all demonstrate the potential benefits of machine learning methods within  
44 empirical finance, it is unclear whether the findings in these papers are easily generalizable to: one,  
45 different training and validation periods; two, different sampling frequencies, which result in stock  
46 returns with significant different characteristics (e.g., daily volatility is significantly higher than  
47 monthly volatility); and three, different loss-measures of predictive accuracy. The answer to such  
48 questions are particularly pertinent given that the machine learning literature has already documented  
49 the difficulties of certain methods, including those references above, in dealing with data that displays  
50 the stylized facts of asset returns. For instance, methods such as penalized regression and tree-based  
51 models assume a form of conditional independence between observations, which is violated by the  
52 state dependence that exists within, and across, asset returns. In addition, it has already noted that  
53 training more standard types of neural networks, such as the feed forward kind considered in Gu  
54 et al, becomes particularly difficult when data displays strong dependence, ([?]). In addition, more  
55 complex machine learning approaches require extremely large amounts of data, as well as specialized  
56 sample splitting and cross-validation schemes, to deal with possible model over-fitting.

57 In some ways, existing applications of machine learning to empirical asset pricing have either over-  
58 looked, downplayed, or simply ignored the importance of the above issues. For example, Messemmer  
59 (2017) and [?] use cross validation as part of their model building procedures, thereby destroying  
60 the temporal ordering of data. In addition, [?] and Messemmer (2017) produce models using training  
61 samples that end much earlier than the data sets which they ultimately produce forecasts for: in the  
62 case of Messemmer (1970), the training period ends in 1981, while the which ends in the 1970s to  
63 ultimately produce forecasts for the most recent 30 years; in the case of [?], the training ends in  
64 the 1970sm, with predictions ultimately produced only for the period of returns from 1987-2016.  
65 This is particularly worrying as the factors driving daily or monthly returns in the 1980s, and starkly  
66 different than those driving returns in, say, 2001 onwards. However, both of these papers suggest that  
67 the training and validation sets used for the various methods does not impact the test set results.

68 While some combination of machine learning methods can undoubtedly lead to better performance  
69 than simpler model-based solutions, a more systematic treatment on the ability of these methods to 1)  
70 accurately detect significant factors; and 2) accurately predict returns according to a range of loss  
71 measures, must be formulated before researchers can rely on such methods in practice. The goal of  
72 this paper is to bridge this gap and thereby provide a systematic, rigorous, realistic, and reproducible,  
73 study on the performance of several machine learning methods that have been used in empirical asset  
74 pricing.

75 First, through a rigour simulation study, which captures the stylized facts of asset reutnrns, we give an  
76 in-depth comparison of several machine learning methods used in the litarture. The simulation study  
77 explicitly explores how different aspects of financial data such as persistence in regressors, cross  
78 sectional correlation and different complexities of data generating process can affect a methods ability  
79 to: 1) accurately predict future returns across a range of loss measures; and 2) correctly identify  
80 the significant factors driving returns. In contrast to existing findings, in this realistic simulation  
81 design, we find that neural nets procedures, such as feedforward nets, LSTM (CITE), and DeepAR  
82 models (CITE), are among the worst performing methods, while tree-based methods and elastic net  
83 are among the best performing methods. We also demonstrqate that this result is consistent across  
84 various levels of volatilty, cross-sectional correlation, return signal, and across differnt loss functions.  
85 In addition, we demonstrate that elastic net and tree-based methods also outperform neural net based  
86 approach in terms of correclytly identifying significant factors.

87 Next, we validate these findings using a empirical data set of asset returns that cosiders monthly  
88 individual price data from CRSP for all firms listed in the NYSE, AMEX and NASDAQ. The starting  
89 period of the data is January first 1957 (starting date of the S&P 500) and the ending date is December  
90 2016, totalling 60 years. A set of 549 possible factors are used to explain the cross-section of returns.  
91 We pay carful attention to the training and test split, and only use the last fourteen years of month

92 returns to evaluate the different machine learning methods. The results found in the empirical study  
 93 agree completely with those in the aforementioned simulation study: across all machine learning  
 94 methods, neural net based procedure perform the worst across various loss functions, while tree-based  
 95 methods and elastic net perform the best.

96 The results of this study suggest that great care and diligence is required if one wishes to implement  
 97 machine learning methods within empirical finance. Indeed, our results suggest that the efficacy of  
 98 machine learning methods within empirical finance depends are highly-dependent on the samples used  
 99 for training and testing, the loss functions used for evaluation, and the specific nature of the data  
 100 series one wishes to predict. As such, while potentially quite useful in empirical finance, machine  
 101 learning methods are not necessarily a panacea to correctly predict future asset prices or to correctly  
 102 disentangle which factors are relevant.

103 The remainder of the paper is organized as follows....

## 104 2 Model and Methods

### 105 2.1 Statistical Model

106 In this section we briefly discuss the statistical model considered for asset returns. Excess monthly  
 107 returns on asset  $i$ ,  $i = 1, \dots, n$ , at time  $t$ ,  $t = 1, \dots, T$ , are assumed to evolve in an additive fashion:

$$r_{i,t+1} = E(r_{i,t+1}|\mathcal{F}_t) + \epsilon_{i,t+1}, \quad E(\epsilon_{i,t+1}|\mathcal{F}_t) = 0 \quad (1)$$

108 where  $\mathcal{F}_t$  denotes the observable information at time  $t$ , and  $\epsilon_{i,t+1}$  is a martingale difference sequence  
 109 (hereafter, mds). We further consider that the conditional mean of returns is an unknown function of  
 110 a  $P$ -dimensional vector of features, assumed measurable at time  $t$ , such that

$$E(r_{i,t+1}|\mathcal{F}_t) = g(z_{i,t}) \quad (2)$$

111 The features, or predictors,  $z_{i,t}$  are assumed to be composed of time- $t$  information, and depends only  
 112 the characteristics of stock  $i$ . It is not assumed that all  $z_{i,t}$  are present within the function  $g(\cdot)$  across  
 113 all  $i$  units. That is, the function  $g(\cdot)$  need not depend on the same  $z_{i,t}$  as  $i$  varies. The assumption  
 114 that the information set can be characterized by the variables  $z_{i,t}$  without dependence on the  $j \neq i$   
 115 return units, is reasonable given that the collection of  $z_{i,t}$  is rich enough.

116 In what follows, we represent the space of possible features as the Kronecker product of two pieces

$$z_{i,t} = x_t \otimes c_{i,t} \quad (3)$$

117 where the variables  $c_{i,t}$  represent a  $P_c \times 1$  vector of individual-level characteristics for return  $i$ , and  
 118  $x_t$  represents a  $P_x \times 1$  vector of macroeconomic predictors, and  $\otimes$  represents the Kronecker product.  
 119 Thus, for  $P = P_c \cdot P_x$ ,  $z_{i,t}$  represents a  $P \times 1$  feature space that can be used to approximate the  
 120 unknown function  $g(\cdot)$ .

### 121 2.2 Methods

122 Given features  $z_{i,t}$ , the goal of any machine learning method is to approximate the unknown function  
 123  $g(\cdot)$  in 1. Broadly speaking, how different ML methods choose to approximate this function depends  
 124 on three components:

- 125 1. the model used to make predictions,<sup>1</sup>
- 126 2. the regularization mechanism employed to mitigate over-fitting;
- 127 3. a loss function that penalized poor predictions.

128 To ensure the results of ML different methods will be comparable, we fix both the regularization  
 129 mechanisms and loss functions used within each method, and allow only the models used for  
 130 prediction to vary. This approach seeks to ensure that performances in one method, relative to another,  
 131 are based on the model structure and not to some feature of how the models were fit. To this end, we  
 132 first discuss points 2. and 3. above, and then briefly present the models used for our comparison.

<sup>1</sup>The model used by the ML method need not correspond to the statistical models assumed to describe the data. Herein, our goal will not be to assess the ‘‘accuracy’’ of the statistical model, but to determine how different ML methods accurately determine the salient features of this model.

133 **Loss functions:** The choice of loss function used to fit the ML methods is instrumental in the  
 134 methods ultimate performance. Herein, we consider two separate loss functions: Mean Absolute  
 135 Error (MAE) and Mean Squared Error (MSE):

$$\text{MAE} = \frac{1}{n} \sum_{j=i}^n |y_j - \hat{y}_j| \text{ and } \text{MSE} = \frac{1}{n} \sum_{j=i}^n (y_j - \hat{y}_j)^2,$$

136 We consider both loss functions since MAE is less sensitive to outliers in the data which financial  
 137 returns are known to exhibit, and which are caused by extreme market movements. Given this, we  
 138 expect MAE to produce predictive results that are more robust to such outlier events.

139 **Mitigating over-fitting:** ML methods guard against over-fitting by emphasizing out-of-sample  
 140 performance. To this end, observed data is split into “training”, “validation” and “test” set. Since  
 141 returns data is intrinsically dependent, when constructing such a split we must consider a schema that  
 142 respects this dependence structure.

143 Throughout our experiments/applications, to balance computation and accuracy, we use a hybrid  
 144 “rolling window” and “recursive” approach to training/validation/test splits: for each model refit, the  
 145 training set is increased by one year observations, i.e., 12 monthly observations; the validation set is  
 146 fixed at one year and moves forward (by one year) with each model refit; predictions are generated  
 147 using that model for the subsequent year.

148 **Models** The remaining specification for the ML methods is the chosen model used to generate  
 149 predictions. Herein, we consider a host of different models: including elastic net (Hastie et al., XXX),  
 150 Random forest (XXX), feed-forward neural nets (XXX), LSTM (XXX), DeepAR models (XXX).  
 151 To keep the details as brief as possible, we give full details on each model and certain features of its  
 152 implementation used in this work in the appendix. For each of the different methods, we consider  
 153 two variants, one based on the MAE loss and one based on the MSE.

## 154 2.3 Model evaluation measures

155 **Predictive accuracy** Predictive performance for individual excess returns are assessed using Mean  
 156 Absolute Error (MAE), Mean Squared Error (MSE) (evaluated over the test set) and an out-of-sample  
 157  $R^2$  measure. While out-of-sample  $R^2$  is a common measure, there is no universally agreed-upon  
 158 definition. As such, we explicitly state the version employed herein as

$$R_{OOS}^2 = 1 - \frac{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \hat{r}_{i,t+1})^2}{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \bar{r}_{i,t+1})^2} \quad (4)$$

159 where  $\mathcal{T}_3$  indicates that the fits are only assessed on the test sub-sample, which is never used for  
 160 training or tuning.

161 Since  $R^2$  is based on in-sample-fit of a linear model, this measure is less meaningful for most of the  
 162 ML methods considered in this paper. However, we report this measure since this measure has also  
 163 been considered in other applications of ML to empirical finance (see, e.g., Gu et al., 2019).

164 **Factor Selection** An important aspect of empirical finance is the understanding of which features  
 165 drive risk. That is, which features are explicitly represented within  $z_{i,t}$  and can thus be used to help  
 166 price risk using equation 1. To this end, we define a simple variable importance (VI) measure to be  
 167 applied across all ML methods in this research. To this end, we mirror the measure produced in [?]  
 168 and define  $VI_j$  as the reduction in predictive  $R^2$  from setting all values of predictor  $j$  to 0, while  
 169 holding the remaining model estimates fixed. Each  $VI_j$  is then normalized to sum to 1.

170 However, as  $VI_j$  can sometimes be negative, we shift  $VI_j$  by the smallest  $VI_j$  plus a small constant,  
 171 then dividing by this sum to alleviate numerical issues<sup>2</sup>. The resulting VI measure is then.

$$VI_{j,norm} = \frac{VI_j + \min(VI_j) + o}{\sum VI_j + \min(VI_j) + o} \quad ; \quad o = 10^{-100} \quad (5)$$

<sup>2</sup>This mechanism was chosen because the other popular normalization mechanism “softmax” was observed to be unable to preserve the distances between each original  $VI_j$ , making discernment between each  $VI_j$  difficult.

### 3 Simulation Study

We begin with the simulation study as a way to explore how machine learning performs with specific regards to the characteristics present in financial data. By specifying a simulation design with the desired characteristics, a controlled and well understood environment is available for explicitly testing the predictive performance and factor selection capabilities of the machine learning models considered. In particular, a simulation design which incorporated the following characteristics of financial data was needed:

- Low signal to noise ratio
- Stochastic volatility in errors (including large random external shocks and volatility clustering)
- Persistence in regressors
- Cross sectional correlation (multicollinearity) in regressors

#### 3.0.1 Overall Design

We first construct our simulation design by considering the design considered by [?], which consists of three overall components:

- Simulate individual firm and macroeconomic factors
- Enter these factors into a latent, true data generating process
- Produce the overall returns process as this latent generative process plus an error process

However, [?]'s specification has two main issues: the factors which enter the return equation are uncorrelated across each stock  $i$ , and the error process specified amounts to a white noise, constant volatility specification. As noted by [?] and many others, this is not what is observed in practice.

We therefore simulate a latent factor model with a stochastic volatility process for excess returns  $r_{t+1}$ , for  $t = 1, \dots, T$ :

$$r_{i,t+1} = g(z_{i,t}) + \beta_{i,t+1}v_{t+1} + e_{i,t+1}; \quad z_{i,t} = (1, x_t)' \otimes c_{i,t}, \quad \beta_{i,t} = (c_{i1,t}, c_{i2,t}, c_{i3,t}) \quad (6)$$

$$e_{i,t+1} = \sigma_{i,t+1}\varepsilon_{i,t+1}; \quad (7)$$

$$\log(\sigma_{i,t+1}^2) = \omega + \gamma \log(\sigma_t^2) + \sigma_u u; \quad u \sim N(0, 1) \quad (8)$$

Let  $v_{t+1}$  be a  $3 \times 1$  vector of errors, and  $w_{t+1} \sim N(0, 1)$  and  $\varepsilon_{i,t+1} \sim N(0, 1)$  scalar error terms.

The matrix  $C_t$  is an  $N \times P_c$  matrix of latent factors, where the first three columns correspond to  $\beta_{i,t}$ , across the  $1 \leq i \leq N$  dimensions, while the remaining  $P_c - 3$  factors do not enter the return equation. The  $P_x \times 1$  vector  $x_t$  is a  $3 \times 1$  multivariate time series, and  $\varepsilon_{t+1}$  is a  $N \times 1$  vector of idiosyncratic errors.

The parameters of these were tuned such that the annualized volatility of each return series was approximately 22%, as is often observed empirically.

Note that we also reproduce [?]'s error specification as a case where there is no stochastic volatility:

$$v_{t+1} \sim N(0, 0.05^2 \times I_3) \quad (9)$$

$$e_{i,t+1} \sim t_5(0, 0.05^2) \quad (10)$$

### 203 3.0.2 Simulating Characteristics

204 A simulation mechanism for  $C_t$  that gives some correlation across the factors and across time was  
 205 used. We build in correlation across time among factors by drawing normal random numbers for each  
 206  $1 \leq i \leq N$  and  $1 \leq j \leq P_c$ , according to

$$\bar{c}_{ij,t} = \rho_j \bar{c}_{ij,t-1} + \epsilon_{ij,t}; \quad \rho_j \sim \mathcal{U}\left(\frac{1}{2}, 1\right) \quad (11)$$

207 To build in cross sectional correlation, we define the positive-semidefinite matrix  $B$ :

$$B := \Lambda \Lambda' + \frac{1}{10} \mathbb{I}_n, \quad \Lambda_i = (\lambda_{i1}, \dots, \lambda_{i4}), \quad \lambda_{ik} \sim N(0, \lambda_{sd}), \quad k = 1, \dots, 4 \quad (12)$$

208 to serve as a variance covariance matrix with  $\lambda_{sd}$  controlling the density of the matrix, and hence  
 209 degree of cross sectional correlation.  $\lambda_{sd}$  values of 0.01, 0.1 and 1 were used to explore increasing  
 210 degrees of cross sectional correlation.

211 To build this into our  $N \times P_c$  characteristics matrix  $\bar{C}_t$ , we simulate characteristics according to

$$\hat{C}_t = L \bar{C}_t; \quad B = LL' \quad (13)$$

212 where  $L$  represents the lower triangle matrix of  $B$  using the Cholesky decomposition.

213 Finally, the "observed" characteristics for each  $1 \leq i \leq N$  and for  $j = 1, \dots, P_c$  are constructed  
 214 according to:

$$c_{ij,t} = \frac{2}{n+1} \text{rank}(\hat{c}_{ij,t}) - 1. \quad (14)$$

215 with the rank transformation normalizing all predictors to be within  $[-1, 1]$ .

### 216 3.0.3 Simulating Macroeconomic Series

217 For simulation of  $x_t$ , a  $3 \times 1$  multivariate time series, we consider a Vector Autoregression (VAR)  
 218 model<sup>3</sup>:

$$x_t = Ax_{t-1} + u_t; \quad A = \begin{pmatrix} .95 & 0 & 0 \\ 0 & .95 & 0 \\ 0 & 0 & .95 \end{pmatrix} \quad u_t \sim N(\mu = (0, 0, 0)', \Sigma = I_3)$$

### 219 3.0.4 Simulating Return Series

220 We consider three different functions for  $g(z_{i,t})$ :

$$(1) \quad g_1(z_{i,t}) = (c_{i1,t}, c_{i2,t}, c_{i3,t} \times x'_t[3,]) \theta_0 \quad (15)$$

$$(2) \quad g_2(z_{i,t}) = (c_{i1,t}^2, c_{i1,t} \times c_{i2,t}, \text{sgn}(c_{i3,t} \times x'_t[3,])) \theta_0 \quad (16)$$

$$(3) \quad g_3(z_{i,t}) = (1[c_{i3,t} > 0], c_{i2,t}^3, c_{i1,t} \times c_{i2,t} \times 1[c_{i3,t} > 0], \text{logit}(c_{i3,t})) \theta_0 \quad (17)$$

221 where  $x'_t[3,]$  denotes the third element of the  $x'_t$  vector.

222  $g_1(z_{i,t})$  allows the characteristics to enter the return equation linearly, and  $g_2(z_{i,t})$  allows the  
 223 characteristics to enter the return equation interactively and non-linearly. The true underlying

<sup>3</sup>Other more complex and interactive matrix specifications of  $A$  were briefly explored, but these did not appear to have a significant impact on results. More complex designs were observed to only affect the variable importance metrics, but to an insignificant degree

regressors for these specifications are  $(c_{i1,t}, c_{i2,t}, c_{i3,t} \times x'_t[3, ])$ . These two specifications correspond to the simulation design used by [?].

$g_3(z_{i,t})$  allows the characteristics to enter in a complex and non-linear fashion. The true underlying regressors for this specification are  $(c_{i1,t}, c_{i2,t}, c_{i3,t})$ .

It should be noted however, that because  $g_2(z_{i,t})$  has a large part of its signal entering through a  $\text{sgn}$  function, this should make it the most difficult to estimate given the regressors and resulting returns process.


$\theta^0$  was tuned such that the predictive  $R^2$  was approximately 5%.

The simulation design results in  $3 \times 3 = 12$  different simulated datasets, each with  $N = 200$  stocks,  $T = 180$  periods and  $P_c = 100$  characteristics. Each design was simulated 10 times to assess the robustness of machine learning algorithms. The number of simulations was kept low for computational feasibility.

### 3.0.5 Sample Splitting

If viewed as monthly periods,  $T = 180$  corresponds to 15 years. A data splitting scheme similar to the scheme to be used in the empirical data study was used: a training:validation length ratio of approximately 1.5 to begin, and a test set that is 1 year in length. We employ the hybrid growing window approach as described earlier in section ?? (see Figure 13 for a graphical representation).

Figure 1: Sample Splitting Procedure



Other schemes in the forecasting literature such as using an “inner” rolling window validation loop to find the best hyperparameters on average, finally aggregating them in an “outer” loop for a more robust error were considered but not implemented for a variety of reasons. Firstly, many of the models were computationally too intensive for this to be feasible. More importantly, during the model fitting process it was observed that the optimal hyperparameters for the different rolling windows were highly unstable (see Appendix). Thus, this would have made the selection of the best hyperparameters on average across all windows significantly less meaningful.

## 3.1 Simulation Study Results

Overall, in the simulation study we observe that in general elastic nets are the best performing model, followed closely by random forests, then neural networks. All machine learning models were unaffected by cross sectional correlation in terms of prediction performance, and had better performance when fitted with respect to quantile loss, in stark contrast to linear models. The random forest only outperformed the elastic nets on highly non-linear specifications. The neural network models were not observed to outperform any of the machine learning models. We note that most of these results contradict the sparse literature, and in particular, the results reported by [?], even on their proposed simulation design.

### 3.1.1 Prediction Performance

Looking at the prediction performance of different models, we find that in general, penalized linear models performed the best, followed extremely closely by random forests and then neural networks, which all outperform the baseline linear models. Most importantly, we observe that for machine learning models, cross sectional correlation does not seem to affect prediction performance by much.

262 This is in stark contrast to the linear models, whose prediction performance is severely affected by  
 263 both non-linearities, and increasing cross sectional correlation. This result is consistent across all loss  
 264 metrics, and is most obvious when looking at the out-of-sample R-squared metrics.

265 Machine learning models fitted with respect to minimizing MAE (quantile loss) generally perform  
 266 better, even when evaluated against MSE loss metrics. This is not a surprising result, especially  
 267 considering the stochastic error design which introduces significant shocks to the returns process,  
 268 leading to large outliers which the mean squared error metric is more sensitive to. Though the actual  
 269 difference between the loss metrics between the penalized linear models, random forests and neural  
 270 networks are very small, when considering the consistency of the results across numerous Monte  
 271 Carlo simulations, the differences in prediction performance, though small, is robust and significant.

Figure 2: Simulation Test MAE

Figure 3: Simulation Test MSE, see Figure ?? for naming conventions

Figure 4: Simulation Test MSE, see Figure ?? for naming conventions

272 Across all specifications with a stochastic volatility component, we observe a decrease in prediction  
 273 performance as the sample size increased according to the expanding window approach implemented.  
 274 This is likely due a larger sample having a higher chance to experiences external shocks due to the  
 275 stochastic volatility process, and thus a higher chance to experience large outliers in the training  
 276 sample, leading to worse prediction performance. This indicates that machine learning performs  
 277 poorly when the training data supplied contains more large outliers, and is still sensitive to such  
 278 outliers even with the use of regularization and robust loss functions.

279 Focusing on the neural networks, we clearly see that they do not outperform any of the other  
 280 machine learning models, even when the underlying data generating process is non-linear. This  
 281 directly contradicts the result which [?] find, even when considering the design with no cross  
 282 sectional correlation and stochastic volatility (top row in graphs), which corresponds to their exact  
 283 specification but with a multivariate macroeconomic series. We also find consistent evidence that  
 284 deeper architectures provide better prediction performance, another result which contradicts [?]'s  
 285 conclusions that shallow learning may be better.

### 286 3.1.2 Factor Importance

287 Focusing on the factor selection capabilities of different models by looking at variable importance  
 288 metrics however reveals a more interesting comparison between models. We can clearly observe  
 289 that the elastic net outperforms all other models consistently in terms of assigning the correct  
 290 relative importance to the true underlying regressors<sup>4</sup>, even in settings with very high cross sectional  
 291 correlation.

292 In terms of finding the correct underlying regressors, we find that the penalized linear models perform  
 293 the best at identifying the true data generating regressors, and that this appears to be mostly robust  
 294 regardless of the amount of cross sectional correlation in the factor set. The penalized linear models  
 295 are not perfect and their ability to correctly identify underlying regressors worsens as the data  
 296 generating process becomes more non-linear. On these more difficult specifications, the penalized  
 297 linear models are generally very conservative, sometimes only identifying a single covariate as  
 298 important. This is most apparent on the  $g_2$  specification. Occasionally, the elastic nets identified the  
 299 incorrect covariates completely, though the relative importance assigned to them was small.

300 The random forests and to a lesser extent the neural networks also correctly identified the correct  
 301 underlying regressors, but struggled with adequately discerning relative importance among correlated  
 302 regressors. This was became more apparent as the degree of cross sectional correlation increased (see  
 303 decreasing relative importance of true underlying regressors in Figures ?? and ?? in Appendix). In  
 304 the case of the random forests, this is to be expected, likely due to how the random forest algorithms  
 305 work. The random forest algorithm is an ensemble of tree models, with each tree model only having  
 306 access to a subset of all available predictors. If this subset does not include the true data generating

<sup>4</sup>( $c_1$ .constant,  $c_2$ .constant and  $c_3.x_3$  for  $g_1$  and  $g_2$  specifications, and  $c_1$ .constant,  $c_2$ .constant and  $c_3$ .constant for  $g_3$ )



307 predictor, then that particular tree will likely select the predictors which have the highest correlation  
308 with the true data generating predictor instead. Thus, the resulting ensemble model is likely to believe  
309 that cross sectionally correlated predictors are important, relative to the true underlying regressor.  
310 Due to the complexity of the neural networks, there does not exist a similar intuitive explanation for  
311 their factor selection ability.

Figure 5: Simulation Variable Importance averaged across all test samples

Note: faceted by data generating specification at the top, and degree of cross-sectional correlation on the right. Model is on the y axis and follows the naming convention in Figure ??.  
Factor is on the x axis, and only contains the top 30 important factors, as measured by average importance across all samples and models. Note that the four true underlying factors are bolded.

The linear models, unsurprisingly, struggled with factor significance analysis with respect to both increasing cross sectional correlation and increasing non-linearities. This highlights the non-robustness and ineffectiveness of using traditional linear regression as documented by the literature; linear models were consistently observed to identify irrelevant regressors as important, especially as the degree of cross sectional correlation increased. Considering that the graphs represent the averaged variable importance metrics over different simulation realisations which each have random patterns of cross sectional correlation, this means that on a single simulation realization, the performance of linear models is significantly worse.

The overall high performance of the elastic net models may be somewhat surprising given its relative simplicity compared to other machine learning models. However, when recalling that elastic net models are the only machine learning models which are specifically noted to perform well on datasets with high degrees of multicollinearity, the result is perhaps less surprising. Indeed, random forest and neural network models are only noted to be better at capturing non-linear relationships in independent and identically distributed data, a property which we observe on the non-linear specification, and only for the random forests.

Of particular note are the instability of the machine learning models' hyperparameters across different training samples. For the elastic nets, the optimal value for  $\alpha$  is generally 1 (corresponding to LASSO and thus a sparse representation), but it was not uncommon to observe  $\alpha$  values swinging between values close to 0 (corresponding to ridge regression, and thus a dense representation) to 1 as the training sample moved forwards in time. As the penalized linear models consistently performed the best and still remained able to correctly identify the true covariates this is not a large issue, but it should be noted that this can lead to interpretation issues. For the random forests, it was similarly observed that the optimal value for *mtry* (the number of variables subsetted) and *nodesize* was highly non-robust. Again, given that the final prediction performance was consistent this is not a large issue, but can lead to some interpretation issues.

## 4 Empirical analysis

We conduct an empirical study as a final way to corroborate the findings of the properties of machine learning models which we observed in the simulation study. Though our simulation study was aimed at capturing the main features of observed data, the underlying data generating process for empirical returns is unknown. This study thus acts as a robustness check as to how machine learning performs on real world data, which can be significantly more complex and noisy than simulated contexts. Our two studies together can be thought of a repeated sampling exercise in exploring how machine learning methods perform on datasets which feature the "stylized facts" of empirical returns. This empirical study also acts as a final validation against what has been reported in the literature.

Importantly, we find that our findings from the simulation study are largely corroborated for empirical returns data.

### 4.1 Data

We begin by obtaining monthly individual price data from CRSP for all firms listed in the NYSE, AMEX and NASDAQ, starting from 1957 (starting date of the S&P 500) and ending in December 2016, totalling 60 years. To build individual factors, we construct a factor set based on the cross section of returns literature. This data was sourced from and is the same data used in [?]. Like our initial returns sample, it begins in March 1957 and ends in December 2016, totalling 60 years. It contains 94 stock level characteristics: 61 updated annually, 13 updated quarterly and 20 updated monthly, in addition to 74 industry dummies corresponding the the first two digits of the Standard Industrial Classification (SIC) codes. It is noted that this dataset so far contains all securities traded, including those with a CRSP share code other than 10 or 11 and thus includes instruments such as REITs and mutual funds, and those with a share price of less than \$5.

We detail our cleaning procedure of this dataset. To reduce the size of the dataset and increase feasibility, the dataset was filtered such that only stocks traded primarily on NASDAQ were included (using the PRIMEXCH variable from WRDS). Then, penny stocks (also referred to as microcaps in the literature) with a stock price of less than \$5 were filtered out, as is commonly done in the literature to reduce variability. Stocks without a share code of 10 or 11 (referring to equities) were

Table 1: Macroeconomic Factors, ([?])

No.	Acronym	Macroeconomic Factor
1	macro_dp	Dividend Price Ratio
2	macro_ep	Earnings Price Ratio
3	macro_bm	Book to Market Ratio
4	macro_ntis	Net Equity Expansion
5	macro_tbl	Treasury Bill Rate
6	macro_tms	Term Spread
7	macro_dfy	Default Spread
8	macro_svar	Stock Variance

364 filtered out, so that securities that are not equities were not included (such as REITs and trust funds).  
 365 The dataset is provided in a monthly format, which means that many of the factors which are updated  
 366 only quarterly or annually have very low levels of variability, which can lead to misleading results in  
 367 the model fitting process. To achieve a balance between having a dataset with enough data points and  
 368 variability among factors, the dataset was converted to a quarterly format. Quarterly returns were  
 369 then constructed using the PRC variable according to actual returns (ie not logged differences):

$$RET_t = \frac{PRC_t - PRC_{t-1}}{PRC_{t-1}} \quad (18)$$

370 We allow all stocks which have a quarterly return to enter the dataset, even if they disappear from the  
 371 dataset for certain periods, as opposed to only keeping stocks which appear continuously throughout  
 372 the entire period. This was primarily done to reduce survivorship bias in the dataset, which can be  
 373 very prevalent in financial data, and also allows for stocks which were unlisted and relisted again  
 374 to feature in the dataset. This has the obvious drawback of introducing some bias in the dataset, as  
 375 attrition in the dataset is likely to be non-random and correlated with the stocks' returns.

376 The sic2 variable, corresponding to the stocks' Standard Industrial Classification (SIC) codes was  
 377 also dropped. The SIC code system suffers from inconsistent logic in classifying companies, and as  
 378 a system built for pre-1970s traditional industries has been slow in recognizing new and emerging  
 379 industries. Indeed, WRDS explicitly cautions the use of SIC codes beyond the use of rough grouping  
 380 of industries, warning that SIC codes are not strictly enforced by government agencies for accuracy,  
 381 in addition to most large companies belonging to multiple SIC codes over time. Because of this  
 382 latter point in particular, there can be inconsistencies on the correct SIC code for the same company  
 383 depending on the data source. Dropping the sic2 variable also reduced the dimensionality of the  
 384 dataset by 74 columns, significant increasing computational feasibility.

385 There existed a significant amount of missing data in the dataset. The dataset's columns were first  
 386 examined, and any characteristics that had over 20% of their data were removed. However, as the  
 387 amount of missing data increases dramatically going further back in time, a balance between using  
 388 more periods at the cost of removing more characteristics versus using less periods but keeping more  
 389 characteristics was needed. 1993 Q3 was determined to be a reasonable time frame to begin the  
 390 dataset, as there was a noticeable increase in data availability and quality after this time. Missing  
 391 characteristics were then imputed using their cross sectional medians for each year.


392 We then follow [?] and construct eight macroeconomic factors following the variable definitions  
 393 in [?]: dividend-price ratio (dp), earnings-price ratio (ep), book-to-market ratio (bm), net equity  
 394 expansion (ntis), Treasury-bill rate (tbl), term spread (tms), default spread (dfy) and stock variance  
 395 (svar). These factors were lagged by one period so as to be used to predict one period ahead quarterly  
 396 returns. The treasury bill rate was also used from this source to proxy for the risk free rate in order to  
 397 construct excess quarterly returns.

398 The two sets of factors were then combined to form a baseline set of covariates, which we define  
 399 throughout all methods and analysis as:

$$z_{i,t} = (1, x_t)' \otimes c_{i,t} \quad (19)$$

400 where  $c_{i,t}$  is a  $P_c$  matrix of characteristics for each stock  $i$ , and  $(1, x_t)'$  is a  $P_x \times 1$  vector of  
 401 macroeconomic predictors, , and  $\otimes$  represents the Kronecker product.  $z_{i,t}$  is therefore a  $P_x P_c$  vector  
 402 of features for predicting individual stock returns and includes interactions between stock level  
 403 characteristics and macroeconomic variables. The total number of covariates in this baseline set is  
 404  $61 \times (8 + 1) = 549^5$ .  
 405 The dataset was not normalized for all methods, as only penalized regression and neural networks  
 406 are sensitive to normalization. For these two methods, the dataset was normalized such that each  
 407 predictor column had 0 mean and 1 variance.  
 408 The final dataset spanned from 1993 Q3 to 2016 Q4 with 202, 066 individual observations.  
 409 We mimic the procedure used in the simulation study. For the sample splitting procedure, the  
 410 dataset was split such that the training and validation sets were split such that the training set was  
 411 approximately 1.5 times the length of the validation set, in order to predict a test set that is one year  
 412 in length.

Figure 6: Empirical Data Sample Splitting Procedure



<sup>5</sup>As the individual and macroeconomic factors can have similar names, individual and macroeconomic factors were prefixed with ind\_ and macro\_ respectively.

## 4.2 Empirical Data Results

In general, the empirical results are in remarkable agreement with the those obtained in the simulation study. In particular, we similarly see that the penalized linear models generally performing the best, with the random forest models offering slightly worse performance, occasionally outperforming penalized linear models. Machine learning models fitted with respect to median quantile loss were similarly observed to typically offer improvements across all machine learning models across all loss metrics.

### 4.2.1 Prediction Accuracy

In terms of prediction accuracy, we can see that in general the results of the simulation study were repeated: the elastic net models perform the best, followed by the random forests, then the neural networks, and finally the linear models. We note that the differences between each model using the MSE and MAE loss metrics are much more pronounced on empirical data. Even so, the predictive performance between the elastic net models and the quantile random forests is not particularly large, and we observe the quantile random forests outperforming the elastic nets in the first data sample. We similarly see that machine learning models perform better when fitted with respect to quantile loss instead of MSE. Most notably, we start to see the neural network models performing poorly on the empirical data, a direct contradiction to what has been reported in the literature.

Figure 7: Empirical Test MSE

Figure 8: Empirical Test MAE

Figure 9: Empirical Test Predictive R-Squared

Focusing on the neural networks specifically, their non-robustness is amplified on the empirical dataset, with some neural networks in some samples even performing worse than linear models. This was observed to be somewhat more common on neural networks fitted with respect to MSE, suggesting that they are indeed very sensitive to outliers in training data. We similarly observe some evidence that deeper neural networks perform better, though this result is less apparent due to the lower robustness on empirical data (see ?? in Appendix for results).

Interestingly, we do not observe worsening performance as the training sample increases as we did in the simulation study. This suggests that the simulation design may have been too volatile when compared to the specific empirical time periods examined.

### 4.2.2 Factor Importance

As the data generating process for empirical returns is unknown, the variable importance results cannot be directly compared with those of the simulation study. Even so, we see similar results: the elastic net and random forest models tend to agree on the same subset of predictors, but the random forest struggles to discern between highly correlated regressors. Similar to the prediction performance results, neural networks perform poorly.

Figure 10: Empirical Individual Factor Importance, averaged across all training samples

Figure 11: Empirical Macroeconomic Factor Importance, faceted by training sample

If we focus on the two top performing models of elastic net and random forest, we see that they consistently pick out the 1 month and 6 month momentum factors out of the individual characteristics as important, and the book-to-market and default yield spread factors out of the macroeconomic factors are important. In general though, the variable importance metrics are less consistent for the random forests, and it should be noted in particular that the random forest tends to determine factors highly correlated with momentum as important, such as change in moment, dollar trading volume and return volatility. Looking at the macroeconomic factors, penalized linear models tend to identify the average book to market ratio and the default spread as the most important macroeconomic factors.

On the macroeconomic factor set, the random forests were inconsistent with the elastic nets, and tended to assign very similar variable importance metrics to most macroeconomic factors.

The neural networks tended to believe that the market value factor was the most important among the individual factors, a result not repeated by any of the other models considered. Within the macroeconomic factors, the neural networks identified the dividend-price ratio and earnings-price ratio as the most important among the macroeconomic factors, though these results were highly non-robust across different architectures, loss functions and training samples.

Interestingly, we find that the linear models assign the controversial dividend price ratio macroeconomic factor as highly important, a result mirrored only with the neural networks. Their variable importance for individual factors across different training samples is highly non-robust, with the important variables almost completely changing year to year. The linear models consistently identified the controversial dividend-price ratio as important, a result that was somewhat consistent with the neural networks.

The overall results again contradict the results of [?], who conclude that all of the machine methods agree on the same subset of important factors. Indeed, we only see consistency in variable importance between the elastic nets and random forests on the individual factors only - all other variable importance metrics were either inconsistent between different models, or non-robust.

All models considered typically preferred sparse parameterizations. That is, most if not all of the individual factors had little to no importance across all models.<sup>6</sup>

## 5 Conclusion

Our findings demonstrate that the field of machine learning may offer certain tools to improve stock prediction and identification of true underlying factors. Penalized linear models and to a lesser extent, random forests are the best performing methods in the analysis undertaken.

Importantly, we find that the feed-forward neural network architectures considered by [?] fail in the context of stock return prediction, at both prediction performance and variable importance analysis. This result is consistent across a variety of simulated datasets, as well as empirical data. We also find weak evidence that deeper neural networks with more hidden layers perform better, though this result, like the performance of neural networks in general, is non-robust.

Lastly, we find that the top performing models - the penalized linear models and random forests, tend to agree and correctly identify the correct underlying regressors in simulated contexts, as well as agree on the same subset of factors which are important in empirical contexts. We find that of all the most considered, the penalized linear models are the most consistent at identifying true underlying regressors through the simulation study. We find that in the empirical setting, among the individual factors the individual 1 and 6 month momentum factors are the most powerful predictors of stock returns, according to the penalized linear models and random forests.

Across all models except for linear models, we find that minimizing quantile loss yields better prediction performance.

The overall findings of this paper differ from the sparse literature on machine learning methods in empirical finance. However, the performance of the penalized linear models with respect to both out of sample prediction performance and variable importance analysis is promising, and our findings show that machine learning provides some tools which may aid in the problems of stock return prediction and risk factor selection in the financial world.

### 5.1 Retrieval of style files

The style files for NeurIPS and other conference information are available on the World Wide Web at

<http://www.neurips.cc/>

---

<sup>6</sup>Note that because the variable importance here was not evaluated explicitly for each pairwise interaction term, some of the individual factors appear as slightly important. This is because setting an individual factor to zero also sets some of the macroeconomic pairwise terms to zero, increasing its apparent importance.

498 The file `neurips_2020.pdf` contains these instructions and illustrates the various formatting re-  
499 quirements your NeurIPS paper must satisfy.

500 The only supported style file for NeurIPS 2020 is `neurips_2020.sty`, rewritten for  $\text{\LaTeX}$  2 $\epsilon$ .  
501 **Previous style files for  $\text{\LaTeX}$  2.09, Microsoft Word, and RTF are no longer supported!**

502 The  $\text{\LaTeX}$  style file contains three optional arguments: `final`, which creates a camera-ready copy,  
503 `preprint`, which creates a preprint for submission to, e.g., arXiv, and `nonatbib`, which will not  
504 load the `natbib` package for you in case of package clash.

505 **Preprint option** If you wish to post a preprint of your work online, e.g., on arXiv, using the  
506 NeurIPS style, please use the `preprint` option. This will create a nonanonymized version of your  
507 work with the text “Preprint. Work in progress.” in the footer. This version may be distributed as  
508 you see fit. Please **do not** use the `final` option, which should **only** be used for papers accepted to  
509 NeurIPS.

510 At submission time, please omit the `final` and `preprint` options. This will anonymize your  
511 submission and add line numbers to aid review. Please *do not* refer to these line numbers in your  
512 paper as they will be removed during generation of camera-ready copies.

513 The file `neurips_2020.tex` may be used as a “shell” for writing your paper. All you have to do is  
514 replace the author, title, abstract, and text of the paper with your own.

515 The formatting instructions contained in these style files are summarized in Sections 6, 7, and 8  
516 below.

## 517 6 General formatting instructions

518 The text must be confined within a rectangle 5.5 inches (33 picas) wide and 9 inches (54 picas) long.  
519 The left margin is 1.5 inch (9 picas). Use 10 point type with a vertical spacing (leading) of 11 points.  
520 Times New Roman is the preferred typeface throughout, and will be selected for you by default.  
521 Paragraphs are separated by  $\frac{1}{2}$  line space (5.5 points), with no indentation.

522 The paper title should be 17 point, initial caps/lower case, bold, centered between two horizontal  
523 rules. The top rule should be 4 points thick and the bottom rule should be 1 point thick. Allow  $\frac{1}{4}$  inch  
524 space above and below the title to rules. All pages should start at 1 inch (6 picas) from the top of the  
525 page.

526 For the final version, authors’ names are set in boldface, and each name is centered above the  
527 corresponding address. The lead author’s name is to be listed first (left-most), and the co-authors’  
528 names (if different address) are set to follow. If there is only one co-author, list both author and  
529 co-author side by side.

530 Please pay special attention to the instructions in Section 8 regarding figures, tables, acknowledgments,  
531 and references.

## 532 7 Headings: first level

533 All headings should be lower case (except for first word and proper nouns), flush left, and bold.

534 First-level headings should be in 12-point type.

### 535 7.1 Headings: second level

536 Second-level headings should be in 10-point type.

#### 537 7.1.1 Headings: third level

538 Third-level headings should be in 10-point type.

539 **Paragraphs** There is also a `\paragraph` command available, which sets the heading in bold, flush  
540 left, and inline with the text, with the heading followed by 1 em of space.



## 8 Citations, figures, tables, references

These instructions apply to everyone.

### 8.1 Citations within the text

The `natbib` package will be loaded for you by default. Citations may be author/year or numeric, as long as you maintain internal consistency. As to the format of the references themselves, any style is acceptable as long as it is used consistently.

The documentation for `natbib` may be found at

<http://mirrors.ctan.org/macros/latex/contrib/natbib/natnotes.pdf>

Of note is the command `\citet`, which produces citations appropriate for use in inline text. For example,

```
\citet{hasselmo} investigated\dotso
```

produces

Hasselmo, et al. (1995) investigated...

If you wish to load the `natbib` package with options, you may add the following before loading the `neurips_2020` package:

```
\PassOptionsToPackage{options}{natbib}
```

If `natbib` clashes with another package you load, you can add the optional argument `nonatbib` when loading the style file:

```
\usepackage[nonatbib]{neurips_2020}
```

As submission is double blind, refer to your own published work in the third person. That is, use “In the previous work of Jones et al. [4],” not “In our previous work [4].” If you cite your other papers that are not widely available (e.g., a journal paper under review), use anonymous author names in the citation, e.g., an author of the form “A. Anonymous.”

### 8.2 Footnotes

Footnotes should be used sparingly. If you do require a footnote, indicate footnotes with a number<sup>7</sup> in the text. Place the footnotes at the bottom of the page on which they appear. Precede the footnote with a horizontal rule of 2 inches (12 picas).

Note that footnotes are properly typeset *after* punctuation marks.<sup>8</sup>

### 8.3 Figures

All artwork must be neat, clean, and legible. Lines should be dark enough for purposes of reproduction. The figure number and caption always appear after the figure. Place one line space before the figure caption and one line space after the figure. The figure caption should be lower case (except for first word and proper nouns); figures are numbered consecutively.

You may use color figures. However, it is best for the figure captions and the paper body to be legible if the paper is printed in either black/white or in color.

### 8.4 Tables


All tables must be centered, neat, clean and legible. The table number and title always appear before the table. See Table 2.

---

<sup>7</sup>Sample of the first footnote.

<sup>8</sup>As in this example.

Figure 12: Sample figure caption.



Macroeconomic Factors shown on x axis (see Table 3 for definitions)

Table 2: Sample table title

Part		
Name	Description	Size ( $\mu\text{m}$ )
Dendrite	Input terminal	$\sim 100$
Axon	Output terminal	$\sim 10$
Soma	Cell body	up to $10^6$

579 Place one line space before the table title, one line space after the table title, and one line space after  
580 the table. The table title must be lower case (except for first word and proper nouns); tables are  
581 numbered consecutively.

582 Note that publication-quality tables *do not contain vertical rules*. We strongly suggest the use of the  
583 booktabs package, which allows for typesetting high-quality, professional tables:

584 <https://www.ctan.org/pkg/booktabs>

585 This package was used to typeset Table 2.

## 586 9 Final instructions

587 Do not change any aspects of the formatting parameters in the style files. In particular, do not modify  
588 the width or length of the rectangle the text should fit into, and do not change font sizes (except  
589 perhaps in the **References** section; see below). Please note that pages should be numbered.

## 590 10 Preparing PDF files

591 Please prepare submission files with paper size “US Letter,” and not, for example, “A4.”

592 Fonts were the main cause of problems in the past years. Your PDF file must only contain Type 1 or  
593 Embedded TrueType fonts. Here are a few instructions to achieve this.

- 594 • You should directly generate PDF files using `pdflatex`.
- 595 • You can check which fonts a PDF files uses. In Acrobat Reader, select the menu  
596 Files>Document Properties>Fonts and select Show All Fonts. You can also use the program  
597 `pdf fonts` which comes with `xpdf` and is available out-of-the-box on most Linux machines.
- 598 • The IEEE has recommendations for generating PDF files whose fonts are also ac-  
599 ceptable for NeurIPS. Please see [http://www.emfield.org/icuwb2010/downloads/](http://www.emfield.org/icuwb2010/downloads/IEEE-PDF-SpecV32.pdf)  
600 [IEEE-PDF-SpecV32.pdf](http://www.emfield.org/icuwb2010/downloads/IEEE-PDF-SpecV32.pdf)
- 601 • `xfig` “patterned” shapes are implemented with bitmap fonts. Use “solid” shapes instead.

- The `\bbold` package almost always uses bitmap fonts. You should use the equivalent AMS Fonts:

```
\usepackage{amsfonts}
```

followed by, e.g., `\mathbb{R}`, `\mathbb{N}`, or `\mathbb{C}` for  $\mathbb{R}$ ,  $\mathbb{N}$  or  $\mathbb{C}$ . You can also use the following workaround for reals, natural and complex:

```
\newcommand{\RR}{I\!\!R} %real numbers
\newcommand{\Nat}{I\!\!N} %natural numbers
\newcommand{\CC}{I\!\!C} %complex numbers
```

Note that `amsfonts` is automatically loaded by the `amssymb` package.

If your file contains type 3 fonts or non embedded TrueType fonts, we will ask you to fix it.

## 10.1 Margins in L<sup>A</sup>T<sub>E</sub>X

Most of the margin problems come from figures positioned by hand using `\special` or other commands. We suggest using the command `\includegraphics` from the `graphicx` package. Always specify the figure width as a multiple of the line width as in the example below:

```
\usepackage[pdftex]{graphicx} ...
\includegraphics[width=0.8\linewidth]{myfile.pdf}
```

See Section 4.4 in the graphics bundle documentation (<http://mirrors.ctan.org/macros/latex/required/graphics/grfguide.pdf>)

A number of width problems arise when L<sup>A</sup>T<sub>E</sub>X cannot properly hyphenate a line. Please give LaTeX hyphenation hints using the `\-` command when necessary.

## Broader Impact

Authors are required to include a statement of the broader impact of their work, including its ethical aspects and future societal consequences. Authors should discuss both positive and negative outcomes, if any. For instance, authors should discuss a) who may benefit from this research, b) who may be put at disadvantage from this research, c) what are the consequences of failure of the system, and d) whether the task/method leverages biases in the data. If authors believe this is not applicable to them, authors can simply state this.

Use unnumbered first level headings for this section, which should go at the end of the paper. **Note that this section does not count towards the eight pages of content that are allowed.**

## References

References follow the acknowledgments. Use unnumbered first-level heading for the references. Any choice of citation style is acceptable as long as you are consistent. It is permissible to reduce the font size to `small` (9 point) when listing the references. **Note that the Reference section does not count towards the eight pages of content that are allowed.**

[1] Alexander, J.A. & Mozer, M.C. (1995) Template-based algorithms for connectionist rule extraction. In G. Tesauro, D.S. Touretzky and T.K. Leen (eds.), *Advances in Neural Information Processing Systems 7*, pp. 609–616. Cambridge, MA: MIT Press.

[2] Bower, J.M. & Beeman, D. (1995) *The Book of GENESIS: Exploring Realistic Neural Models with the GEneral NEural Simulation System*. New York: TELOS/Springer-Verlag.

[3] Hasselmo, M.E., Schnell, E. & Barkai, E. (1995) Dynamics of learning and recall at excitatory recurrent synapses and cholinergic modulation in rat hippocampal region CA3. *Journal of Neuroscience* **15**(7):5249-5262.

## A Additional details: dimulation design

In this section, we give additional features of the simulation design required to impenent our results. All code and data can be found at XXXX.

## 646 A.1 Simulation Design

647 We begin with the simulation study as a way to explore how machine learning performs with specific regards to  
 648 the characteristics present in financial data. By specifying a simulation design with the desired characteristics, a  
 649 controlled and well understood environment is available for explicitly testing the predictive performance and  
 650 factor selection capabilities of the machine learning models considered. In particular, a simulation design which  
 651 incorporated the following characteristics of financial data was needed:

- 652 • Low signal to noise ratio
- 653 • Stochastic volatility in errors (including large random external shocks and volatility clustering)
- 654 • Persistence in regressors
- 655 • Cross sectional correlation (multicollinearity) in regressors

### 656 A.1.1 Overall Design

657 We first construct our simulation design by considering the design considered by [?], which consists of three  
 658 overall components:

- 659 • Simulate individual firm and macroeconomic factors
- 660 • Enter these factors into a latent, true data generating process
- 661 • Produce the overall returns process as this latent generative process plus an error process

662 However, [?]'s specification has two main issues: the factors which enter the return equation are uncorrelated  
 663 across each stock  $i$ , and the error process specified amounts to a white noise, constant volatility specification. As  
 664 noted by [?] and many others, this is not what is observed in practice.

665 We therefore simulate a latent factor model with a stochastic volatility process for excess returns  $r_{t+1}$ , for  
 666  $t = 1, \dots, T$ :

$$r_{i,t+1} = g(z_{i,t}) + \beta_{i,t+1}v_{t+1} + e_{i,t+1}; \quad z_{i,t} = (1, x_t)' \otimes c_{i,t}, \quad \beta_{i,t} = (c_{i1,t}, c_{i2,t}, c_{i3,t}) \quad (20)$$

$$e_{i,t+1} = \sigma_{i,t+1}\varepsilon_{i,t+1}; \quad (21)$$

$$\log(\sigma_{i,t+1}^2) = \omega + \gamma \log(\sigma_t^2) + \sigma_u u; \quad u \sim N(0, 1) \quad (22)$$

667 Let  $v_{t+1}$  be a  $3 \times 1$  vector of errors, and  $w_{t+1} \sim N(0, 1)$  and  $\varepsilon_{i,t+1} \sim N(0, 1)$  scalar error terms.

668 The matrix  $C_t$  is an  $N \times P_c$  matrix of latent factors, where the first three columns correspond to  $\beta_{i,t}$ , across the  
 669  $1 \leq i \leq N$  dimensions, while the remaining  $P_c - 3$  factors do not enter the return equation. The  $P_x \times 1$  vector  
 670  $x_t$  is a  $3 \times 1$  multivariate time series, and  $\varepsilon_{t+1}$  is a  $N \times 1$  vector of idiosyncratic errors.

671 The parameters of these were tuned such that the annualized volatility of each return series was approximately  
 672 22%, as is often observed empirically.

673 Note that we also reproduce [?]'s error specification as a case where there is no stochastic volatility:

$$v_{t+1} \sim N(0, 0.05^2 \times I_3) \quad (23)$$

$$e_{i,t+1} \sim t_5(0, 0.05^2) \quad (24)$$

### 674 A.1.2 Simulating Characteristics

675 A simulation mechanism for  $C_t$  that gives some correlation across the factors and across time was used. We  
 676 build in correlation across time among factors by drawing normal random numbers for each  $1 \leq i \leq N$  and  
 677  $1 \leq j \leq P_c$ , according to

$$\bar{c}_{ij,t} = \rho_j \bar{c}_{ij,t-1} + \epsilon_{ij,t}; \quad \rho_j \sim \mathcal{U}\left(\frac{1}{2}, 1\right) \quad (25)$$

678 To build in cross sectional correlation, we define the positive-semidefinite matrix  $B$ :

$$B := \Lambda \Lambda' + \frac{1}{10} \mathbb{I}_n, \quad \Lambda_i = (\lambda_{i1}, \dots, \lambda_{i4}), \quad \lambda_{ik} \sim N(0, \lambda_{sd}), \quad k = 1, \dots, 4 \quad (26)$$

to serve as a variance covariance matrix with  $\lambda_{sd}$  controlling the density of the matrix, and hence degree of cross sectional correlation.  $\lambda_{sd}$  values of 0.01, 0.1 and 1 were used to explore increasing degrees of cross sectional correlation.

To build this into our  $N \times P_c$  characteristics matrix  $\bar{C}_t$ , we simulate characteristics according to

$$\hat{C}_t = L\bar{C}_t; \quad B = LL' \quad (27)$$

where  $L$  represents the lower triangle matrix of  $B$  using the Cholesky decomposition.

Finally, the "observed" characteristics for each  $1 \leq i \leq N$  and for  $j = 1, \dots, P_c$  are constructed according to:

$$c_{ij,t} = \frac{2}{n+1} \text{rank}(\hat{c}_{ij,t}) - 1. \quad (28)$$

with the rank transformation normalizing all predictors to be within  $[-1, 1]$ .

### A.1.3 Simulating Macroeconomic Series

For simulation of  $x_t$ , a  $3 \times 1$  multivariate time series, we consider a Vector Autoregression (VAR) model<sup>9</sup>:

$$x_t = Ax_{t-1} + u_t; \quad A = \begin{pmatrix} .95 & 0 & 0 \\ 0 & .95 & 0 \\ 0 & 0 & .95 \end{pmatrix} \quad u_t \sim N(\mu = (0, 0, 0)', \Sigma = I_3)$$

### A.1.4 Simulating Return Series

We consider three different functions for  $g(z_{i,t})$ :

$$(1) \quad g_1(z_{i,t}) = (c_{i1,t}, c_{i2,t}, c_{i3,t} \times x'_t[3,]) \theta_0 \quad (29)$$

$$(2) \quad g_2(z_{i,t}) = (c_{i1,t}^2, c_{i1,t} \times c_{i2,t}, \text{sgn}(c_{i3,t} \times x'_t[3,])) \theta_0 \quad (30)$$

$$(3) \quad g_3(z_{i,t}) = (1[c_{i3,t} > 0], c_{i2,t}^3, c_{i1,t} \times c_{i2,t} \times 1[c_{i3,t} > 0], \text{logit}(c_{i3,t})) \theta_0 \quad (31)$$

where  $x'_t[3,]$  denotes the third element of the  $x'_t$  vector.

$g_1(z_{i,t})$  allows the characteristics to enter the return equation linearly, and  $g_2(z_{i,t})$  allows the characteristics to enter the return equation interactively and non-linearly. The true underlying regressors for these specifications are  $(c_{i1,t}, c_{i2,t}, c_{i3,t} \times x'_t[3,])$ . These two specifications correspond to the simulation design used by [?].

$g_3(z_{i,t})$  allows the characteristics to enter in a complex and non-linear fashion. The true underlying regressors for this specification are  $(c_{i1,t}, c_{i2,t}, c_{i3,t})$ .

It should be noted however, that because  $g_2(z_{i,t})$  has a large part of its signal entering through a  $\text{sgn}$  function, this should make it the most difficult to estimate given the regressors and resulting returns process.

$\theta^0$  was tuned such that the predictive  $R^2$  was approximately 5%.

The simulation design results in  $3 \times 3 = 12$  different simulated datasets, each with  $N = 200$  stocks,  $T = 180$  periods and  $P_c = 100$  characteristics. Each design was simulated 10 times to assess the robustness of machine learning algorithms. The number of simulations was kept low for computational feasibility.

### A.1.5 Sample Splitting

If viewed as monthly periods,  $T = 180$  corresponds to 15 years. A data splitting scheme similar to the scheme to be used in the empirical data study was used: a training:validation length ratio of approximately 1.5 to begin, and a test set that is 1 year in length. We employ the hybrid growing window approach as described earlier in section ?? (see Figure 13 for a graphical representation).

Other schemes in the forecasting literature such as using an "inner" rolling window validation loop to find the best hyperparameters on average, finally aggregating them in an "outer" loop for a more robust error were

<sup>9</sup>Other more complex and interactive matrix specifications of  $A$  were briefly explored, but these did not appear to have a significant impact on results. More complex designs were observed to only affect the variable importance metrics, but to an insignificant degree

Figure 13: Sample Splitting Procedure

Set No.															
3															
2															
1															
Year	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
								Training							
								Validation							
								Test							

considered but not implemented for a variety of reasons. Firstly, many of the models were computationally too intensive for this to be feasible. More importantly, during the model fitting process it was observed that the optimal hyperparameters for the different rolling windows were highly unstable (see Appendix). Thus, this would have made the selection of the best hyperparameters on average across all windows significantly less meaningful.

## A.2 Simulation Study Results

Overall, in the simulation study we observe that in general elastic nets are the best performing model, followed closely by random forests, then neural networks. All machine learning models were unaffected by cross sectional correlation in terms of prediction performance, and had better performance when fitted with respect to quantile loss, in stark contrast to linear models. The random forest only outperformed the elastic nets on highly non-linear specifications. The neural network models were not observed to outperform any of the machine learning models. We note that most of these results contradict the sparse literature, and in particular, the results reported by [?], even on their proposed simulation design.

### A.2.1 Prediction Performance

Looking at the prediction performance of different models, we find that in general, penalized linear models performed the best, followed extremely closely by random forests and then neural networks, which all outperform the baseline linear models. Most importantly, we observe that for machine learning models, cross sectional correlation does not seem to affect prediction performance by much. This is in stark contrast to the linear models, whose prediction performance is severely affected by both non-linearities, and increasing cross sectional correlation. This result is consistent across all loss metrics, and is most obvious when looking at the out-of-sample R-squared metrics.

Machine learning models fitted with respect to minimizing MAE (quantile loss) generally perform better, even when evaluated against MSE loss metrics. This is not a surprising result, especially considering the stochastic error design which introduces significant shocks to the returns process, leading to large outliers which the mean squared error metric is more sensitive to. Though the actual difference between the loss metrics between the penalized linear models, random forests and neural networks are very small, when considering the consistency of the results across numerous Monte Carlo simulations, the differences in prediction performance, though small, is robust and significant.

Figure 14: Simulation Test MAE

Figure 15: Simulation Test MSE, see Figure ?? for naming conventions

Figure 16: Simulation Test MSE, see Figure ?? for naming conventions

Across all specifications with a stochastic volatility component, we observe a decrease in prediction performance as the sample size increased according to the expanding window approach implemented. This is likely due a larger sample having a higher chance to experience external shocks due to the stochastic volatility process, and thus a higher chance to experience large outliers in the training sample, leading to worse prediction performance. This indicates that machine learning performs poorly when the training data supplied contains more large outliers, and is still sensitive to such outliers even with the use of regularization and robust loss functions.

Focusing on the neural networks, we clearly see that they do not outperform any of the other machine learning models, even when the underlying data generating process is non-linear. This directly contradicts the result which [?] find, even when considering the design with no cross sectional correlation and stochastic volatility

(top row in graphs), which corresponds to their exact specification but with a multivariate macroeconomic series. We also find consistent evidence that deeper architectures provide better prediction performance, another result which contradicts [?]'s conclusions that shallow learning may be better.

## A.2.2 Factor Importance

Focusing on the factor selection capabilities of different models by looking at variable importance metrics however reveals a more interesting comparison between models. We can clearly observe that the elastic net outperforms all other models consistently in terms of assigning the correct relative importance to the true underlying regressors<sup>10</sup>, even in settings with very high cross sectional correlation.

In terms of finding the correct underlying regressors, we find that the penalized linear models perform the best at identifying the true data generating regressors, and that this appears to be mostly robust regardless of the amount of cross sectional correlation in the factor set. The penalized linear models are not perfect and their ability to correctly identify underlying regressors worsens as the data generating process becomes more non-linear. On these more difficult specifications, the penalized linear models are generally very conservative, sometimes only identifying a single covariate as important. This is most apparent on the  $g_2$  specification. Occasionally, the elastic nets identified the incorrect covariates completely, though the relative importance assigned to them was small.

The random forests and to a lesser extent the neural networks also correctly identified the correct underlying regressors, but struggled with adequately discerning relative importance among correlated regressors. This was became more apparent as the degree of cross sectional correlation increased (see decreasing relative importance of true underlying regressors in Figures ?? and ?? in Appendix). In the case of the random forests, this is to be expected, likely due to how the random forest algorithms work. The random forest algorithm is an ensemble of tree models, with each tree model only having access to a subset of all available predictors. If this subset does not include the true data generating predictor, then that particular tree will likely select the predictors which have the highest correlation with the true data generating predictor instead. Thus, the resulting ensemble model is likely to believe that cross sectionally correlated predictors are important, relative to the true underlying regressor. Due to the complexity of the neural networks, there does not exist a similar intuitive explanation for their factor selection ability.

---

<sup>10</sup>( $c_1$ .constant,  $c_2$ .constant and  $c_3.x_3$  for  $g_1$  and  $g_2$  specifications, and  $c_1$ .constant,  $c_2$ .constant and  $c_3$ .constant for  $g_3$ )

Figure 17: Simulation Variable Importance averaged across all test samples

Note: faceted by data generating specification at the top, and degree of cross-sectional correlation on the right. Model is on the y axis and follows the naming convention in Figure ??.  
Factor is on the x axis, and only contains the top 30 important factors, as measured by average importance across all samples and models. Note that the four true underlying factors are bolded.



773 The linear models, unsurprisingly, struggled with factor significance analysis with respect to both increasing  
 774 cross sectional correlation and increasing non-linearities. This highlights the non-robustness and ineffectiveness  
 775 of using traditional linear regression as documented by the literature; linear models were consistently observed  
 776 to identify irrelevant regressors as important, especially as the degree of cross sectional correlation increased.  
 777 Considering that the graphs represent the averaged variable importance metrics over different simulation  
 778 realisations which each have random patterns of cross sectional correlation, this means that on a single simulation  
 779 realization, the performance of linear models is significantly worse.

780 The overall high performance of the elastic net models may be somewhat surprising given its relative simplicity  
 781 compared to other machine learning models. However, when recalling that elastic net models are the only machine  
 782 learning models which are specifically noted to perform well on datasets with high degrees of multicollinearity,  
 783 the result is perhaps less surprising. Indeed, random forest and neural network models are only noted to be  
 784 better at capturing non-linear relationships in independent and identically distributed data, a property which we  
 785 observe on the non-linear specification, and only for the random forests.

786 Of particular note are the instability of the machine learning models' hyperparameters across different training  
 787 samples. For the elastic nets, the optimal value for  $\alpha$  is generally 1 (corresponding to LASSO and thus a sparse  
 788 representation), but it was not uncommon to observe  $\alpha$  values swinging between values close to 0 (corresponding  
 789 to ridge regression, and thus a dense representation) to 1 as the training sample moved forwards in time. As  
 790 the penalized linear models consistently performed the best and still remained able to correctly identify the  
 791 true covariates this is not a large issue, but it should be noted that this can lead to interpretation issues. For the  
 792 random forests, it was similarly observed that the optimal value for  $mtry$  (the number of variables subsetting)  
 793 and  $nodesize$  was highly non-robust. Again, given that the final prediction performance was consistent this is  
 794 not a large issue, but can lead to some interpretation issues.

## 795 A Additional details: Empirical analysis

796 We conduct an empirical study as a final way to corroborate the findings of the properties of machine learning  
 797 models which we observed in the simulation study. Though our simulation study was aimed at capturing the  
 798 main features of observed data, the underlying data generating process for empirical returns is unknown. This  
 799 study thus acts as a robustness check as to how machine learning performs on real world data, which can be  
 800 significantly more complex and noisy than simulated contexts. Our two studies together can be thought of a  
 801 repeated sampling exercise in exploring how machine learning methods perform on datasets which feature the  
 802 "stylized facts" of empirical returns. This empirical study also acts as a final validation against what has been  
 803 reported in the literature.

804 Importantly, we find that our findings from the simulation study are largely corroborated for empirical returns  
 805 data.

### 806 A.1 Data

807 We begin by obtaining monthly individual price data from CRSP for all firms listed in the NYSE, AMEX and  
 808 NASDAQ, starting from 1957 (starting date of the S&P 500) and ending in December 2016, totalling 60 years.  
 809 To build individual factors, we construct a factor set based on the cross section of returns literature. This data  
 810 was sourced from and is the same data used in [?]. Like our initial returns sample, it begins in March 1957  
 811 and ends in December 2016, totalling 60 years. It contains 94 stock level characteristics: 61 updated annually,  
 812 13 updated quarterly and 20 updated monthly, in addition to 74 industry dummies corresponding the the first  
 813 two digits of the Standard Industrial Classification (SIC) codes. It is noted that this dataset so far contains all  
 814 securities traded, including those with a CRSP share code other than 10 or 11 and thus includes instruments such  
 815 as REITs and mutual funds, and those with a share price of less than \$5.

816 We detail our cleaning procedure of this dataset. To reduce the size of the dataset and increase feasibility, the  
 817 dataset was filtered such that only stocks traded primarily on NASDAQ were included (using the PRIMEXCH  
 818 variable from WRDS). Then, penny stocks (also referred to as microcaps in the literature) with a stock price  
 819 of less than \$5 were filtered out, as is commonly done in the literature to reduce variability. Stocks without a  
 820 share code of 10 or 11 (referring to equities) were filtered out, so that securities that are not equities were not  
 821 included (such as REITs and trust funds). The dataset is provided in a monthly format, which means that many  
 822 of the factors which are updated only quarterly or annually have very low levels of variability, which can lead to  
 823 misleading results in the model fitting process. To achieve a balance between having a dataset with enough data  
 824 points and variability among factors, the dataset was converted to a quarterly format. Quarterly returns were  
 825 then constructed using the PRC variable according to actual returns (ie not logged differences):

$$RET_t = \frac{PRC_t - PRC_{t-1}}{PRC_{t-1}} \quad (32)$$

Table 3: Macroeconomic Factors, ([?])

No.	Acronym	Macroeconomic Factor
1	macro_dp	Dividend Price Ratio
2	macro_ep	Earnings Price Ratio
3	macro_bm	Book to Market Ratio
4	macro_ntis	Net Equity Expansion
5	macro_tbl	Treasury Bill Rate
6	macro_tms	Term Spread
7	macro_dfy	Default Spread
8	macro_svar	Stock Variance

We allow all stocks which have a quarterly return to enter the dataset, even if they disappear from the dataset for certain periods, as opposed to only keeping stocks which appear continuously throughout the entire period. This was primarily done to reduce survivorship bias in the dataset, which can be very prevalent in financial data, and also allows for stocks which were unlisted and relisted again to feature in the dataset. This has the obvious drawback of introducing some bias in the dataset, as attrition in the dataset is likely to be non-random and correlated with the stocks' returns.

The sic2 variable, corresponding to the stocks' Standard Industrial Classification (SIC) codes was also dropped. The SIC code system suffers from inconsistent logic in classifying companies, and as a system built for pre-1970s traditional industries has been slow in recognizing new and emerging industries. Indeed, WRDS explicitly cautions the use of SIC codes beyond the use of rough grouping of industries, warning that SIC codes are not strictly enforced by government agencies for accuracy, in addition to most large companies belonging to multiple SIC codes over time. Because of this latter point in particular, there can be inconsistencies on the correct SIC code for the same company depending on the data source. Dropping the sic2 variable also reduced the dimensionality of the dataset by 74 columns, significant increasing computational feasibility.

There existed a significant amount of missing data in the dataset. The dataset's columns were first examined, and any characteristics that had over 20% of their data were removed. However, as the amount of missing data increases dramatically going further back in time, a balance between using more periods at the cost of removing more characteristics versus using less periods but keeping more characteristics was needed. 1993 Q3 was determined to be a reasonable time frame to begin the dataset, as there was a noticeable increase in data availability and quality after this time. Missing characteristics were then imputed using their cross sectional medians for each year.

We then follow [?] and construct eight macroeconomic factors following the variable definitions in [?]: dividend-price ratio (dp), earnings-price ratio (ep), book-to-market ratio (bm), net equity expansion (ntis), Treasury-bill rate (tbl), term spread (tms), default spread (dfy) and stock variance (svar). These factors were lagged by one period so as to be used to predict one period ahead quarterly returns. The treasury bill rate was also used from this source to proxy for the risk free rate in order to construct excess quarterly returns.

The two sets of factors were then combined to form a baseline set of covariates, which we define throughout all methods and analysis as:

$$z_{i,t} = (1, x_t)' \otimes c_{i,t} \quad (33)$$

where  $c_{i,t}$  is a  $P_c$  matrix of characteristics for each stock  $i$ , and  $(1, x_t)'$  is a  $P_x \times 1$  vector of macroeconomic predictors, , and  $\otimes$  represents the Kronecker product.  $z_{i,t}$  is therefore a  $P_x P_c$  vector of features for predicting individual stock returns and includes interactions between stock level characteristics and macroeconomic variables. The total number of covariates in this baseline set is  $61 \times (8 + 1) = 549$ <sup>11</sup>.


The dataset was not normalized for all methods, as only penalized regression and neural networks are sensitive to normalization. For these two methods, the dataset was normalized such that each predictor column had 0 mean and 1 variance.

The final dataset spanned from 1993 Q3 to 2016 Q4 with 202, 066 individual observations.

We mimic the procedure used in the simulation study. For the sample splitting procedure, the dataset was split such that the training and validation sets were split such that the training set was approximately 1.5 times the length of the validation set, in order to predict a test set that is one year in length.

<sup>11</sup> As the individual and macroeconomic factors can have similar names, individual and macroeconomic factors were prefixed with ind\_ and macro\_ respectively.

Figure 18: Empirical Data Sample Splitting Procedure



## A.2 Empirical Data Results

In general, the empirical results are in remarkable agreement with the those obtained in the simulation study. In particular, we similarly see that the penalized linear models generally performing the best, with the random forest models offering slightly worse performance, occasionally outperforming penalized linear models. Machine learning models fitted with respect to median quantile loss were similarly observed to typically offer improvements across all machine learning models across all loss metrics.

### A.2.1 Prediction Accuracy

In terms of prediction accuracy, we can see that in general the results of the simulation study were repeated: the elastic net models perform the best, followed by the random forests, then the neural networks, and finally the linear models. We note that the differences between each model using the MSE and MAE loss metrics are much more pronounced on empirical data. Even so, the predictive performance between the elastic net models and the quantile random forests is not particularly large, and we observe the quantile random forests outperforming the elastic nets in the first data sample. We similarly observe some evidence that machine learning models perform better when fitted with respect to quantile loss instead of MSE. Most notably, we start to see the neural network models performing poorly on the empirical data, a direct contradiction to what has been reported in the literature.

Figure 19: Empirical Test MSE

Figure 20: Empirical Test MAE

Figure 21: Empirical Test Predictive R-Squared

Focusing on the neural networks specifically, their non-robustness is amplified on the empirical dataset, with some neural networks in some samples even performing worse than linear models. This was observed to be somewhat more common on neural networks fitted with respect to MSE, suggesting that they are indeed very sensitive to outliers in training data. We similarly observe some evidence that deeper neural networks perform better, though this result is less apparent due to the lower robustness on empirical data (see ?? in Appendix for results).

Interestingly, we do not observe worsening performance as the training sample increases as we did in the simulation study. This suggests that the simulation design may have been too volatile when compared to the specific empirical time periods examined.

### A.2.2 Factor Importance

As the data generating process for empirical returns is unknown, the variable importance results cannot be directly compared with those of the simulation study. Even so, we see similar results: the elastic net and random forest models tend to agree on the same subset of predictors, but the random forest struggles to discern between highly correlated regressors. Similar to the prediction performance results, neural networks perform poorly.

Figure 22: Empirical Individual Factor Importance, averaged across all training samples

Figure 23: Empirical Macroeconomic Factor Importance, faceted by training sample

If we focus on the two top performing models of elastic net and random forest, we see that they consistently pick out the 1 month and 6 month momentum factors out of the individual characteristics as important, and the

896 book-to-market and default yield spread factors out of the macroeconomic factors are important. In general  
 897 though, the variable importance metrics are less consistent for the random forests, and it should be noted in  
 898 particular that the random forest tends to determine factors highly correlated with momentum as important,  
 899 such as change in moment, dollar trading volume and return volatility. Looking at the macroeconomic factors,  
 900 penalized linear models tend to identify the average book to market ratio and the default spread as the most  
 901 important macroeconomic factors. On the macroeconomic factor set, the random forests were inconsistent with  
 902 the elastic nets, and tended to assign very similar variable importance metrics to most macroeconomic factors.

903 The neural networks tended to believe that the market value factor was the most important among the individual  
 904 factors, a result not repeated by any of the other models considered. Within the macroeconomic factors, the  
 905 neural networks identified the dividend-price ratio and earnings-price ratio as the most important among the  
 906 macroeconomic factors, though these results were highly non-robust across different architectures, loss functions  
 907 and training samples.

908 Interestingly, we find that the linear models assign the controversial dividend price ratio macroeconomic factor as  
 909 highly important, a result mirrored only with the neural networks. Their variable importance for individual factors  
 910 across different training samples is highly non-robust, with the important variables almost completely changing  
 911 year to year. The linear models consistently identified the controversial dividend-price ratio as important, a result  
 912 that was somewhat consistent with the neural networks.

913 The overall results again contradict the results of [?], who conclude that all of the machine methods agree on the  
 914 same subset of important factors. Indeed, we only see consistency in variable importance between the elastic nets  
 915 and random forests on the individual factors only - all other variable importance metrics were either inconsistent  
 916 between different models, or non-robust.

917 All models considered typically preferred sparse parameterizations. That is, most if not all of the individual  
 918 factors had little to no importance across all models.<sup>12</sup>

---

<sup>12</sup>Note that because the variable importance here was not evaluated explicitly for each pairwise interaction term, some of the individual factors appear as slightly important. This is because setting an individual factor to zero also sets some of the macroeconomic pairwise terms to zero, increasing its apparent importance.