

Evaluation of Machine Learning in Empirical Asset Pricing

Ze Yu Zhong

Thursday 17th October, 2019

Abstract

We explore the performance of some common machine learning algorithms in empirical asset pricing with regards to prediction performance and causal inference through a simulation study and empirical study, with special focus on how well they deal with the problems in financial returns data. We evaluate prediction performance according to MSE, MAE and OOS R-Squared, and causal inference with a variable importance metric. Across both the simulation and empirical, we find consistent results that contradict those documented in the literature, and conclude that elastic net models perform the best both in terms of prediction performance and causal ability, followed by random forest models. We take particular note that neural networks fail to outperform other machine learning models, and highlight their non-robustness. Almost all machine learning models considered outperformed traditional regression by a large margin, and we show that machine learning methods offers some potential tools for the field of empirical finance.

Contents

1	Introduction	4
1.1	Topic	4
1.2	Background Literature and Motivations	4
2	Model Specification and Methodology	8
2.1	Model Specification	8
2.2	Methodology	8
2.2.1	Sample Splitting	9
2.2.2	Loss Function	10
2.2.3	Models	10
2.2.4	Linear Models	10
2.2.5	Penalized Linear Models	11
2.2.6	Classification and Regression Trees	12
2.2.7	Random Forests	12
2.2.8	Neural Networks	13
2.3	Model Evaluation	16
3	Simulation Study	18
3.1	Simulation Design	18
3.1.1	Overall Design	18
3.1.2	Simulating Characteristics	19
3.1.3	Simulating Macroeconomic Series	20
3.1.4	Simulating Return Series	20
3.1.5	Sample Splitting	20
3.2	Simulation Study Results	21
4	Empirical Study	29
4.1	Data	29
4.2	Empirical Data Results	32
5	Conclusion	37
6	Appendix	38
6.1	Data	38
6.2	Additional Results	43
6.2.1	Simulation Study	43
6.2.2	Empirical Study	50
6.3	Computational Details	54
6.3.1	Linear Models	54
6.3.2	Penalized Linear	54

6.3.3	Classification and Regression Trees	55
6.3.4	Random Forest	55
6.3.5	Neural Networks	56
6.3.6	Tuning of Simulated Datasets	57

1. Introduction

1.1. Topic

This thesis aims to evaluate the application of machine learning algorithms in empirical asset pricing. While there has been significant recent interest in applying machine learning to the problem of predicting asset returns, there is little literature that focuses on how well these algorithms are at capturing true underlying variables in determining stock returns. 12 different simulated datasets ranging from linear to highly non-linear data generating processes incorporating observed phenomena of cross sectional correlation, persistence, and stochastic volatility, in addition to real world data will be used to assess the performance of linear models, elastic net models, random forests and neural networks. Model performance will be assessed according to their predictive performance, measured by Mean Absolute Error, Mean Squared Error, and Predictive R^2 , and their causal inference ability, measured by a variable importance metric.

1.2. Background Literature and Motivations

Empirical finance can be broadly classified into two main goals: improving the prediction of future excess returns, and understanding the underlying causal relationship in regressors which drive excess returns. These two goals are noted to be very challenging for traditional regression methods, and this is largely due to the challenges present in the regressors of financial data, also known as “factors.”

We define “factors” with the more contemporary definition as suggested by [Harvey et al. \(2016\)](#): a collection of regressors to be used in pricing excess returns that can be used to proxy for unknown underlying risk factors due to their correlation with cross sectional returns.¹ [Harvey et al. \(2016\)](#) further group factors into the two broad categories of “common” and individual firm “characteristics.” “Common” factors under this definition can be viewed as proxies for sources of risk constant across all firms, such as macroeconomic variables. Examples of this include the Fama-French factors, which are portfolios of assets sorted by certain characteristics to proxy for them.

Because of this less strict definition factors are often noted for exhibiting properties which make them highly unsuitable for inclusion in traditional regression, ([Harvey et al. \(2016\)](#)). Factors have been empirically observed to have high persistence, high cross sectional correlation (multicollinearity), non-stationarity or have very low time series variation due to being pre-known or suffering from a lower sampling frequency. The consequences of including such regressors into traditional regression models are well documented. The resulting coefficient estimates and t-statistics can be biased, directly leading to incorrect statistical inference. Additionally, the variances for the the coefficient estimates can be very high, resulting in imprecise coefficient estimates. This will result in any out of sample predictions, which are a direct function of these

¹Most notably, their definition rejects the more strict view that risk factors should be variables that have unpredictable variation through time, and that they should be able explain cross sectional return patterns.

imprecise coefficients, also imprecise, particularly when the multicollinearity between variables changes over time.

Several real world examples of this have been documented in the literature, and perhaps the most infamous examples are the dividend ratio factors. Dividend ratio factors were initially very popular in the 1990s due to its inclusion in regression improving in sample performance. However, later literature such as [Goetzmann and Jorion \(1993\)](#) and [Ang and Bekaert \(2006\)](#) note the persistence present in dividend ratio factors. This means that movements in dividend ratios are dominated by movements in price and therefore dividend ratios are correlated with lagged dependent variables on the right hand side of the regression equation. This violates the assumptions of exogeneous regressors (independent from the error term) required for traditional regression models (ordinary least squares) to be unbiased, resulting in t statistics which are biased upwards and increase with time horizon due to autocorrelated errors. Importantly, [Goetzmann and Jorion \(1993\)](#) show that corrections to t statistics using the Generalized Method of Moments and Newey-West standard errors also appear to be biased upwards, making them unreliable.

[Goyal and Welch \(2003\)](#) provide a more comprehensive study on the performance of lagged dividend price ratios, with a specific focus on out of sample predictive performance both in terms of R^2 and forecast errors. They conclude that while models incorporating dividend related factors were able to achieve higher in sample performance prior to 1990 than the historical mean, they could not have outperformed the historical mean *out of sample*. [Goyal and Welch \(2003\)](#) attribute this to the increasing persistence and non-stationarity of dividend ratios, noting that they have become like random walks as of 2001. This mirrors the sentiment of ([Lettau and Ludvigson \(2001\)](#), [Schwert \(2003\)](#) and others) who conclude that models incorporating dividend ratios seemed to break down in the 2000s due to a changing economic environment, despite having performed well in the 1990s.

Despite the controversy, the prevailing tone within the literature was that various factors such as dividend ratios, earnings price ratio, interest and inflation and other financial indicators were able to predict excess returns, with [Lettau and Ludvigson \(2001\)](#) remarking that this was now “widely accepted.” However, [Welch and Goyal \(2008\)](#) extend upon the work of [Goyal and Welch \(2003\)](#) by including a more comprehensive set of variables and time horizons. They conclude that not a single variable had any statistical forecasting power. Crucially, they demonstrate the non-robustness of models incorporating these factors by showing that the significance values of some factors change with the choice of sample periods.

Despite this, the literature has continued to produce more factors: quantitative trading firms were using 81 factor models as the norm by 2014 ([Hsu and Kalesnik, 2014](#)), and [Harvey and Liu \(2019\)](#) currently document well over 600 different factors suggested in the literature.

The dramatic increase in the number of factors alone poses challenges to traditional statistical techniques. [Harvey et al. \(2016\)](#) detail the false discovery problem when the number of potential factors is extremely high. The significance of a factor in a traditional regression

setting is determined by a single hypothesis test, which by construction carries a level of significance α controlling the type I error rate: the probability of rejecting a "null" hypothesis that a factor is not important, and concluding that it is significant when the true factor is indeed insignificant. When the number of potential factors is large, it is much more likely that a factor will be concluded as significant by pure chance. For example, a factor model consisting of 600 factors would find around 30 factors significant by chance at the 5% significance level. [Harvey et al. \(2016\)](#) produce a multiple testing framework to mitigate this, and conclude that many of the historically discovered would have been deemed significant by chance.

Furthermore, [Feng et al. \(2019\)](#) note that the number of potential factors discovered in the literature has increased to the same scale as, if not greater, than the number of stocks considered in a typical portfolio, or the time horizon, producing highly inefficient covariances in a standard cross sectional regression. Moreover, when the number of factors exceeds the sample size, traditional cross sectional regressions become infeasible and do not produce solutions altogether.

It does not help that many factors are cross sectionally correlated, meaning that factors which are discovered to be significant may simply be so because they are correlated with a true, underlying factor and do not provide independent information themselves, a concern which [Cochrane \(2011\)](#) calls the multidimensional challenge. [Freyberger et al. \(2017\)](#) notes that this is especially challenging for traditional regression models, which make strong functional form assumptions and are sensitive to outliers.

More recently, machine learning methods have been used within the literature and appear to be well suited for use in empirical financial data, particularly as a means to manage the recent factor explosion within the literature. Machine learning requires less strict assumptions on the functional forms of returns, offers explicit "regularization" methods aimed to guard against over-fitting, and methods which emphasize variable selection and dimensional reduction techniques. All of these properties make machine learning potentially well suited to producing more robust models that outperform traditional regression.

The financial literature already has some applications of machine learning methods to assist with both returns prediction and causal inference. For example, [Kozak et al. \(2017\)](#), [Rapach and Zhou \(2013\)](#) and [Freyberger et al. \(2017\)](#) apply shrinkage and selection methods from machine learning to assist with the problem of factor selection.

Most recently, there is literature that suggests machine learning methods can result in vastly better prediction performance ([Gu et al. \(2018\)](#), [Hsu and Kalesnik \(2014\)](#) and [Feng et al. \(2018\)](#)) in terms of out of sample predictive R^2 . [Gu et al. \(2018\)](#) attribute this to machine learning's ability to evaluate and consider non-linear complexities among factors that cannot be feasibly achieved using traditional techniques.

However, there is little work done on how machine learning actually recognises and deals with the challenges of returns prediction documented in the literature. Prior work has been done by [Gu et al. \(2018\)](#), however; only basic simulation designs which were not representative of real financial data were considered. In particular, the performance of machine learning in

empirical financial contexts with specific regards to characteristics such as stochastic and time varying volatility, long term dependence and cross sectionally correlated regressors has not been adequately explored.

This particularly problematic, given that the machine learning literature has documented that machine learning methods struggle with a lot of the characteristics observed within financial data. Several methods such as penalized regression and tree based models assume some form of conditional independence between each observation, a feature that is expressly violated by state dependent data such as empirical returns. Neural networks in particular, have been extensively noted to be extremely difficult to train and struggle with capturing long term memory effects in state dependent data, (Bengio et al. (1994)). Furthermore, many of the more complex machine learning models are still noted to be prone to overfitting, and as such mostly see applications in contexts with extremely large amounts of data, which may not be available with empirical returns data due to the limited time frame in which relevant data is available.

Furthermore, there have been some methodology issues observed within the sparse literature on the topic. Feng et al. (2018) use cross validation as part of their model building procedure, destroying the temporal aspect of returns data, in addition to only using a handful of factors. Gu et al. (2018) produce models using a training sample which ends in the 1970s to ultimately produce forecasts for the most recent 30 years.

Given the sparse literature on the topic, the methodology issues with aforementioned literature, and the noted difficulties with applying machine learning to data such as empirical returns, it stands to reason then, that more research should be done into the application of machine learning methods to empirical data. Specifically, the prediction and causal inference ability of machine learning regards to datasets exhibiting characteristics such as multicollinearity, regressor persistence, stochastic volatility, high dimensionality low signal to noise ratio and potentially low number of observations as is the case with financial data should be explored.

This paper is the first to document how machine learning algorithms perform in the context of empirical returns data. We explore this through a simulation study and an empirical study. The simulation study explicitly explores how different aspects of financial data such as persistence in regressors, cross sectional correlation and different complexities of data generating process can affect machine learning's ability in terms of both prediction performance, and causal inference. These results are then validated against an empirical dataset, filtered to be more recent and representative of recent returns. These two studies together offer a clearer glimpse into how machine learning algorithms deal with the challenges present in financial data, and crucially, we find that the results from both studies mostly consistent, suggesting a robust set of results.

2. Model Specification and Methodology

2.1. Model Specification

All asset excess monthly returns denoted as $r_{i,t+1}$ are modelled as an additive prediction error model conditional on the true and unobservable information set available to market participants up to and including time t , \mathcal{F}_t):

$$r_{i,t+1} = E(r_{i,t+1}|\mathcal{F}_t) + \epsilon_{i,t+1} \quad (1)$$

where

$$E(r_{i,t+1}|\mathcal{F}_t) = g^*(z_{i,t}) \quad (2)$$

with stocks indexed as $i = 1, \dots, N$ and time domain by $t = 1, \dots, T$. $g^*(z_{i,t})$ represents the model approximation using the P dimensional predictor set $z_{i,t}$. We allow $g^*(z_{i,t})$ to be a flexible function of the predictor set $z_{i,t}$, and most notably, not depend on i or t directly. This means that we do not re-estimate a model for each time period, or independently estimate a model for each stock. Note that $g^*(z_{i,t})$ only contains information in time t for individual stock i , meaning that while the model and its parameters will be estimated using \mathcal{F}_t for stock i , predictions for $r_{i,t+1}$ will only use information at time t as an input, analogous to using variables lagged by one period.

All machine learning methods are designed to approximate the empirical model $E_t(r_{i,t+1}) = g^*(z_{i,t})$ defined in equation (2). We define the baseline set of stock-level covariates $z_{i,t}$ as:

$$z_{i,t} = x_t \otimes c_{i,t} \quad (3)$$

where $c_{i,t}$ is a $P_c \times 1$ matrix of characteristics for each stock i , and x_t is a $P_x \times 1$ vector of macroeconomic predictors (and are this common to all stocks, including a constant). Thus $z_{i,t}$ is a $P \times 1$ vector of features for predicting individual stock returns ($P = P_c P_x$) and includes interactions between individual characteristics and macroeconomic characteristics.

2.2. Methodology

The definition of machine learning can be vague and is often context specific; [Hastie et al. \(2009\)](#) in *An Introduction to Statistical Learning* describes statistical (machine) learning as a vast set of tools for understanding data, and *supervised* learning specifically as the process of building a statistical model for the prediction or estimation of an output based on input(s). In the context of asset pricing, we use the term to refer to a diverse collection of:

1. high-dimensional models for statistical prediction,
2. the “regularization” methods for model selection and mitigation of over-fitting input data,

3. and the efficient systematic methods for searching potential model specifications.

The high dimensional and hence flexible nature of machine learning brings more hope to approximating unknown and likely complex data generating processes that underlie excess returns. The flexibility however, comes at a cost of potentially over-fitting in sample data (referred to as training data in the machine learning literature), generalizing poorly and producing poor forecasts. The regularization aspect of machine learning explicitly guards against over-fitting problems and emphasizes out of sample performance. The most explicit example of regularization is the splitting of the dataset into an explicit "training" set used for model fitting, and a "test" set withheld and used solely for evaluating out of sample performance. Finally, machine learning offers tools which are designed to produce an optimal model specification from all possible models with manageable computational cost, all in a systematically consistent way.

All machine learning methodologies can be broadly broken down into three components: sample splitting, the choice of loss function, and finally the specification and fitting of models.

2.2.1. Sample Splitting

Imperative to any machine learning technique is the establishment of how the dataset is to be split into training, validation and test sets. The training set is used to initially build the model and provide initial estimates of parameters, whereas the validation set is used to tune model parameters to optimise out of sample performance, thus preventing overfitting. The validation set acts as a simulation of out of sample testing, whereas the test set is used only for evaluation, and is thus truly out of sample.

There are three main approaches to splitting temporal data (such as financial data).

The first is to decide arbitrarily on a single training, validation and test set. This method is straightforward and the least computationally intensive, but is limited and inflexible in evaluating how models perform when more recent data is provided for training.

The second method is a "rolling window" method, where a fixed size or "window" for the training and validation set is first chosen. This window then incrementally move forwards in time to include more recent data, with a set of forecasts for the test sets made for all possible windows.

The third is a "recursive" method, which is the same as the rolling window method, but different in that the training set always contains previous data, with only the validation set staying fixed in size and "rolling" forwards. Hence, it is also referred to as a "growing window."

Both the rolling window and recursive schemes are very computationally intensive. Therefore, a hybrid of the rolling and recursive schemes was considered: the training set is increased by one year with each refit, the validation set remains one year in length but moves forward by one year, and forecasts are made using that model for the subsequent year. The "traditional" cross validation method of randomly sampling to determine a train and validation set was not done to maintain the temporal ordering of the data.

2.2.2. Loss Function

The choice of the loss function used in models is imperative to machine learning. The loss functions considered are Mean Absolute Error (MAE) and Mean Squared Error².

The mean absolute error (MAE) is simply the average magnitude of errors. Because of this, it places equal weighting to all magnitudes of errors and is more robust to outliers.

$$\text{MAE} = \frac{1}{n} \sum_{j=i}^n |y_j - \hat{y}_j| \quad (4)$$

It should be noted that minimizing the MAE criterion is equivalent to minimizing 0.5 quantile loss.

The mean squared error (MSE) and root mean squared error (RMSE) are quadratic scoring methods. This means that they place higher weight on large errors. Models that minimize this metric are therefore more sensitive to outliers.

$$\text{MSE} = \frac{1}{n} \sum_{j=i}^n (y_j - \hat{y}_j)^2 \quad (5)$$

2.2.3. Models

We focus our methodology on four common machine learning models: linear models, penalized linear models, random forests, and neural networks. Each model will be presented and explained so that a reader without any machine learning background can understand the basic idea behind each model.

2.2.4. Linear Models

The least complex model considered is the simple linear regression model. The simple linear model assumes that the underlying conditional expectation $g^*(z_{i,t})$ can be modelled as a linear function of the predictors and the parameter vector θ :

$$g(z_{i,t}; \theta) = z'_{i,t} \theta \quad (6)$$

Computing this model with respect to minimizing the mean squared error yields the pooled ordinary least squares estimator (OLS), while minimizing the mean absolute error corresponds to quantile regression predicting the 0.5th quantile, also known as the Least Absolute Deviation (LAD) estimator.

The OLS estimator is known to be consistent when the regressors are exogeneous, and is optimal in the class of linear unbiased estimators when the errors are homoscedastic and serially uncorrelated. This unbiasedness consequently means that linear models have a large amount of variance due to the bias-variance trade-off, which can lead to poorer predictive ability as

²Also referred to as l_1 and l_2 loss respectively within the machine learning literature

the estimator's variance will be very high and therefore inefficient. Additionally, OLS struggles when multicollinearity (also known as cross sectional correlation in empirical finance settings) is present, as the coefficients of different variables can behave erratically to small changes in the data or model. This impreciseness in the coefficient estimates further leads to poor predictive performance. The LAD estimator attempts to improve upon the OLS estimator through the use of loss function that is more robust to outliers. The mean absolute error weights all residuals equally, compared to mean squared error which through considering squared residuals gives more emphasis to large residuals.

Linear models can capture non-linearities only if the predictor set $z_{i,t}^*$ contains specified non-linear transformations or interaction terms. Despite being expected to perform poorly, linear models were implemented as a “control.”

2.2.5. Penalized Linear Models

Penalized linear models attempt to improve upon linear models through regularization, and are particularly well suited to addressing problems of multicollinearity, which can commonly occur in models with a large number of regressors. They achieve this by lowering the variance of their estimates (thus increasing efficiency) in exchange for introducing a tolerable amount of bias. This can be viewed as the model mechanically decreasing its in sample performance, in hopes that the model will overfit less, particularly to noise in the dataset, thus correctly preserving its fit to true underlying regressors.

Penalized linear models have the same underlying statistical model as simple linear models, but differ in their addition of a new penalty term in the loss function:

$$\mathcal{L}(\theta; \cdot) = \underbrace{\mathcal{L}(\theta)}_{\text{Loss Function}} + \underbrace{\phi(\theta; \cdot)}_{\text{Penalty Term}} \quad (7)$$

Several choices exist for the choice of penalty function $\phi(\theta; \cdot)$. We restrict our scope to the popular “elastic net” penalty (Zou and Hastie, 2005):

$$\phi(\theta; \lambda, \rho) = \lambda(1 - \rho) \sum_{j=1}^P |\theta_j| + \frac{1}{2} \lambda \rho \sum_{j=1}^P \theta_j^2 \quad (8)$$

The elastic net has two hyperparameters: λ , which controls the overall magnitude of the loss, and ρ , which controls the shape of the penalization.

The $\rho = 1$ case corresponds to ridge regression proposed by Hoerl and Kennard (1970), which uses l_2 penalty that shrinks all coefficients closer to 0, but not to 0. Ridge regression is therefore a shrinkage method which prevents coefficients from becoming too large and overpowering.

The $\rho = 0$ case corresponds to the popular LASSO and uses absolute (l_1) parameter penalization proposed by Tibshirani (1996), which geometrically allows the coefficients to be shrunk to 0. This allows it to impose sparsity, and can be thought of as a variable selection tool.

For $0 < \rho < 1$, the elastic net aims to produce parsimonious models through both shrinkage

and selection by combining the properties of LASSO and ridge regression.

The hyperparameters λ and ρ are both tuned using the validation sample (see 6.2.2).

2.2.6. Classification and Regression Trees

Classification and regression trees are fully non-parametric models that can capture complex multi-way interactions. A tree "grows" in a series of iterations. With each iteration, a split ("branch") is made along one predictor such that it is the best split available at that stage with respect to minimizing the loss function. These steps are continued until each observation is its own node, or more commonly until the stopping criterion is met. The eventual model slices the predictor space into rectangular partitions, and predicts the unknown function $g^*(z_{i,t})$ with the average value of the outcome variable in each partition.

The prediction of a tree, \mathcal{T} , with K "leaves" (terminal nodes), and depth L is

$$g(z_{i,t}; \theta, K, L) = \sum_{k=1}^K \theta_k \mathbf{1}_{z_{i,t} \in C_k(L)} \quad (9)$$

where $C_k(L)$ is one of the K partitions in the model.

For this study, only recursive binary trees (the most common and easy to implement) are considered. Though trees were originally proposed and fit with respect to minimizing mean squared error, they can be grown with respect to a variety of loss functions, including mean absolute error, mean squared error, where the loss within each C partition is denoted by $H(\theta, C)$:

$$H(\theta, C) = \frac{1}{|C|} \sum_{z_{i,t} \in C} L(r_{i,t+1} - \theta) \quad (10)$$

where $|C|$ denotes the number of observations in set C (partition). Given C , it is clear that the optimal choice for minimising the loss function when it is mean squared error is simply $\theta = \frac{1}{|C|} \sum_{z_{i,t} \in C} r_{i,t+1}$ i.e. the average of the partition, and the median of the partition when the loss function is mean absolute error.

Trees, grown to a deep enough level, are highly unbiased and flexible, as each partition can potentially predict a single, or low number of observations. The trade-off is their high variance and instability. Thus, an ensemble method called "Random Forest" was proposed by Breiman (2001) to regularize trees by combining many different trees into a single prediction.

2.2.7. Random Forests

Random Forests are an extension of regression trees that attempt to address some of their problems, proposed by Breiman (2001). A random forest algorithm creates B different bootstrap samples from the training dataset, fits an overfit (and hence low bias) regression tree to each dataset using only a random subset m size from all available predictors (also known as dropout), and then averages their forecasts as the final output. The dropout procedure in particular ensures that trees will be unable to use the same predictors when considering splits, lowering

the correlation between each tree and thus further reducing the variance across the ensemble model. Thus, the ensemble consisting of only overfit (and hence low bias) which have low uncorrelation is a low bias, yet stable model. Specific details of the random forest algorithm are detailed in the appendix.

2.2.8. Neural Networks

Neural networks have theoretical underpinnings as “universal approximators” for any function, (Hornik et al. (1989)). They are arguably the most complex type of model available, able to capture several non-linear interactions through their many layers, hence its other name “deep learning.” On the flipside, their high flexibility often means that they are among the most parameterized and least interpretable models, earning them the reputation as a black box model.

Per this analysis, we focus on traditional “feed-forward” networks. The feed forward network consists of an “input layer” of scaled data inputs, one or more “hidden layers” which interact and non-linearly transform the inputs, and finally an output layer that aggregates the hidden layers and transform them a final time for the final output.

More specifically, a neural network consists of layers denoted by $l = 0, 1, \dots, L$, with $l = 0$ denoting the input layer and $l = L$ denoting the output layer. The input layer is defined by the scaled predictor set, $x^{(0)} = (1, z_1, \dots, z_N)'$. The model adds complexity through the use of one or more hidden layer, each containing $K^{(l)}$ “neurons”. Each neuron linearly aggregates the values of the previous layer, and applies some non-linear “activation function” which we denote as α to its aggregated signal before sending its output to the next layer. The output of neuron k in layer l is then $x_k^{(l)}$. Next, define the vector of outputs for this layer as $x^{(l)} = (1, x_1^{(l)}, \dots, x_{K^{(l)}}^{(l)})'$. The recursive output formula for the neural network at each neuron in layer $l > 0$ is then:

$$x_k^{(l)} = \alpha(x^{(l-1)'} \theta_k^{l-1}), \quad (11)$$

where $\alpha()$ represents the activation function for that layer with the final output ³

$$g(z; \theta) = x^{(L-1)'} \theta^{L-1} \quad (12)$$

The neural network’s weight and bias parameters for each layer are estimated by minimizing the loss function with respect to the parameters, i.e. by calculating the partial derivative with respect to a specific weight or bias element.

Due to the complexity and hence non-existent analytical form for this solution, this is typically found via backpropagation, an algorithm which exploits the chain rule of the partial derivative and iteratively finds a local optimum using a first order gradient based algorithm, also known as “gradient descent.” The gradient descent algorithm minimizes some function (such as the loss function in the context of machine learning) by iteratively moving in the direction

³Note that the specification of a constant “1” at the beginning of each layer is the same as specifying a bias term as is popular in other parametrizations.

of steepest descent, defined as the negative gradient. Formally, for a loss function $L(x)$ that is defined and has a gradient defined in the neighbourhood of the parameter set a , the updating algorithm is:

$$a_{n+1} = a_n - \gamma \Delta F(a_n) \quad (13)$$

where γ controls the size of each update. This γ parameter is known as the learning rate in neural network training, and controlling this is critical for good performance. As the loss functions of neural networks can be very complex with many local minima, the learning rate should be high enough such that the parameter updates are large enough to skip or jump over them. Too large of a learning rate however, and the neural may fail to converge to a solution at all. Due to computational limitations, we tune the learning rate manually, and employ the ADAM algorithm to apply allow different learning rates for each parameter.

Stochastic gradient descent is a variation of traditional gradient descent which assists with computational feasibility and producing solutions which generalize better, and hence better performance. Instead of optimizing the weight parameters with respect to the entire training sample, a small random subset is instead used, the size of which is called the batch size. This has the effect of potentially increasing computational feasibility due to smaller memory requirements. Less understood, but empirically evident is that smaller batch sizes tend to produce better solutions, (Keskar et al. (2016)). There is however, a trade-off: smaller batch sizes can result in substantial noise due to the randomness of subsampling, and the batch size should be large enough such that this noise does not overpower weight updates. It is therefore a critical hyperparameter to tune.

We also employ l_1 penalty to impose some sparse restrictions on the weight terms, aiding with the production of better generalization. This simply adds the absolute value of each weight parameter, multiplied by a scaling factor which is a hyperparameter to be tuned. It was observed that some l_1 penalty was necessary to ensure good convergence of the neural networks.

“Batch normalization” is a technique for addressing a phenomenon known as internal covariate shift, which occurs when the distributions of each layers’ inputs change as the parameters of the previous layer change (Ioffe and Szegedy, 2015). This occurs due to how neural networks repeatedly apply the activation function many times over different hidden layers. For activation functions such as tanh, this results in weight parameters getting pushed and thus “saturated” towards -1 and 1, leading to difficult and slow training. Batch normalization addresses this by normalizing (de-meaning and variance standardizing) the outputs of each layer, hence restoring the representative power of each neuron.

Model Architecture and Specification Neural networks with up to 5 hidden layers were considered, each named NNX where X represents the number of hidden layers. The number of neurons in each layer was chosen according to the geometric pyramid rule (Masters, 1993): NN1 has 32 neurons, NN2 has 32 and 16 neurons in the first and second hidden layers respectively,

NN3 has 32, 16, and 8 neurons, NN4 has 32, 16, 8, and 4 neurons, and NN5 has 32, 16, 8, 4, 2 neurons respectively. All units are fully connected; that is, each neurons receives input from all neurons the layer before it (see Figure 1). This mimics the methodology in [Gu et al. \(2018\)](#).

Several choices of activation functions exist in the literature. We use the hyperbolic tangent function:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \quad (14)$$

for all hidden layers due to its observed performance, and ability to retain weight terms and hence always receive weight updates during training. This point is particularly important, because the popular ReLU activation function (see [Lecun et al. \(2015\)](#) and [Ramachandran et al. \(2017\)](#), among others):

$$\text{ReLU}(x) = \max(0, x) \quad (15)$$

was considered, but observed to suffer tremendously from the "dying-ReLU problem". This is where ReLU neurons receive weight updates that fail to activate (output a 0), hence making it unable to receive further weight updates and learn. Networks trained using this activation function were observed to output the same value for the majority of, if not all inputs. Variants of the ReLU activation function aimed at addressing this issue such as the leaky ReLU, which allows a small, non-zero gradient for negative weights were explored, but still suffered from poor convergence.

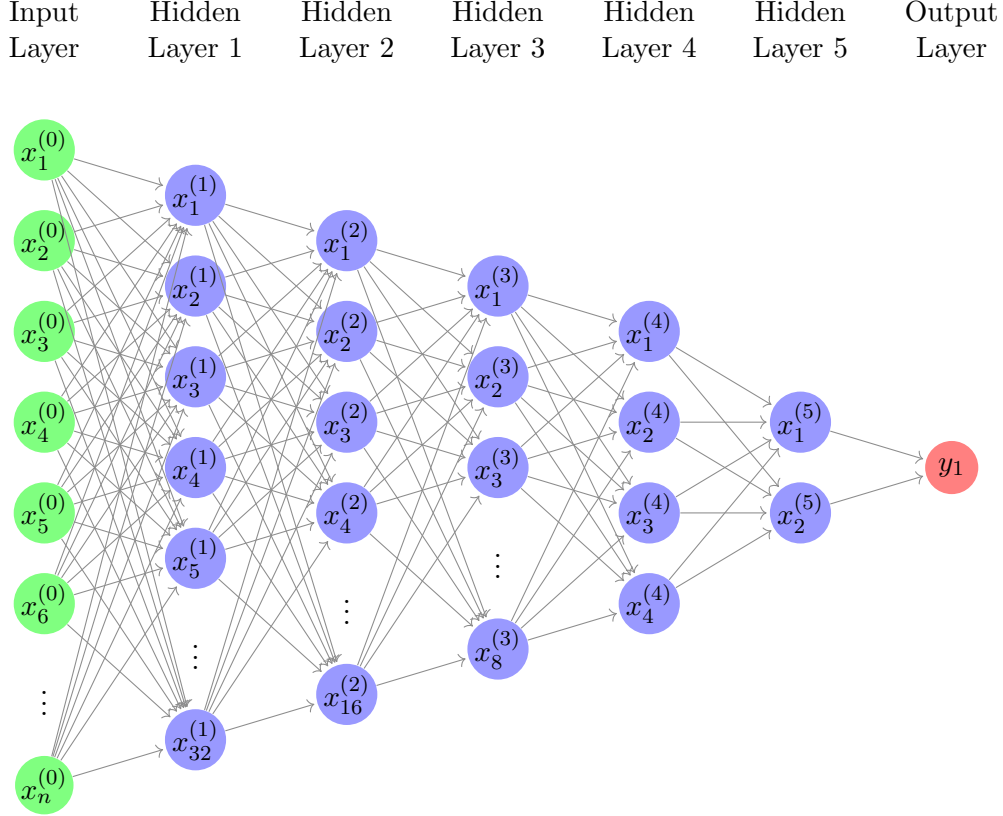


Figure 1. Neural Network 5 (most complex considered), without bias terms drawn

2.3. Model Evaluation

Loss Metrics Predictive performance for individual excess stock returns were assessed using popular loss metrics in the literature: Mean Absolute Error (MAE), Mean Squared Error (MSE) and an out-of-sample R-squared metric.

MSE is noted to be traditionally very popular within both the machine learning and econometrics literature. However, there is also some literature which argues against the use of MSE (and all square error measures) due to its inaccuracy and potential biasedness in time series settings, and advocates for the use of MAE instead. There is no consensus on an “ideal” loss metric, and for this reason we produce both of them. These loss metrics also have a one-to-one correspondence to the loss functions considered for fitting the models.

An out of sample R^2 metric was also reported, as is popular in the empirical finance literature. Due to the lack of consensus as to how this metric is defined, we provide our formulation:

$$\text{MAE} = \frac{1}{n} \sum_{j=i}^n |y_j - \hat{y}_j| \quad (16)$$

$$\text{MSE} = \frac{1}{n} \sum_{j=i}^n (y_j - \hat{y}_j)^2 \quad (17)$$

$$R_{OOS}^2 = 1 - \frac{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \hat{r}_{i,t+1})^2}{\sum_{(i,t) \in \mathcal{T}_3} (r_{i,t+1} - \bar{r}_{i,t+1})^2} \quad (18)$$

where \mathcal{T}_3 indicates that the fits are only assessed on the test subsample, which is never used for training or tuning.

Note that because R^2 measures were originally formulated for assessing in sample fit for linear models, the interpretation of this metric is less meaningful in the contexts of forecasting out of sample and non-linear models. This metric was nevertheless produced because of its property to exaggerate the differences in performance across different models.

Importantly, we see that the choice of loss metrics chosen does not affect the overall conclusion of which model performs the best.

Variable Importance To explore the causal inference capabilities of machine learning, we define a simple variable importance metric that is able to be consistently applied to all machine learning models. The importance of each predictor j is denoted as VI_j , and is defined as the reduction in predictive R-Squared from setting all values of predictor j to 0, while holding the remaining model estimates fixed, mirroring the procedure of [Gu et al. \(2018\)](#). These were then normalized to sum to 1 within each model as a way to assess the relative importance of each predictor for each model.

As VI_j can sometimes be negative, or in some cases be 0 across most of the factors, VI_j positively shifted by the magnitude of the smallest VI_j plus a minor offset o , then dividing all VI_j by the total to alleviate numerical issues⁴:

$$VI_{j,norm} = \frac{VI_j + \min(VI_j) + o}{\sum VI_j + \min(VI_j) + o} \quad ; \quad o = 10^{-100} \quad (19)$$

⁴This mechanism was chosen because the other popular normalization mechanism "softmax" was observed to be unable to preserve the distances between each original VI_j , making discernment between each VI_j difficult.

3. Simulation Study

3.1. Simulation Design

RS: motivate this section more clearly

We begin with the simulation study as a way to explore how machine learning performs with specific regards to the characteristics present in financial data. By specifying a simulation design with the desired characteristics, a controlled and well understood environment is available for explicitly testing the predictive performance and causal analysis abilities of the machine learning models considered. In particular, a sensible simulation design which incorporated the following characteristics of financial data was needed:

- Low signal to noise ratio
- Stochastic volatility in errors (including large random external shocks and volatility clustering)
- Persistence in regressors
- Cross sectional correlation (multicollinearity) in regressors

3.1.1. Overall Design

We first construct our simulation design by considering the design considered by [Gu et al. \(2018\)](#), which consists of three overall components:

- Simulate individual firm and macroeconomic factors
- Enter these factors into a latent, true data generating process
- Produce the overall returns process as this latent generative process plus an error process

However, [Gu et al. \(2018\)](#)'s specification has two main issues: the factors which enter the return equation are uncorrelated across each stock i , and the error process specified amounts to a white noise, constant volatility specification. As noted by [Harvey et al. \(2016\)](#) and many others, this is not what is observed in practice.

We therefore simulate a latent factor model with a stochastic volatility process for excess returns r_{t+1} , for $t = 1, \dots, T$:

$$r_{i,t+1} = g(z_{i,t}) + \beta_{i,t+1}v_{t+1} + e_{i,t+1}; \quad z_{i,t} = (1, x_t)' \otimes c_{i,t}, \quad \beta_{i,t} = (c_{i1,t}, c_{i2,t}, c_{i3,t}) \quad (20)$$

$$e_{i,t+1} = \sigma_{i,t+1}\varepsilon_{i,t+1}; \quad (21)$$

$$\log(\sigma_{i,t+1}^2) = \omega + \gamma \log(\sigma_t^2) + \sigma_u u; \quad u \sim N(0, 1) \quad (22)$$

Let v_{t+1} be a 3×1 vector of errors, and $w_{t+1} \sim N(0, 1)$ and $\varepsilon_{i,t+1} \sim N(0, 1)$ scalar error terms. The parameters of these are tuned such that the annualized volatility of each return series was approximately 22%, as is often observed empirically.

The matrix C_t is an $N \times P_c$ vector of latent factors, where the first three columns correspond to $\beta_{i,t}$, across the $1 \leq i \leq N$ dimensions, while the remaining $P_c - 3$ factors do not enter the return equation. The $P_x \times 1$ vector x_t is a 3×1 multivariate time series, and ε_{t+1} is a $N \times 1$ vector of idiosyncratic errors.

Note that we also reproduce [Gu et al. \(2018\)](#)'s error specification as a case where there is no stochastic volatility:

$$v_{t+1} \sim N(0, 0.05^2 \times I_3) \quad (23)$$

$$e_{i,t+1} \sim t_5(0, 0.05^2) \quad (24)$$

3.1.2. Simulating Characteristics

A simulation mechanism for C_t that gives some correlation across the factors and across time was used. We build in correlation across time among factors by drawing normal random numbers for each $1 \leq i \leq N$ and $1 \leq j \leq P_c$, according to

$$\bar{c}_{ij,t} = \rho_j \bar{c}_{ij,t-1} + \epsilon_{ij,t}; \quad \rho_j \sim \mathcal{U}\left(\frac{1}{2}, 1\right) \quad (25)$$

To build in cross sectional correlation, we define the positive-semidefinite matrix B :

$$B := \Lambda \Lambda' + \frac{1}{10} \mathbb{I}_n, \quad \Lambda_i = (\lambda_{i1}, \dots, \lambda_{i4}), \quad \lambda_{ik} \sim N(0, \lambda_{sd}), \quad k = 1, \dots, 4 \quad (26)$$

to serve as a variance covariance matrix with λ_{sd} controlling the density of the matrix, and hence degree of cross sectional correlation. λ_{sd} values of 0.01, 0.1 and 1 were used to explore increasing degrees of cross sectional correlation.

To build this into our $N \times P_c$ characteristics matrix \bar{C}_t , we simulate characteristics according to

$$\hat{C}_t = L \bar{C}_t; \quad B = LL' \quad (27)$$

where L represents the lower triangle matrix of B using the Cholesky decomposition.

Finally, the "observed" characteristics for each $1 \leq i \leq N$ and for $j = 1, \dots, P_c$ are constructed according to:

$$c_{ij,t} = \frac{2}{n+1} \text{rank}(\hat{c}_{ij,t}) - 1. \quad (28)$$

with the rank transformation normalizing all predictors to be within $[-1, 1]$.

3.1.3. Simulating Macroeconomic Series

For simulation of x_t , a 3×1 multivariate time series, we consider a Vector Autoregression (VAR) model ⁵:

$$x_t = Ax_{t-1} + u_t; \quad A = \begin{pmatrix} .95 & 0 & 0 \\ 0 & .95 & 0 \\ 0 & 0 & .95 \end{pmatrix} \quad u_t \sim N(\mu = (0, 0, 0)', \Sigma = I_3)$$

3.1.4. Simulating Return Series

We consider three different functions for $g(z_{i,t})$:

$$(1) \ g_1(z_{i,t}) = (c_{i1,t}, c_{i2,t}, c_{i3,t} \times x'_t[3,]) \theta_0 \quad (29)$$

$$(2) \ g_2(z_{i,t}) = (c_{i1,t}^2, c_{i1,t} \times c_{i2,t}, \text{sgn}(c_{i3,t} \times x'_t[3,])) \theta_0 \quad (30)$$

$$(3) \ g_3(z_{i,t}) = (1[c_{i3,t} > 0], c_{i2,t}^3, c_{i1,t} \times c_{i2,t} \times 1[c_{i3,t} > 0], \text{logit}(c_{i3,t})) \theta_0 \quad (31)$$

where $x'_t[3,]$ denotes the third element of the x'_t vector.

$g_1(z_{i,t})$ allows the characteristics to enter the return equation linearly, and $g_2(z_{i,t})$ allows the characteristics to enter the return equation interactively and non-linearly. The true underlying causal regressors for these specifications are $(c_{i1,t}, c_{i2,t}, c_{i3,t} \times x'_t[3,])$. These two specifications correspond to the simulation design used by [Gu et al. \(2018\)](#).

$g_3(z_{i,t})$ allows the characteristics to enter in a complex and non-linear fashion. The true underlying causal regressors for this specification are $(c_{i1,t}, c_{i2,t}, c_{i3,t})$.

It should be noted however, that because $g_2(z_{i,t})$ has a large part of its signal entering through a sgn function, this should make it the most difficult to estimate given the regressors and resulting returns process.

θ^0 was tuned such that the predictive R^2 was approximately 5%.

The simulation design results in $3 \times 3 = 12$ different simulated datasets, each with $N = 200$ stocks, $T = 180$ periods and $P_c = 100$ characteristics. Each design was simulated 10 times to assess the robustness of machine learning algorithms. The number of simulations was kept low for computational feasibility.

3.1.5. Sample Splitting

If viewed as monthly periods, $T = 180$ corresponds to 15 years. A data splitting scheme similar to the scheme to be used in the empirical data study was used: a training:validation length ratio of approximately 1.5 to begin, and a test set that is 1 year in length. We employ the hybrid

⁵Other more complex and interactive matrix specifications of A were briefly explored, but these did not appear to have a significant impact on results. More complex designs were observed to only affect the variable importance metrics, but to an insignificant degree




Figure 2. Sample Splitting Procedure

growing window approach as described earlier in section 2.2.1 (see Figure 2 for a graphical representation).

Other schemes in the forecasting literature such as using an “inner” rolling window validation loop to find the best hyperparameters on average, finally aggregating them in an “outer” loop for a more robust error were considered but not implemented for a variety of reasons. Firstly, many of the models were computationally too intensive for this to be feasible. More importantly, during the model fitting process it was observed that the optimal hyperparameters for the different rolling windows were highly unstable (see Appendix). Thus, this would have made the selection of the best hyperparameters on average across all windows significantly less meaningful.

3.2. Simulation Study Results

Overall, in the simulation study we observe that in general elastic nets are the best performing model, followed closely by random forests, then neural networks. All machine learning models were unaffected by cross sectional correlation in terms of prediction performance, and had better performance when fitted with respect to quantile loss, in stark contrast to linear models. The random forest only outperformed the elastic nets on highly non-linear specifications. The neural network models were not observed to outperform any of the machine learning models. We note that most of these results contradict sparse literature, and in particular, the results reported by Gu et al. (2018), even on their proposed simulation design.

Looking at the prediction performance of different models, we find that in general, penalized linear models performed the best, followed extremely closely by random forests and then neural networks, which all outperform the baseline linear models. Most importantly, we observe that for machine learning models, cross sectional correlation does not seem to affect prediction performance by much. This is in stark contrast to the linear models, whose prediction performance is severely affected by both non-linearities, and increasing cross sectional correlation. This result is consistent across all loss metrics, and is most obvious when looking at the out-of-sample R-squared metrics.

Machine learning models fitted with respect to minimizing MAE (quantile loss) generally perform better, even when evaluated against MSE loss metrics. This is not a surprising result, especially considering the stochastic error design which introduces significant shocks to the

returns process, leading to large outliers which the mean squared error metric is more sensitive to. Though the actual difference between the loss metrics between the penalized linear models, random forests and neural networks are very small, when considering the consistency of the results across numerous Monte Carlo simulations, the differences in prediction performance, though small, is robust and significant.

Across all specifications with a stochastic volatility component, we observe a decrease in prediction performance as the sample size increased according to the expanding window approach implemented. This is likely due a larger sample having a higher chance to experiences external shocks due to the stochastic volatility process, and thus a higher chance to experience large outliers in the training sample, leading to worse prediction performance.

Focusing on the neural networks, we clearly see that they do not outperform any of the other machine learning models, even when the underlying data generating process is non-linear. This directly contradicts the result which [Gu et al. \(2018\)](#) find, even when considering the design with no cross sectional correlation and stochastic volatility (top row in graphs), which corresponds to their exact specification but with a multivariate macroeconomic series. We also find consistent evidence that deeper architectures provide better prediction performance, another result which contradicts [Gu et al. \(2018\)](#)'s conclusions that shallow learning may be better.

Focusing on the causal inference capabilities of different models by looking at variable importance metrics however reveals a more interesting comparison between models. We can clearly observe that the elastic net outperforms all other models consistently in terms of assigning the correct relative importance to the true causal regressors ⁶, even in settings with very high cross sectional correlation.

In terms of finding the correct underlying causal regressors, we find that the penalized linear models perform the best at identifying the true data generating regressors, and that this appears to be mostly robust regardless of the amount of cross sectional correlation in the factor set. The penalized linear models are not perfect and their ability to correctly identify causal regressors worsens as the data generating process becomes more non-linear. On these more difficult specifications, the penalized linear models are generally very conservative, sometimes only identifying a single covariate as important. This is most apparent on the g_2 specification. Occasionally, the elastic nets identified the incorrect covariates completely, though the relative importance assigned to them was small.

The random forests and to a lesser extent the neural networks also correctly identified the correct causal regressors, but struggled with adequately discerning relative importance among correlated regressors. This was became more apparent as the degree of cross sectional correlation increased (see decreasing relative importance of true causal regressors in Figures 3.2 and 3.2). In the case of the random forests, this is to be expected, likely due to how the random forest algorithms work. The random forest algorithm is an ensemble of tree models, with each tree model only having access to a subset of all available predictors. If this subset does not include

⁶(c_1 .constant, c_2 .constant and $c_3.x_3$ for g_1 and g_2 specifications, and c_1 .constant, c_2 .constant and c_3 .constant for g_3)



Figure 3. Simulation Test Loss Metrics, only including best performing Neural Networks




Figure 4. Simulation Test MAE, only including best performing Neural Networks




Figure 5. Simulation Test MSE, only including best performing Neural Networks

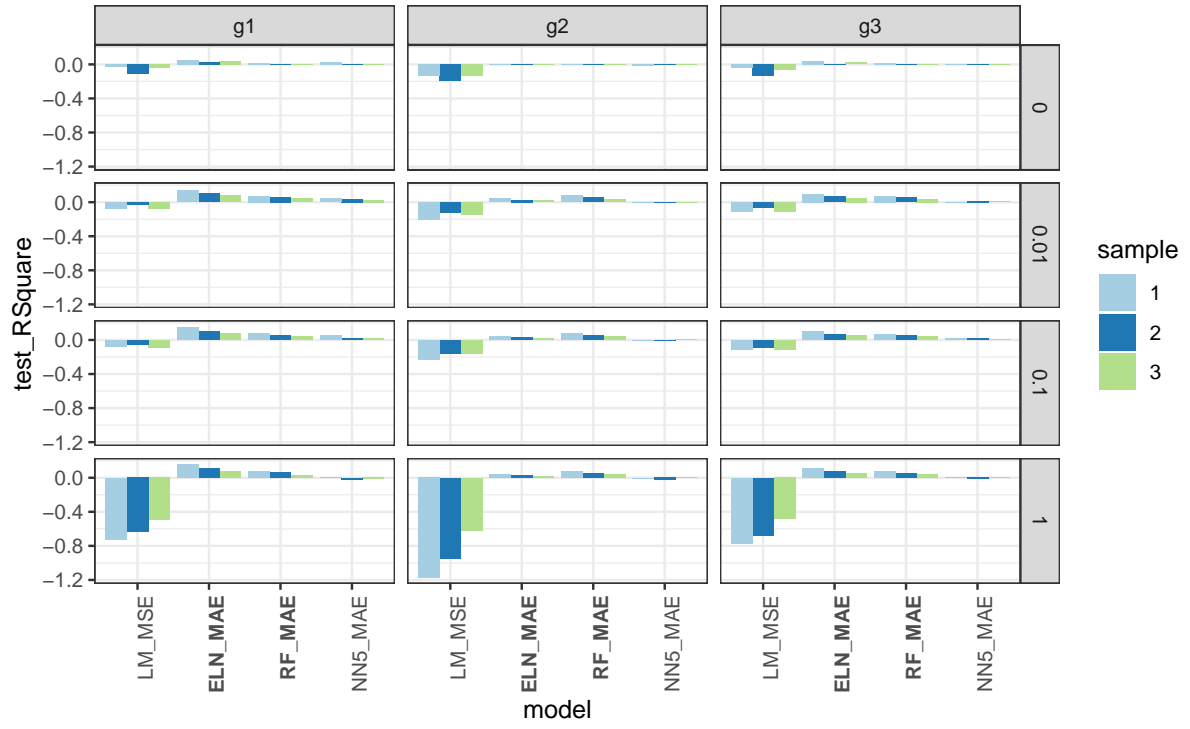


Figure 6. Simulation Test OOS R Squared, only including best performing Neural Networks

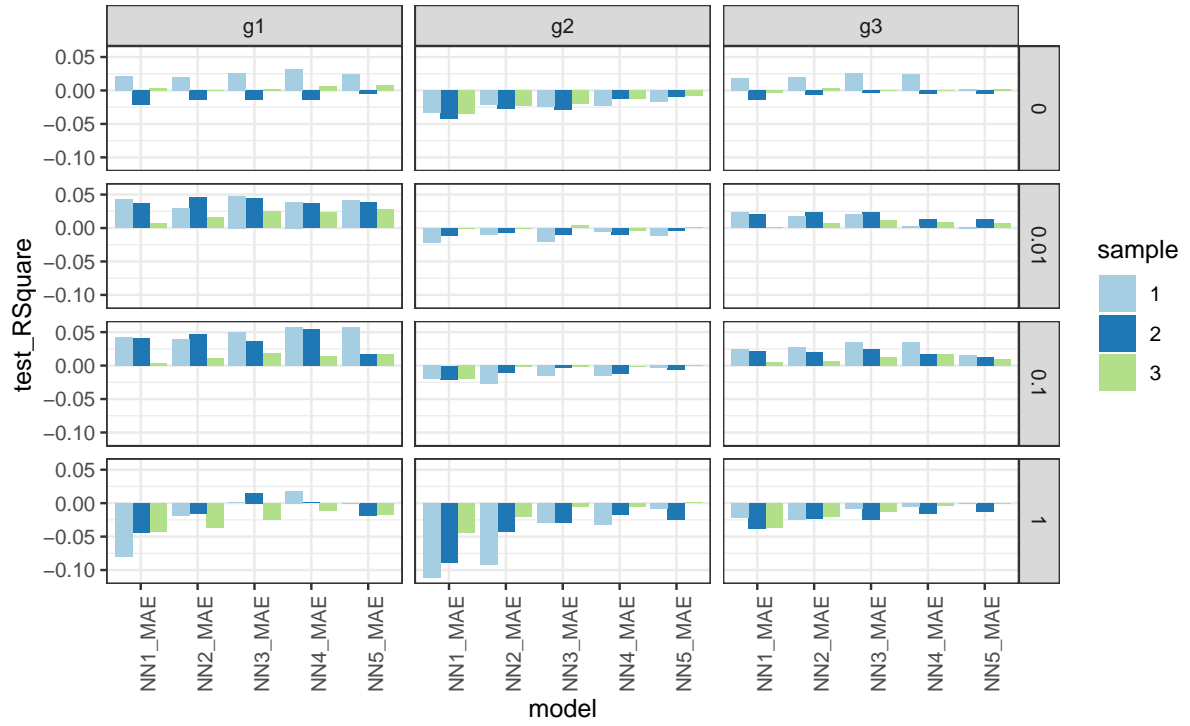



Figure 7. Simulation Test OOS R-Squared for Neural Networks optimized w.r.t. MAE



the true data generating predictor, then that particular tree will likely select the predictors which have the highest correlation with the true data generating predictor instead. Thus, the resulting ensemble model is likely to believe that cross sectionally correlated predictors are important, relative to the true causal regressor. Due to the complexity of the neural networks, there does not exist a similar intuitive explanation for their causal inference ability.

The linear models, unsurprisingly, struggled with causal analysis with respect to both increasing cross sectional correlation and increasing non-linearities. This highlights the non-robustness and ineffectiveness of using traditional linear regression as documented by the literature; linear models were consistently observed to identify non-causal regressors as important, especially as the degree of cross sectional correlation increased. Considering that the graphs represent the averaged variable importance metrics over different simulation realisations which each have random patterns of cross sectional correlation, this means that on a single simulation realization, the performance of linear models is significantly worse.

Of particular note are the instability of the machine learning models’ hyperparameters across different training samples. For the elastic nets, the optimal value for α is generally 1 (corresponding to LASSO and thus a sparse representation), but it was not uncommon to observe α values swinging between values close to 0 (corresponding to ridge regression, and thus a dense representation) to 1 as the training sample moved forwards in time. As the penalized linear models consistently performed the best and still remained able to correctly identify the true




Figure 9. Simulation Variable Importance Plots for Random Forests




Figure 10. Simulation Variable Importance Plots for Neural Networks

covariates this is not a large issue, but it should be noted that this can lead to interpretation issues. For the random forests, it was similarly observed that the optimal value for *mtry* (the number of variables subsetted) and *nodesize* was highly non-robust. Again, given that the final prediction performance was consistent this is not a large issue, but can lead to some interpretation issues.

4. Empirical Study

We then begin the empirical study to validate and check the robustness of the results observed in the simulation study. Though our simulation study aimed at capturing the main characteristics of empirical data, the underlying data generating process for empirical returns is unknown this thus as a robustness check as to how machine learning performs on real world data, which can be significantly more complex and noisy than simulated contexts. This also acts as a final validation against what has been reported in the literature.

Importantly, we find that our results from the simulation study are largely repeated, suggesting a degree of robustness for our simulation study results.

4.1. Data

We mimic the data procedure of [Gu et al. \(2018\)](#). This means that we obtain the dataset provided by Gu on his website. This dataset sample begins in March 1957 (the start date of the S&P 500) and ends in December 2016, totalling 60 years. It contains 94 stock level characteristics: 61 updated annually, 13 updated quarterly and 20 updated monthly, in addition to 74 industry dummies corresponding the the first two digits of the Standard Industrial Classification (SIC) codes. It is noted that this dataset contains all securities traded, including those with a CRSP share code other than 10 or 11 and thus includes instruments such as REITs and mutual funds, and those with a share price of less than \$5.

We detail our cleaning procedure of this dataset. To reduce the size of the dataset and increase feasibility, the dataset was filtered such that only stocks traded primarily on NASDAQ were included (using the PRIMEXCH variable from WRDS). Then, penny stocks (also referred to as microcaps in the literature) with a stock price of less than \$5 were filtered out, as is commonly done in the literature to reduce variability. Stocks without a share code of 10 or 11 (referring to equities) were filtered out, so that securities that are not equities were not included (such as REITs and trust funds). The dataset is provided in a monthly format, which means that many of the factors which are updated only quarterly or annually have very low levels of variability, which can lead to misleading results in the model fitting process. To achieve a balance between having a dataset with enough data points and variability among factors, the dataset was converted to a quarterly format. Quarterly returns were then constructed using the PRC variable according to actual returns (ie not logged differences):

$$RET_t = \frac{PRC_t - PRC_{t-1}}{PRC_{t-1}} \quad (32)$$

We allow all stocks which have a quarterly return to enter the dataset, even if they disappear from the dataset for certain periods, as opposed to only keeping stocks which appear continuously throughout the entire period. This was primarily done to reduce survivorship bias in the dataset, which can be very prevalent in financial data, and also allows for stocks which were unlisted and relisted again to feature in the dataset. This has the obvious drawback of

introducing some bias in the dataset, as attrition in the dataset is likely to be non-random and correlated with the stocks' returns.

The sic2 variable, corresponding to the stocks' Standard Industrial Classification (SIC) codes was also dropped. The SIC code system suffers from inconsistent logic in classifying companies, and as a system built for pre-1970s traditional industries has been slow in recognizing new and emerging industries. Indeed, WRDS explicitly cautions the use of SIC codes beyond the use of rough grouping of industries, warning that SIC codes are not strictly enforced by government agencies for accuracy, in addition to most large companies belonging to multiple SIC codes over time. Because of this latter point in particular, there can be inconsistencies on the correct SIC code for the same company depending on the data source. Dropping the sic2 variable also reduced the dimensionality of the dataset by 74 columns, significant increasing computational feasibility.

There existed a significant amount of missing data in the dataset. The dataset's columns were first examined, and any characteristics that had over 20% of their data were removed. However, as the amount of missing data increases dramatically going further back in time, a balance between using more periods at the cost of removing more characteristics versus using less periods but keeping more characteristics was needed. 1993 Q3 was determined to be a reasonable time frame to begin the dataset, as there was a noticeable increase in data availability and quality after this time. Missing characteristics were then imputed using their cross sectional medians for each year.

We then follow [Gu et al. \(2018\)](#) and construct eight macroeconomic factors following the variable definitions in [Welch and Goyal \(2008\)](#): dividend-price ratio (dp), earnings-price ratio (ep), book-to-market ratio (bm), net equity expansion (ntis), Treasury-bill rate (tbl), term spread (tms), default spread (dfy) and stock variance (svar). These factors were lagged by one period so as to be used to predict one period ahead quarterly returns. The treasury bill rate was also used from this source to proxy for the risk free rate in order to construct excess quarterly returns.

The two sets of factors were then combined to form a baseline set of covariates, which we define throughout all methods and analysis as:

$$z_{i,t} = (1, x_t)' \otimes c_{i,t} \quad (33)$$

where $c_{i,t}$ is a P_c matrix of characteristics for each stock i , and $(1, x_t)'$ is a $P_x \times 1$ vector of macroeconomic predictors. $z_{i,t}$ is therefore a $P_x P_c$ vector of features for predicting individual stock returns and includes interactions between stock level characteristics and macroeconomic variables. The total number of covariates in this baseline set is $61 \times (8 + 1) = 549$ ⁷.

The dataset was not normalized for all methods, as only penalized regression and neural networks are sensitive to normalization. For these two methods, the dataset was normalized

⁷As the individual and macroeconomic factors can have similar names, individual and macroeconomic factors were prefixed with ind_ and macro_ respectively.




Figure 11. Empirical Data Sample Splitting Procedure

such that each predictor column had 0 mean and 1 variance.

The final dataset spanned from 1993 Q3 to 2016 Q4 with 202066 individual observations.

We mimic the procedure used in the simulation study. This means that the dataset was split such that the training and validation sets were split such that the training set was approximately 1.5 times the length of the validation set, in order to predict a test set that is one year in length.




Figure 12. Empirical Test MSE Averaged Across All Samples

4.2. Empirical Data Results

In general, results from the simulation study were reproduced in the empirical study. This in of itself is a significant result, as it suggests that the simulation design itself is a reasonable approximation to empirical data, and that results from the simulation study are robust.

We similarly see that the penalized linear models generally performing the best, with the random forest models offering slightly worse performance, occasionally outperforming penalized linear models. Machine learning models fitted with respect to median quantile loss were similarly observed to typically offer improvements across all machine learning models across all loss metrics.

However, on the empirical data we start to see the neural networks start to fail.

In terms of prediction accuracy, we can see that in general the results of the simulation study were repeated: the elastic net models perform the best, followed by the random forests, then the neural networks, and finally the linear models. We note that the differences between each model using the MSE and MAE loss metrics are much more pronounced on empirical data. We similarly see that machine learning models perform better when fitted with respect to quantile loss instead of MSE.

Focusing on the neural networks specifically, their non-robustness is amplified on the empirical dataset, with some neural networks in some samples even performing worse than linear models. This was observed to be somewhat more common on neural networks fitted with re-




Figure 13. Empirical Test MAE Averaged Across All Samples

spect to MSE, suggesting that they are indeed very sensitive to outliers in training data. We similarly observe some evidence that deeper neural networks perform better, though this result is less apparent to the lower robustness on empirical data.

Interestingly, we do not observe worsening performance as the training sample increases as we did in the simulation study. This suggests that the simulation design may have been too volatile when compared to the specific empirical time periods examined.

As the data generating process for empirical returns is unknown, the variable importance results cannot be directly compared with those of the simulation study. Even so, we see similar results: the elastic net and random forest models tend to agree on the same subset of predictors, but the random forest struggles with adequately discerning between highly correlated regressors. However, we see similarly to the prediction performance results that the neural networks become less robust.

If we focus on the two top performing models of elastic net and random forest, we see that they consistently pick out the 1 month and 6 month momentum factors out of the individual characteristics as important, and the book-to-market and default yield spread factors out of the macroeconomic factors are important. In general though, the variable importance metrics are less consistent for the random forests, and it should be noted in particular that the random forest tends to determine factors highly correlated with momentum as important, such as change in moment, dollar trading volume and return volatility. Looking at the macroeconomic factors,




Figure 15. Empirical Individual Factor Importance, averaged across all training samples




Figure 16. Empirical Macroeconomic Factor Importance, faceted by training sample

models, or non-robust.

All models considered typically preferred sparse parameterizations. That is, most if not all of the individual factors had little to no importance across all models.⁸

⁸Note that because the variable importance here was not evaluated explicitly for each pairwise interaction term, some of the individual factors appear as slightly important. This is because setting an individual factor to zero also sets some of the macroeconomic pairwise terms to zero, increasing its apparent importance.

5. Conclusion

Our findings demonstrate that within the contexts considered, the field of machine learning offers tools to improve stock prediction and identification of true underlying regressors. Penalized linear models and to a lesser extent, random forests are the best performing methods.

Importantly, we find that the feed-forward neural network architectures considered fail in the context of stock return prediction, at both prediction performance and variable importance analysis. This is a result consistent across a variety of simulated datasets, as well as empirical data. We also find weak evidence that deeper neural networks with more hidden layers perform better, though this result, like the performance of neural networks in general, is non-robust.

Lastly, we find that the top performing models - the penalized linear models and random forests, tend to agree and correctly identify the correct causal regressors in simulated contexts, as well as agree on the same subset of factors which are important in empirical contexts. We find that of all the most considered, the penalized linear models are the most consistent at identifying true causal regressors through the simulation study. We find that in the empirical setting, among the individual factors the individual 1 and 6 month momentum factors are the most powerful predictors of stock returns, according to the penalized linear models and random forests.

Across all models except for linear models, we find that minimizing quantile loss yields better prediction performance.

The overall findings of this paper differ from the sparse literature on similar topics. However, the performance of the penalized linear models with respect to both out of sample prediction performance and variable importance analysis is promising, and our findings show that machine learning provides some tools which may aid in the problems of stock return prediction and risk factor selection in the financial world.

6. Appendix

6.1. *Data*

Table 1: Individual Factors

No.	Acronym	Firm Characteristic	Author(s)	Data Source	Frequency
1	ind_absacc ⁸	Absolute Accruals	Bandyopadhyay et al. (2010)	Compustat	Annual
2	ind_acc ⁸	Working capital accruals	Sloan (1996)	Compustat	Annual
3	ind_aeavol	Abnormal earnings announcement volume	Lerman et al. (2008)	Compustat	Quarterly
4	ind_age	# years since first Compustat coverage	Jiang et al. (2005)	Compustat	Annual
5	ind_agr	Asset growth	Cooper et al. (2008)	Compustat	Annual
6	ind_baspread	Bid-ask spread	Amihud and Mendelson (1989)	Compustat	Monthly
7	ind_beta	Beta	Fama and MacBeth (1973)	Compustat	Monthly
8	ind_betasq	Beta squared	Fama and MacBeth (1973)	Compustat	Monthly
9	ind_bm	Book-to-market	Rosenberg et al. (1985)	Compustat	Annual
10	ind_bm_ia	Industry-adjusted book to market	Asness et al. (2000)	Compustat	Quarterly
11	ind_cash	Cash holdings	Palazzo (2012)	Compustat	Annual
12	ind_cashdebt	Cashflow to debt	Ou and Penman (1989)	Compustat	Annual
13	ind_cashpr	Cash productivity	Chandrashekar and Rao (2009)	Compustat	Annual
14	ind_cfp ⁸	Cashflow to price ratio	Desai et al. (2004)	Compustat	Annual
15	ind_cfp_ia ⁸	Industry-adjusted cashflow to price ratio	Asness et al. (2000)	Compustat	Annual
16	ind_chatoia ⁸	Industry-adjusted change in asset turnover	Soliman (2008)	Compustat	Annual
17	ind_chesho	Change in shares outstanding	Pontiff and Woodgate (2008)	Compustat	Annual
18	ind_chempia	Industry-adjusted change in employee	Asness et al. (2000)	Compustat	Annual
19	ind_chinv ⁸	Change in inventory	Thomas and Zhang (2002)	Compustat	Annual
20	ind_chmom	Change in 6-month momentum	Gettleman and Marks (2006)	Compustat	Monthly
21	ind_chpmia ⁸	Industry-adjusted change in profit margin	Soliman (2008)	Compustat	Annual
22	ind_chctx	Change in tax expense	Thomas and Zhang (2011)	Compustat	Quarterly

⁸The factor was included in the original dataset provided by Gu et al. (2018), but was not used due to missing data issues

23	ind_cinvest	Corporate investment	Titman et al. (2004)	Compustat	Quarterly
24	ind_convind	Convertible debt indicator	Valta (2016)	Compustat	Annual
25	ind_currat	Current ratio	Ou and Penman (1989)	Compustat	Annual
26	ind_depr	Depreciation / PP&E	Holthausen and Larcker (1992)	Compustat	Annual
27	ind_divi	Dividend initiation	Michaely et al. (1995)	Compustat	Annual
28	ind_divo	Dividend omission	Michaely et al. (1995)	Compustat	Annual
29	ind_dolvol	Dollar trading volume	Chordia et al. (2001)	Compustat	Monthly
30	ind_dy	Dividend to price	Litzenberger and Ramaswamy (1982)	Compustat	Annual
31	ind_ear	Earnings announcement return	Brandt et al. (2008)	Compustat	Quarterly
32	ind_egr	Growth in common shareholder eq	Richardson et al. (2005)	Compustat	Annual
33	ind_ep	Earnings to price	Basu (1977)	Compustat	Annual
34	ind_gna	Gross profitability	Novy-Marx (2013)	Compustat	Annual
35	ind_grCAPX ⁸	Growth in capital expenditures	Anderson and Garcia-Feijóo (2006)	Compustat	Annual
36	ind_grltnoa ⁸	Growth in long term net operating assets	Fairfield et al. (2003)	Compustat	Annual
37	ind_herf	Industry sales concentration	Hou and Robinson (2006)	Compustat	Annual
38	ind_hire	Employee growth rate	Belo et al. (2014)	Compustat	Annual
39	ind_idiovol	Idiosyncratic return volatility	Ali et al. (2003)	Compustat	Monthly
40	ind_ill	Illiquidity	Amihud (2002)	Compustat	Monthly
41	ind_indmom	Industry momentum	Moskowitz and Grinblatt (1999)	Compustat	Monthly
42	ind_invest ⁸	Capital expenditures and inventory	Chen and Zhang (2010)	Compustat	Annual
43	ind_lev	Leverage	Bhandari (1988)	Compustat	Annual
44	ind_lgr	Growth in long-term debt	Richardson et al. (2005)	Compustat	Annual
45	ind_maxret	Maximum daily return	Bali et al. (2011)	Compustat	Monthly
46	ind_mom12m	12-month momentum	Jegadeesh (1990)	Compustat	Monthly
47	ind_mom1	1-month momentum	Jegadeesh and Titman (1993)	Compustat	Monthly

48	ind_mom36m ⁸	36-month momentum	Jegadeesh and Titman (1993)	Compustat	Monthly
49	ind_mom6m	6-month momentum	Jegadeesh and Titman (1993)	Compustat	Monthly
50	ind_ms	Financial statement score	Mohanram (2005)	Compustat	Quarterly
51	ind_mvell	Size	Banz (1981)	Compustat	Monthly
52	ind_mv_ia	Industry-adjusted size	Asness et al. (2000)	Compustat	Annual
53	ind_nincr	Number of earnings increases	Barth et al. (1999)	Compustat	Quarterly
54	ind_operprof ⁸	Operating profitability	Fama and French (2015)	Compustat	Annual
55	ind_orgcap ⁸	Organizational capital	Eisfeldt and Papanikolaou (2013)	Compustat	Annual
56	ind_pchcapx_ia ⁸	Industry adjusted % change in capital expenditures	Abarbanell and Bushee (1998)	Compustat	Annual
57	ind_pchcurrat ⁸	% change in current ratio	Ou and Penman (1989)	Compustat	Annual
58	ind_pchdepr ⁸	% change in depreciation	Holthausen and Larcker (1992)	Compustat	Annual
59	ind_pchgm_pchsale ⁸	% change in gross margin - % change in sales	Abarbanell and Bushee (1998)	Compustat	Annual
60	ind_pchquick ⁸	% change in quick ratio	Ou and Penman (1989)	Compustat	Annual
61	ind_pchsale_pchinv ⁸	% change in sales - % change in inventory	Abarbanell and Bushee (1998)	Compustat	Annual
62	ind_pchsale_pchrect ⁸	% change in sales - % change in A/R	Abarbanell and Bushee (1998)	Compustat	Annual
63	ind_pchsale_pchxsga ⁸	% change in sales - % change in SG	Abarbanell and Bushee (1998)	Compustat	Annual
64	ind_pchsaleinv ⁸	% change sales-to-inventory	Ou and Penman (1989)	Compustat	Annual
65	ind_pctacc ⁸	Percent accruals	Hafzalla et al. (2011)	Compustat	Annual
66	ind_pricedelay	Price delay	Hou and Moskowitz (2005)	Compustat	Monthly
67	ind_ps	Financial statements score	Piotroski (2000)	Compustat	Annual
68	ind_quick	Quick ratio	Ou and Penman (1989)	Compustat	Annual
69	ind_rd	R&D increase	Eberhart et al. (2004)	Compustat	Annual
70	ind_rd_mv ⁸	R&D to market capitalization	Guo et al. (2006)	Compustat	Annual
71	ind_rd_sale ⁸	R&D to sales	Guo et al. (2006)	Compustat	Annual
72	ind_realestate ⁸	Real estate holdings	Tuzel (2010)	Compustat	Annual

73	ind_retvol	Return volatility	Ang et al. (2006)	Compustat	Monthly
74	ind_roaq	Return on assets	Balakrishnan et al. (2010)	Compustat	Quarterly
75	ind_roavol ⁸	Earnings volatility	Francis et al. (2004)	Compustat	Quarterly
76	ind_roeq	Return on equity	Hou et al. (2015)	Compustat	Quarterly
77	ind_roic	Return on invested capital	Brown and Rowe (2007)	Compustat	Annual
78	ind_rsup	Revenue surprise	Kama (2009)	Compustat	Quarterly
79	ind_salecash	Sales to cash	Ou and Penman (1989)	Compustat	Annual
80	ind_saleinv ⁸	Sales to inventory	Ou and Penman (1989)	Compustat	Annual
81	ind_salerec	Sales to receivables	Ou and Penman (1989)	Compustat	Annual
82	ind_secured ⁸	Secured debt	Valta (2016)	Compustat	Annual
83	ind_securedind	Secured debt indicator	Valta (2016)	Compustat	Annual
84	ind_sgr ⁸	Sales growth	Barbee Jr et al. (1996)	Compustat	Annual
85	ind_sin	Sin stocks	Hong and Kacperczyk (2009)	Compustat	Annual
86	ind_sp	Sales to price	Barbee Jr et al. (1996)	Compustat	Annual
87	ind_std_dolvol	Volatility of liquidity (dollar trading volume)	Chordia et al. (2001)	Compustat	Annual
88	ind_std_turn	Volatility of liquidity (share turnover)	Chordia et al. (2001)	Compustat	Monthly
89	ind_stdacc ⁸	Accrual volatility	Bandyopadhyay et al. (2010)	Compustat	Monthly
90	ind_stdcf ⁸	Cashflow volatility	Huang (2009)	Compustat	Quarterly
91	ind_tang	Debt capacity/rm tangibility	Almeida and Campello (2007)	Compustat	Quarterly
92	ind_tb ⁸	Tax income to book income	Lev and Ohlson (1982)	Compustat	Annual
93	ind_turn	Share turnover	Datar et al. (1998)	Compustat	Monthly
94	ind_zerotrade	Zero trading days	Liu (2006)	Compustat	Monthly

Table 2: Macroeconomic Factors

No.	Acronym	Macroeconomic Factor
1	macro_dp	Dividend Price Ratio
2	macro_ep	Earnings Price Ratio
3	macro_bm	Book to Market Ratio
4	macro_ntis	Net Equity Expansion
5	macro_tbl	Treasury Bill Rate
6	macro_tms	Term Spread
7	macro_dfy	Default Spread
8	macro_svar	Stock Variance

6.2. *Additional Results*

6.2.1. *Simulation Study*




Figure 17. Simulation Test MAE




Figure 18. Simulation Test MSE




Figure 19. Simulation Test R Squared



Figure 20. Simulation Train MAE




Figure 21. Simulation Train MSE




Figure 22. Simulation Train R Squared




Figure 23. Optimal Alpha Values Across Training Samples for Elastic Net, each colour representing a different simulation realisation




Figure 24. Optimal Lambda Values Across Training Samples for Elastic Net, each colour representing a different simulation realisation

6.2.2. Empirical Study



Figure 25. Empirical Data Individual Factor Variable Importance for Sample 1




Figure 26. Empirical Data Individual Factor Variable Importance for Sample 2




Figure 27. Empirical Data Individual Factor Variable Importance for Sample 3

6.3. Computational Details

6.3.1. Linear Models

Ordinary Least Squares (OLS) used for fitting wrt MS.

The **quantreg** package which implements quantile linear regression used for fitting linear models with respect to mean absolute error (noting that minimizing 0.5 quantile loss is equivalent to minimizing mae).

6.3.2. Penalized Linear

The package **hqreg** was used to fit penalized regression models with respect to MSE and MAE.

This package efficiently calculates a regularization path of penalization values given a value for α . This means that it is much more efficient to instead only supply a grid for α , let the algorithm decide its own path of penalization values. The combination of these two parameters which produces the best results on the validation set were then chosen to produce the final model. Note that this is the approach originally suggested by [Zou and Hastie \(2005\)](#).

6.3.3. Classification and Regression Trees

For full details of the Classification and Regression Tree algorithm see [Breiman \(1984\)](#).

Algorithm 1: Classification and Regression Tree

Initialize ;

for d from 1 to L **do**

for i in $C_l(d-1), l = 1, \dots, 2^{d-1}$ **do**

 For each feature $j = 1, 2, \dots, P$, and each threshold level α , define a split as

$s = (j, \alpha)$ which divides $C_l(d-1)$ into C_{left} and C_{right} :

$$C_{left}s = \{z_j \leq \alpha\} \cap C_l(d-1); C_{right}s = \{z_j > \alpha\} \cap C_l(d-1)$$

 Define the impurity function:

$$\mathcal{L}(C, C_{left}, C_{right}) = \frac{|C_{left}|}{|C|} H(C_{left}) + \frac{|C_{right}|}{|C|} H(C_{right})$$

 where

$$H(C) = \frac{1}{|C|} \sum_{z_{i,t} \in C} (r_{i,t+1} - \theta)^2, \theta = \frac{1}{|C|} \sum_{z_{i,t} \in C} r_{i,t+1}$$

 and $|C|$ denotes the number of observations in set C

 Find the optimal split

$$s^* \leftarrow \underset{s}{\operatorname{argmin}} \mathcal{L}(C(s), C_{left}(s), C_{right}(s))$$

 Update nodes (partition the data):

$$C_{2l-1}(d) \leftarrow C_{left}(s^*), C_{2l}(d) \leftarrow C_{right}(s^*)$$

end

end

Result: The prediction of a regression tree is:

$$g(z_{i,t}; \theta, L) = \sum_{k=1}^{2^L} \theta_k \mathbf{1}_{z_{i,t} \in C_k(L)}; \theta_k = \frac{1}{|C_k(L)|} \sum_{z_{i,t} \in C_k(L)} r_{i,t+1}$$

6.3.4. Random Forest

The **randomforestSRC** package was used to fit both random forests for mean regression and quantile regression.

To maintain computational feasibility, all random forest models were grown using 50 trees, and their *mtry* *nodesize* hyperparameters tuned. These correspond to the number of predictors

Hyperparameter	Options Considered	Optimal Choice
Activation Function	ReLU, Leaky ReLU with $\alpha = (0.01, 0.1)$, tanh	tanh
Optimizer	ADAM, NADAM	ADAM
Learning Rate	(0.1, 0.01, 0.001, 0.0001)	0.01
L1 Penalty	(1, 0.1, 0.01)	0.01
Batch Size	(32, 64, 128, 256, 512, 1024, 2048, 4096)	256 (Empirical), 128 (Simulated)
Early Stopping Patience	(10, 20, 30, 40, 50)	40

Table 3: Hyperparameters Considered for Neural Networks

included for each tree, and the depth of each tree respectively.

Rather than use the randomly determined out of bag (OOB) observations for hyperparameter tuning, we explicitly use a separate validation set that is closer to the test set in time, in order to maintain temporal ordering of the data.

For full details of the Random Forest algorithm see [Breiman \(2001\)](#).

Algorithm 2: Random Forest

for b from 1 to B **do**

Draw bootstrap samples $(z_{i,t}, r_{i,t+1}), (i, t) \in \text{Bootstrap}(b)$ from the dataset Grow a tree T_b using Algorithm, using only a random subsample, say \sqrt{P} of all features
Denote the resulting b th tree as

$$\hat{g}_b(z_{i,t}, \hat{\theta}_b, L) = \sum_{k=1}^{2^L} \theta_b^k \mathbf{1}_{z_{i,t} \in C_k(L)}$$

end

Result: The final random forest prediction is given by the output of all trees:

$$\hat{g}_b(z_{i,t}; L, B) = \frac{1}{B} \sum_{b=1}^B \hat{g}_b(z_{i,t}, \hat{\theta}_b, L)$$

6.3.5. Neural Networks

Neural Networks were fit using the **keras** package, using the **tensorflow** backend.

Due to computational limitations, a systematic grid search or random search approach for hyperparameters was not feasible. Instead, common hyperparameters that are usually grid searched were tuned manually.

ADAM algorithm for stochastic gradient descent and learning rate shrinkage as detailed by [Kingma and Ba \(2014\)](#).

Batch Normalization Algorithm as detailed by [Ioffe and Szegedy \(2015\)](#).

Algorithm 3: Early stopping via validation

Initialize $j = 0$, $\epsilon = \infty$ and select the patience parameter p (max iterations)
while $j \leq p$ **do**
 Update θ using the training algorithm Calculate the prediction error from the validation sample, denoted as ϵ'
 if $\epsilon' < \epsilon$ **then**
 $j \leftarrow 0$
 $\epsilon \leftarrow \epsilon'$
 $\theta' \leftarrow \theta$
 else
 $j \leftarrow j + 1$
 end
end
Result: θ' is the final parameter estimate

Algorithm 4: Batch Normalization for one activation over one batch

Input: Values of x for each activation over a batch $\mathcal{B} = x_1, x_2, \dots, x_N$
 $\mu_{\mathcal{B}} \leftarrow \frac{1}{N} \sum_{i=1}^N x_i$
 $\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{N} \sum_{i=1}^N (x_i - \mu_{\mathcal{B}})^2$
 $\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$
 $y_i \leftarrow \gamma \hat{x}_i + \beta := BN_{\gamma, \beta}(x_i)$
Result: $y_i = BN_{\gamma, \beta}(x_i) : i = 1, 2, \dots, N$

6.3.6. Tuning of Simulated Datasets

The simulated datasets were tuned according to the following statistics: average individual time series R squared, average annualized volatility, and cross sectional R squared, and true/predictive R squared.

The methodology for evaluating average time series R squared and cross sectional R squared is consistent with that detailed by [Cochrane \(2005\)](#). The steps are reproduced here for reference; for complete details refer to [Cochrane \(2005\)](#).

First evaluate the following OLS model:

$$R_{it} = \alpha_i + \beta_i' f_{it} + \epsilon_{it} \quad (34)$$

where f_{it} represents the *true* factors in the returns process. The corresponding R-squared value for this time series regression is calculated across all stocks and averaged to yield the average time series R-Squared.

A cross sectional regression for the risk premia is then run across assets of average returns on the factor coefficients:

$$\bar{R}_{it} = \alpha_i = \beta_i' \lambda \quad (35)$$

where the β_i' are the estimated coefficients from each time series regressions run previously. The corresponding R squared for this regression is the cross sectional R squared.

The true R squared is a measure of signal to noise ratio i.e. how much of the returns data is due to This is simply calculated by running a pooled ordinary least squares regression on the entire panel, using the underlying $g()$ as the "predictions." The resulting R squared value is therefore a measure of how much of the panel can be explained by the $g()$ term exclusive of any noise or error terms.

References

- Abarbanell, J. S., Bushee, B. J., 1998. Abnormal returns to a fundamental analysis strategy. *Accounting Review* pp. 19–45.
- Ali, A., Hwang, L.-S., Trombley, M. A., 2003. Arbitrage risk and the book-to-market anomaly. *Journal of Financial Economics* 69, 355–373.
- Almeida, H., Campello, M., 2007. Financial constraints, asset tangibility, and corporate investment. *The Review of Financial Studies* 20, 1429–1460.
- Amihud, Y., 2002. Illiquidity and stock returns: cross-section and time-series effects. *Journal of financial markets* 5, 31–56.
- Amihud, Y., Mendelson, H., 1989. The effects of beta, bid-ask spread, residual risk, and size on stock returns. *The Journal of Finance* 44, 479–486.
- Anderson, C. W., Garcia-Feijóo, L., 2006. Empirical evidence on capital investment, growth options, and security returns. *The Journal of Finance* 61, 171–194.
- Ang, A., Bekaert, G., 2006. Stock return predictability: Is it there? *The Review of Financial Studies* 20, 651–707.
- Ang, A., Hodrick, R. J., Xing, Y., Zhang, X., 2006. The Cross-Section of Volatility and Expected Returns. *The Journal of Finance* 61, 259–299.
- Asness, C. S., Porter, R. B., Stevens, R. L., 2000. Predicting stock returns using industry-relative firm characteristics. Available at SSRN 213872 .
- Balakrishnan, K., Bartov, E., Faurel, L., 2010. Post loss/profit announcement drift. *Journal of Accounting and Economics* 50, 20–41.
- Bali, T. G., Cakici, N., Whitelaw, R. F., 2011. Maxing out: Stocks as lotteries and the cross-section of expected returns. *Journal of Financial Economics* 99, 427–446.
- Bandyopadhyay, S. P., Huang, A. G., Wirjanto, T. S., 2010. The accrual volatility anomaly. Unpublished Manuscript, University of Waterloo .
- Banz, R. W., 1981. The relationship between return and market value of common stocks. *Journal of financial economics* 9, 3–18.

- Barbee Jr, W. C., Mukherji, S., Raines, G. A., 1996. Do sales–price and debt–equity explain stock returns better than book–market and firm size? *Financial Analysts Journal* 52, 56–60.
- Barth, M. E., Elliott, J. A., Finn, M. W., 1999. Market rewards associated with patterns of increasing earnings. *Journal of Accounting Research* 37, 387–413.
- Basu, S., 1977. Investment performance of common stocks in relation to their price-earnings ratios: A test of the efficient market hypothesis. *The journal of Finance* 32, 663–682.
- Belo, F., Lin, X., Bazdresch, S., 2014. Labor hiring, investment, and stock return predictability in the cross section. *Journal of Political Economy* 122, 129–177.
- Bengio, Y., Simard, P., Frasconi, P., 1994. Learning long-term dependencies with gradient descent is difficult. *IEEE transactions on neural networks* 5, 157–166.
- Bhandari, L. C., 1988. Debt/equity ratio and expected common stock returns: Empirical evidence. *The journal of finance* 43, 507–528.
- Brandt, M. W., Kishore, R., Santa-Clara, P., Venkatachalam, M., 2008. Earnings announcements are full of surprises. SSRN eLibrary .
- Breiman, L., 1984. *Classification and Regression Trees*. Routledge.
- Breiman, L., 2001. Random forests. *Machine learning* 45, 5–32.
- Brown, D. P., Rowe, B., 2007. The productivity premium in equity returns. Available at SSRN 993467 .
- Chandrashekar, S., Rao, R. K., 2009. The productivity of corporate cash holdings and the cross-section of expected stock returns. *McCombs Research Paper Series No. FIN-03-09* .
- Chen, L., Zhang, L., 2010. A better three-factor model that explains more anomalies. *Journal of Finance* 65, 563–595.
- Chordia, T., Subrahmanyam, A., Anshuman, V. R., 2001. Trading activity and expected stock returns. *Journal of Financial Economics* 59, 3–32.
- Cochrane, J. H., 2005. *Asset pricing*. Princeton Univ. Press, Princeton, NJ, rev. ed ed., oCLC: 265653065.
- Cochrane, J. H., 2011. Presidential Address: Discount Rates. *The Journal of Finance* 66, 1047–1108.
- Cooper, M. J., Gulen, H., Schill, M. J., 2008. Asset Growth and the Cross-Section of Stock Returns. *The Journal of Finance* 63, 1609–1651.
- Datar, V. T., Naik, N. Y., Radcliffe, R., 1998. Liquidity and stock returns: An alternative test. *Journal of Financial Markets* 1, 203–219.

- Desai, H., Rajgopal, S., Venkatachalam, M., 2004. Value-glamour and accruals mispricing: One anomaly or two? *The Accounting Review* 79, 355–385.
- Eberhart, A. C., Maxwell, W. F., Siddique, A. R., 2004. An examination of long-term abnormal stock returns and operating performance following R&D increases. *The Journal of Finance* 59, 623–650.
- Eisfeldt, A. L., Papanikolaou, D., 2013. Organization capital and the cross-section of expected returns. *The Journal of Finance* 68, 1365–1406.
- Fairfield, P. M., Whisenant, J. S., Yohn, T. L., 2003. Accrued earnings and growth: Implications for future profitability and market mispricing. *The accounting review* 78, 353–371.
- Fama, E. F., French, K. R., 2015. A five-factor asset pricing model. *Journal of financial economics* 116, 1–22.
- Fama, E. F., MacBeth, J. D., 1973. Risk, return, and equilibrium: Empirical tests. *Journal of political economy* 81, 607–636.
- Feng, G., Giglio, S., Xiu, D., 2019. Taming the Factor Zoo: A Test of New Factors. Tech. Rep. w25481, National Bureau of Economic Research, Cambridge, MA.
- Feng, G., He, J., Polson, N. G., 2018. Deep Learning for Predicting Asset Returns. arXiv:1804.09314 [cs, econ, stat] ArXiv: 1804.09314.
- Francis, J., LaFond, R., Olsson, P. M., Schipper, K., 2004. Costs of equity and earnings attributes. *The accounting review* 79, 967–1010.
- Freyberger, J., Neuhierl, A., Weber, M., 2017. Dissecting characteristics nonparametrically. Tech. rep., National Bureau of Economic Research.
- Gettleman, E., Marks, J. M., 2006. Acceleration strategies. Tech. rep., Working Paper, Seton Hall University.
- Goetzmann, W. N., Jorion, P., 1993. Testing the predictive power of dividend yields. *The Journal of Finance* 48, 663–679.
- Goyal, A., Welch, I., 2003. Predicting the equity premium with dividend ratios. *Management Science* 49, 639–654.
- Gu, S., Kelly, B., Xiu, D., 2018. Empirical asset pricing via machine learning. Tech. rep., National Bureau of Economic Research.
- Guo, R.-J., Lev, B., Shi, C., 2006. Explaining the Short-and Long-Term IPO Anomalies in the US by R&D. *Journal of Business Finance & Accounting* 33, 550–579.
- Hafzalla, N., Lundholm, R., Matthew Van Winkle, E., 2011. Percent accruals. *The Accounting Review* 86, 209–236.

- Harvey, C. R., Liu, Y., 2019. A Census of the Factor Zoo. Social Science Research Network p. 7.
- Harvey, C. R., Liu, Y., Zhu, H., 2016. ... and the Cross-Section of Expected Returns. *The Review of Financial Studies* 29, 5–68.
- Hastie, T., Tibshirani, R., Friedman, J. H., 2009. The elements of statistical learning: data mining, inference, and prediction. Springer series in statistics, Springer, New York, NY, second ed.
- Hoerl, A. E., Kennard, R. W., 1970. Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics* 12, 55–67.
- Holthausen, R. W., Larcker, D. F., 1992. The prediction of stock returns using financial statement information. *Journal of accounting and economics* 15, 373–411.
- Hong, H., Kacperczyk, M., 2009. The price of sin: The effects of social norms on markets. *Journal of Financial Economics* 93, 15–36.
- Hornik, K., Stinchcombe, M., White, H., 1989. Multilayer feedforward networks are universal approximators. *Neural networks* 2, 359–366.
- Hou, K., Moskowitz, T. J., 2005. Market frictions, price delay, and the cross-section of expected returns. *The Review of Financial Studies* 18, 981–1020.
- Hou, K., Robinson, D. T., 2006. Industry concentration and average stock returns. *The Journal of Finance* 61, 1927–1956.
- Hou, K., Xue, C., Zhang, L., 2015. Digesting Anomalies: An Investment Approach. *Review of Financial Studies* 28, 650–705.
- Hsu, J., Kalesnik, V., 2014. Finding smart beta in the factor zoo. *Research Affiliates* (July) .
- Huang, A. G., 2009. The cross section of cashflow volatility and expected stock returns. *Journal of Empirical Finance* 16, 409–429.
- Ioffe, S., Szegedy, C., 2015. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *arXiv preprint arXiv:1502.03167* .
- Jegadeesh, N., 1990. Evidence of predictable behavior of security returns. *The Journal of finance* 45, 881–898.
- Jegadeesh, N., Titman, S., 1993. Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of finance* 48, 65–91.
- Jiang, G., Lee, C. M., Zhang, Y., 2005. Information uncertainty and expected returns. *Review of Accounting Studies* 10, 185–221.

- Kama, I., 2009. On the market reaction to revenue and earnings surprises. *Journal of Business Finance & Accounting* 36, 31–50.
- Keskar, N. S., Mudigere, D., Nocedal, J., Smelyanskiy, M., Tang, P. T. P., 2016. On Large-Batch Training for Deep Learning: Generalization Gap and Sharp Minima. arXiv:1609.04836 [cs, math] ArXiv: 1609.04836.
- Kingma, D. P., Ba, J., 2014. Adam: A method for stochastic optimization. arXiv preprint arXiv:1412.6980 .
- Kozak, S., Nagel, S., Santosh, S., 2017. Shrinking the cross section. Tech. rep., National Bureau of Economic Research.
- Lecun, Y., Bengio, Y., Hinton, G., 2015. Deep learning. *Nature* 521, 436–444.
- Lerman, A., Livnat, J., Mendenhall, R. R., 2008. The High-Volume Return Premium and Post-Earnings Announcement Drift. *SSRN Electronic Journal* .
- Lettau, M., Ludvigson, S., 2001. Consumption, Aggregate Wealth, and Expected Stock Returns. *The Journal of Finance* 56, 815–849.
- Lev, B., Ohlson, J. A., 1982. Market-based empirical research in accounting: A review, interpretation, and extension. *Journal of Accounting research* pp. 249–322.
- Litzenberger, R. H., Ramaswamy, K., 1982. The effects of dividends on common stock prices tax effects or information effects? *The Journal of Finance* 37, 429–443.
- Liu, W., 2006. A liquidity-augmented capital asset pricing model. *Journal of financial Economics* 82, 631–671.
- Masters, T., 1993. *Practical Neural Network Recipes in C++*. Morgan Kaufmann, google-Books-ID: 7Ez_Pq0sp2EC.
- Michael, R., Thaler, R. H., Womack, K. L., 1995. Price reactions to dividend initiations and omissions: Overreaction or drift? *The Journal of Finance* 50, 573–608.
- Mohanram, P. S., 2005. Separating winners from losers among lowbook-to-market stocks using financial statement analysis. *Review of accounting studies* 10, 133–170.
- Moskowitz, T. J., Grinblatt, M., 1999. Do industries explain momentum? *The Journal of finance* 54, 1249–1290.
- Novy-Marx, R., 2013. The other side of value: The gross profitability premium. *Journal of Financial Economics* 108, 1–28.
- Ou, J. A., Penman, S. H., 1989. Financial statement analysis and the prediction of stock returns. *Journal of accounting and economics* 11, 295–329.

- Palazzo, B., 2012. Cash holdings, risk, and expected returns. *Journal of Financial Economics* 104, 162–185.
- Piotroski, J. D., 2000. Value investing: The use of historical financial statement information to separate winners from losers. *Journal of Accounting Research* 38, 1–52.
- Pontiff, J., Woodgate, A., 2008. Share issuance and cross-sectional returns. *The Journal of Finance* 63, 921–945.
- Ramachandran, P., Zoph, B., Le, Q. V., 2017. Searching for Activation Functions. arXiv:1710.05941 [cs] ArXiv: 1710.05941.
- Rapach, D., Zhou, G., 2013. Forecasting Stock Returns. In: *Handbook of Economic Forecasting*, Elsevier, vol. 2, pp. 328–383.
- Richardson, S. A., Sloan, R. G., Soliman, M. T., Tuna, I., 2005. Accrual reliability, earnings persistence and stock prices. *Journal of accounting and economics* 39, 437–485.
- Rosenberg, B., Reid, K., Lanstein, R., 1985. Persuasive evidence of market inefficiency. *The Journal of Portfolio Management* 11, 9–16.
- Schwert, G. W., 2003. Anomalies and market efficiency. *Handbook of the Economics of Finance* 1, 939–974.
- Sloan, R. G., 1996. Do Stock Prices Fully Reflect Information in Accruals and Cash Flows about Future Earnings? *The Accounting Review* 71, 289–315.
- Soliman, M. T., 2008. The use of DuPont analysis by market participants. *The Accounting Review* 83, 823–853.
- Thomas, J., Zhang, F. X., 2011. Tax expense momentum. *Journal of Accounting Research* 49, 791–821.
- Thomas, J. K., Zhang, H., 2002. Inventory changes and future returns. *Review of Accounting Studies* 7, 163–187.
- Tibshirani, R., 1996. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B (Methodological)* 58, 267–288.
- Titman, S., Wei, K. J., Xie, F., 2004. Capital investments and stock returns. *Journal of financial and Quantitative Analysis* 39, 677–700.
- Tuzel, S., 2010. Corporate real estate holdings and the cross-section of stock returns. *The Review of Financial Studies* 23, 2268–2302.
- Valta, P., 2016. Strategic default, debt structure, and stock returns. *Journal of Financial and Quantitative Analysis* 51, 197–229.

- Welch, I., Goyal, A., 2008. A Comprehensive Look at The Empirical Performance of Equity Premium Prediction. *Review of Financial Studies* 21, 1455–1508.
- Zou, H., Hastie, T., 2005. Regularization and variable selection via the Elastic Net. *Journal of the Royal Statistical Society, Series B* 67, 301–320.